

Class Pointers

**Code for joining Microsoft Teams
for the class:**

Form to be filled

Syllabus

Architecture: computer organization, single-core optimizations including exploiting cache hierarchy and vectorization, parallel architectures including multi-core, shared memory, distributed memory and GPU architectures

Algorithms and Data Structures: algorithmic analysis, overview of trees and graphs, algorithmic strategies, concurrent data structures

Parallelization Principles: motivation, challenges, metrics, parallelization steps, data distribution, PRAM model

Parallel Programming Models and Languages: OpenMP, MPI, CUDA;

Distributed Computing: Commodity cluster and cloud computing;
Distributed Programming: MapReduce/Hadoop model.

Syllabus

Architecture: computer organization, single-core optimizations including exploiting cache hierarchy and vectorization, parallel architectures including multi-core, shared memory, distributed memory and GPU architectures

Algorithms and Data Structures: algorithmic analysis, overview of trees and graphs, algorithmic strategies, concurrent data structures

Parallelization Principles: motivation, challenges, metrics, parallelization steps, data distribution, PRAM model

Parallel Programming Models and Languages: OpenMP, MPI, CUDA;

Distributed Computing: Commodity cluster and cloud computing;
Distributed Programming: MapReduce/Hadoop model.

Reference

Bryant, O'Hallaron. Computer Systems – A Programmer's Perspective

Culler, Singh. Parallel Computing Architecture. A Hardware/Software Approach

Quinn. Parallel Computing. Theory and Practice

Sahni. Data Structures, Algorithms, and Applications in C++

Grama, Gupta, Karypis, Kumar. Introduction to Parallel Computing

Pacheco. An Introduction to Parallel Programming

Hwang, Dongarra, Fox. Distributed and Cloud Computing: From Parallel Processing to the Internet of Things

Lin, Dyer. Data-Intensive Text Processing with MapReduce

Reference

Bryant, O'Hallaron. Computer Systems – A Programmer's Perspective,
Pearson Education Limited 2016, 3rd Global Edition

Culler, Singh. Parallel Computing Architecture. A Hardware/Software
Approach

Quinn. Parallel Computing. Theory and Practice

Sahni. Data Structures, Algorithms, and Applications in C++

Grama, Gupta, Karypis, Kumar. Introduction to Parallel Computing


Pacheco. An Introduction to Parallel Programming

Hwang, Dongarra, Fox. Distributed and Cloud Computing: From
Parallel Processing to the Internet of Things

Lin, Dyer. Data-Intensive Text Processing with MapReduce

Computer Organization: Memory Hierarchy and Cache Memories

Basic Computer Organization




Inside the Processor...

- Control hardware: Hardware to manage instruction execution
- ALU: Arithmetic and Logical Unit (hardware to do arithmetic, logical operations)


Inside the Processor...

- Control hardware: Hardware to manage instruction execution
- ALU: Arithmetic and Logical Unit (hardware to do arithmetic, logical operations)
- Registers: small units of memory to hold data/instructions temporarily during execution



Inside the Processor...

- Control hardware: Hardware to manage instruction execution
- ALU: Arithmetic and Logical Unit (hardware to do arithmetic, logical operations)
- Registers: small units of memory to hold data/instructions temporarily during execution
- Two kinds of registers
 1. Special purpose registers
 2. General purpose registers



*Hardware that
can remember
things*

General Purpose Registers

- Available for use by programmer, possibly for keeping frequently used data
- Why? Since there is a large speed disparity between processor and main memory
 - 2 GHz Processor: 0.5 nanosecond time scale
 - Main memory: \sim 50-100 nsec time scale
- Machine instruction operands can come from registers or from main memory
- But CPUs do not provide a large number of general purpose registers

Problem: Slow Speed of Main Memory

- Main Memory is much slower (around 100x) than the CPU and only a few CPU registers
 - CPU will be waiting for data most of the time
- Solution: Cache Memory
 - Fast memory that is part of CPU
 - Design principle: Locality of Reference
 - Temporal locality: least recently accessed memory locations are least likely to be referenced in the near future
 - Spatial locality: neighbours of recently accessed memory locations are most likely to be referenced in the near future

Memory Address and Cache Composition

- When a CPU refers to a data or instruction, it gives a memory address where the data/instruction is present
- The memory address is used to check if the contents are in the cache.
- But how?


Memory Address and Cache Composition

- Memory is divided into multiple ***blocks*** of addresses or words
- A cache is decomposed into multiple ***sets***
- Each set consists of multiple ***cache lines***
- Each cache line consists of
 - A ***valid*** bit
 - A ***tag*** (set of bits)
 - A data ***block of B bytes***

Cache Lookup

- Memory address divided into tag, index, offset
- Index – to identify the set number
- Then, all the cache lines in the particular set are searched
- The tag in the cache line checked with the tag part of the memory address
- If valid bit is set and the tag matches, then *cache hit*, else *cache miss*
- Offset used to fetch a particular word from the data block in the cache line
- Depending on the number of sets and lines, caches can be of different kinds

Cache Lookup and Access



Direct-Mapped Caches


- Number of cache lines in a set = 1
- Refer figures 6.27-6.29 in book
- Disadvantages of direct-mapped caches?
- Conflict misses are more.
- Why?
- Problem is with a single cache line: A simple dot product example.

1. Alternative to Direct Mapping

- Set associative mapping
 - e.g., 2 way set associative: Number of cache lines = 2
 - Idea: A given memory block can be present in either of 2 lines of the cache

Set Associative Cache

- A memory block can be loaded into any cache block within a unique set of cache blocks




e.g., 4-way Set Associative Cache

- Assume 16KB cache, 32B block size
- $16\text{KB}/32\text{B} = 512$ blocks
- $512/4 = 128$ sets of blocks
- $\log_2 128 = 7$ set index bits
- $\log_2 32 = 5$ offset bits

Tag 20 bits	Set Index 7 bits	Offset 5 bits
----------------	------------------------	------------------

e.g., 2-way Set Associative Cache



Set Associative Caches


- Number of cache lines, m , in a set greater than 1
- Called *m-way associative* cache
- Refer figures 6.32-6.34 in the book

Fully Associative Caches

- Only one set having all cache lines
- Refer figures 6.35-6.37 in the book

Cache and Programming

- Objective: Learn how to assess cache related performance issues for important parts of our programs
- Will look at some examples of programs
- Will consider only data cache, assuming separate instruction and data caches
- Data cache configuration:
 - Direct mapped 16 KB with 32B block size



Example 1: Vector Sum Reduction

```
double A[2048], sum=0.0;
```

```
for (i=0; i<2048, i++) sum = sum +A[i];
```

- To do analysis, must view program close to machine code form (to see loads/stores)
 - Will assume that both loop index i and variable sum are implemented in registers
 - Will consider only accesses to array elements

Example 1: Reference Sequence

- load A[0] load A[1] load A[2] ... load A[2047]
- Assume base address of A (i.e., address of A[0]) is 0xA000

Example 1: Reference Sequence

- load A[0] load A[1] load A[2] ... load A[2047]
- Assume base address of A (i.e., address of A[0]) is 0xA000, 1010 0000 0000 0000
 - Cache index bits: 100000000 (value = 256)
- Size of an array element (double) = 8B
- So, 4 consecutive array elements fit into each cache block (block size is 32B)
 - A[0] – A[3] have index of 256
 - A[4] – A[7] have index of 257 and so on

Example 1: Cache Misses and Hits

A[0]	0xA000	256	Miss	Cold Start
A[1]	0xA008	256	Hit	
A[2]	0xA010	256	Hit	
A[3]	0xA018	256	Hit	
A[4]	0xA020	257	Miss	Cold Start

```
for (i=0; i<2048; i+=4) tmp=A[i];
```

```
for (i=0; i<2048, i++)
```

```
    sum = sum +A[i];
```

Cold start: we assume that the cache is initially empty

Hit ratio of our loop is 75% -- there are 1536 hits out of 2048 memory accesses

This is entirely due to spatial locality of reference.

What if we precede the loop by a loop that accesses all relevant memory blocks?

Hit ratio of our loop would then be 100%. 25% due to temporal locality and 75% due to spatial locality

Example 1 with double A[4096]

Why should it make a difference?

- Consider the case where the loop is preceded by another loop that accesses all array elements in order
- The entire array no longer fits into the cache – cache size: 16KB, array size: 32KB
- After execution of the previous loop, the second half of the array will be in cache
- Analysis: our loop will see misses as we had calculated

Example 1: Vector Sum Reduction

```
double A[2048], sum=0.0;  
for (i=0; i<2048, i++) sum = sum +A[i];
```

- To estimate data cache hit rate
 - we ignored accesses to sum, i
 - assumed address of A[0] is 0xA000
 - assumed only load/store instructions reference memory operands (others take their operands from registers)

Example 2: Vector Dot Product

```
double A[2048], B[2048], sum=0.0;
```

```
for (i=0; i<2048, i++) sum = sum +A[i] * B[i];
```

- Reference sequence:
 - load A[0] load B[0] load A[1] load B[1] ...
- Assume base addresses of A and B are 0xA000 and 0xE000
- Again, size of array elements is 8B so that 4 consecutive array elements fit into each cache block

Example 2: Vector Dot Product

Base addresses 0xA000 and 0xE000

.....10₁₀₀₀₀₀₀₀₀₀₀₀₀₀₀₀

Index: 256

.....11₁₀₀₀₀₀₀₀₀₀₀₀₀₀₀₀

Index: 256

Example 2: Cache Hits and Misses

A[0]	0xA000	256	Miss	Cold Start
B[0]	0xE000	256	Miss	Conflict
A[1]	0xA008	256	Miss	Conflict
B[1]	0xE008	256	Miss	Conflict
A[2]	0xA010	256	Miss	Conflict

Conflict: A miss due to conflict in cache block requirements caused by memory accesses of the same program

Hit ratio for our program:
0%

Source of the problem: the elements of arrays A and B are accessed in order and have the same cache index

Hit ratio would be better if the base address of array A was different from that of array B

Is this a contrived example?

```
double A[2048], B[2048], sum=0.0;
```

```
for (i=0; i<2048, i++) sum = sum +A[i] * B[i];
```

- How are variable addresses assigned?
- Start with some address, say 0xA000
- Assign addresses to variables in order of their declarations
- Array A: starting at 0xA000
- Array B: starting at 0xA000+ 2048*8
= 0xE000

1010	0000	0000	0000
100	0000	0000	0000
<hr/>			
1110	0000	0000	0000

Example 2: Cache Hits and Misses

A[0]	0xA000	256	Miss	Cold Start
B[0]	0xE000	256	Miss	Conflict
A[1]	0xA008	256	Miss	Conflict
B[1]	0xE008	256	Miss	Conflict
A[2]	0xA010	256	Miss	Conflict

Conflict: A miss due to conflict in cache block requirements caused by memory accesses of the same program

Hit ratio for our program:
0%

Source of the problem: the elements of arrays A and B are accessed in order and have the same cache index

Hit ratio would be better if the base address of array A was different from that of array B

Example 2 with Packing

- Assume that addresses are assigned as variables are encountered in declarations
- Our objective: to shift base address of B enough to make cache index of B[0] different from that of A[0]
double A[2052], B[2048];
- Base address of B is now 0xE020
 - 0xE020 is 1110 0000 0010 0000
 - Cache index of B[0] is 257; B[0] and A[0] do not conflict for the same cache block
- Hit ratio of our loop will rise to 75%

Example 2 with Array Merging

Alternatively, declare the arrays as

```
struct {double A, B;} array[2048];
for (i=0; i<2048, i++)
    sum += array[i].A*array[i].B;
```

Hit ratio: 75%

Example 3: DAXPY

- Double precision $Y = aX + Y$, where X and Y are vectors and a is a scalar

```
double X[2048], Y[2048], a;
```

```
for (i=0; i<2048;i++) Y[i] = a*X[i]+Y[i];
```

- Reference sequence
 - load $X[0]$ load $Y[0]$ store $Y[0]$ load $X[1]$ load $Y[1]$ store $Y[1]$...
- Hits and misses: Assuming that base addresses of X and Y don't conflict in cache, hit ratio of 83.3%

Example 4: 2-d Matrix Sum

```
double A[1024][1024], B[1024][1024];  
for (j=0;j<1024;j++)  
    for (i=0;i<1024;i++)  
        B[i][j] = A[i][j] + B[i][j];
```

- Reference Sequence:
 - load A[0,0]
 - load B[0,0]
 - store B[0,0]
 - load A[1,0]
 - load B[1,0]
 - store B[1,0]
 - ...
- Question: In what order are the elements of a multidimensional array stored in memory?

Storage of Multi-dimensional Arrays

■ Row major order

- Example: for a 2-dimensional array, the elements of the first row of the array are followed by those of the 2nd row of the array, then the 3rd row, and so on
- This is what is used in C

■ Column major order

- A 2-dimensional array is stored column by column in memory
- Used in FORTRAN

Example 4: 2-d Matrix Sum


```
double A[1024][1024], B[1024][1024];  
for (j=0;j<1024;j++)  
    for (i=0;i<1024;i++)  
        B[i][j] = A[i][j] + B[i][j];
```

- Reference Sequence:
 - load A[0,0]
 - load B[0,0]
 - store B[0,0]
 - load A[1,0]
 - load B[1,0]
 - store B[1,0]
 - ...
- Question: In what order are the elements of a multidimensional array stored in memory?


Example 4: Hits and Misses

- Reference order and storage order for our arrays are not the same
- Our loop will show no spatial locality
 - Assume that packing has been done to eliminate conflict misses due to base addresses
 - Miss(cold), Miss(cold), Hit for each array element
 - Hit ratio: 33.3%
 - Question: Will A[0,1] be in the cache when required later in the loop?

A



B



Example 4 with Loop Interchange

```
double A[1024][1024], B[1024][1024];  
for (i=0;i<1024;i++)  
    for (j=0;j<1024;j++)  
        B[i][j] = A[i][j] + B[i][j];
```

■ Reference Sequence:

load A[0,0] load B[0,0] store B[0,0]
load A[0,1] load B[0,1] store B[0,1]

Hit ratio: 83.3%

Is Loop Interchange Always Safe?

for ($j=1$; $j < 2048$; $j++$)

 for ($i=1$; $i < 2048$; $i++$)

$A[i][j] = A[i+1][j-1] + A[i][j-1];$

$$A[1,1] = A[2,0]+A[1,0]$$

$$A[2,1] = A[3,0]+A[2,0]$$

...

$$A[1,2] = A[2,1]+A[1,1]$$

Is Loop Interchange Always Safe?

for ($i=1$; $i < 2048$; $i++$) / interchanged

for ($j=1$; $j < 2048$; $j++$)

$A[i][j] = A[i+1][j-1] + A[i][j-1];$ **NO!**

$$A[1,1] = A[2,0]+A[1,0]$$

$$A[2,1] = A[3,0]+A[2,0]$$

...

$$A[1,2] = A[2,1]+A[1,1]$$

$$A[1,1] = A[2,0]+A[1,0]$$

$$A[1,2] = A[2,1]+A[1,1]$$

...

$$A[2,1] = A[3,0]+A[2,0]$$

Example 5: Matrix Multiplication

```
double X[N][N], Y[N][N], Z[N][N];  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        for (k=0; k<N; k++)  
            X[i][j] += Y[i][k] * Z[k][j];
```

Example 5: Matrix Multiplication

```
double X[N][N], Y[N][N], Z[N][N], tmp;
```

```
for (i=0; i<N; i++)
```

```
    for (j=0; j<N; j++) {
```

```
        tmp = 0;
```

```
        for (k=0; k<N; k++)
```

```
            tmp += Y[i][k] * Z[k][j];
```

X[i][j] = tmp; / Dot product inner loop

```
}
```

Y[0,0], Z[0,0], Y[0,1], Z[1,0], Y[0,2], Z[2,0] ... X[0,0],
Y[1,0], Z[0,1], Y[1,1], Z[1,1], Y[1,2], Z[2,1] ... X[0,1],
Y[2,0], Z[0,2], Y[2,1], Z[1,2], Y[2,2], Z[2,2] ... X[0,2],
...

Example 5: Matrix Multiplication

```
double X[N][N], Y[N][N], Z[N][N], tmp;
```

```
for (i=0; i<N; i++)
```

```
    for (j=0; j<N; j++){
```

```
        tmp = 0;
```

```
        for (k=0; k<N; k++)
```

```
            tmp += Y[i][k] * Z[k][j];
```

X[i][j] = tmp; / Dot product inner loop


```
}
```

Y[0,0], Z[0,0], Y[0,1], Z[1,0], Y[0,2], Z[2,0] ... X[0,0],

Y[1,0], Z[0,1], Y[1,1], Z[1,1], Y[1,2], Z[2,1] ... X[0,1],

Y[2,0], Z[0,2], Y[2,1], Z[1,2], Y[2,2], Z[2,2] ... X[0,2],

...



Matmul: Loop Interchange

- We can interchange the 3 loops
- Example: Interchange i and k loops - make the loops “kji” instead of “ijk”


double X[N][N], Y[N][N], Z[N][N];

for (k=0; k<N; k++)

 for (j=0; j<N; j++)

 for (i=0; i<N; i++)

 X[i][j] += Y[i][k] * Z[k][j];



- For the innermost loop: $Z[k][j]$ can be loaded into register once for each (k,j) , reducing the number of memory references

Analysis of loop interchange

- Will the result change?
- Will the total operations remain the same?
- Will the number of times X and Y are read remain the same?
- What about performance?
- Assumptions:
 - Elements are double elements – 8 bytes
 - Cache has 32-byte block size ($B=32$)

3 loops (i,j,k) - Can come up with 6 different versions ijk variant

```
for(i=0; i<N; i++)  
    for(j=0; j<N; j++)  
        sum=0.0;  
        for(k=0;k<N; k++)  
            sum += A[i][k]*B[k][j]  
        C[i][j] = sum;
```

- Loads per iteration – 2
- Stores per iteration – 0
- A misses per iteration – 0.25 (stride?)
- B misses per iteration – 1.00 (stride)
- C misses per iteration – 0.00
- Total misses per iteration – 1.25

jik variant

```
for(j=0; j<N; j++)
    for(i=0; i<N; i++)
        sum=0.0;
        for(k=0;k<N; k++)
            sum += A[i][k]*B[k][j]
        C[j][i] = sum;
```

- Loads per iteration – 2
- Stores per iteration – 0
- A misses per iteration – 0.25
- B misses per iteration – 1.00
- C misses per iteration – 0.00
- Total misses per iteration – 1.25
- Same as ijk variant

JKI variant

```
for(j=0; j<N; j++)
    for(k=0; k<N; k++)
        r=B[k][j];
        for(i=0;i<N; i++)
            C[i][j] += A[i][k]*r
```

- Loads per iteration – 2
- Stores per iteration – 1
- A misses per iteration – 1.00 (stride?)
- B misses per iteration – 0.00
- C misses per iteration – 1.00 (stride?)
- Total misses per iteration – 2.00

kji variant

```
for(k=0; k<N; k++)
    for(j=0; j<N; j++)
        r=B[k][j];
        for(i=0;i<N; i++)
            C[i][j] += A[i][k]*r
```

- Loads per iteration – 2
- Stores per iteration – 1
- A misses per iteration – 1.00
- B misses per iteration – 0.00
- C misses per iteration – 1.00
- Total misses per iteration – 2.00
- Same as jki variant

kij variant

```
for(k=0; k<N; k++)
    for(i=0; i<N; i++)
        r=A[i][k];
        for(j=0;j<N; j++)
            C[i][j] += r*B[k][j]
```

- Loads per iteration – 2
- Stores per iteration – 1
- A misses per iteration – 0.00
- B misses per iteration – 0.25 (stride?)
- C misses per iteration – 0.25 (stride?)
- Total misses per iteration – 0.50

ikj variant

```
for(i=0; i<N; i++)
    for(k=0; k<N; k++)
        r=A[i][k];
        for(j=0;j<N; j++)
            C[i][j] += r*B[k][j]
```

- Loads per iteration – 2
- Stores per iteration – 1
- A misses per iteration – 0.00
- B misses per iteration – 0.25 (stride?)
- C misses per iteration – 0.25 (stride?)
- Total misses per iteration – 0.50
- Same as kij variant

Summary

Finding the best performance involves trade-off between

Cache performance

Number of memory accesses

How to improve the cache hits and performance further?

Loop unrolling and blocking

“Loop Unrolling”

```
double X[10];  
for (i=0; i<10; i++)  
    X[i] = X[i] - 1;
```

Unrolled once:

```
for (i=0; i<10; i+=2){  
    X[i] = X[i] - 1;  
    X[i+1] = X[i+1] - 1;  
}
```

Fully unrolled:

```
X[0] = X[0] - 1;  
X[1] = X[1] - 1;  
X[2] = X[2] - 1;  
...  
X[9] = X[9] - 1;
```

Unrolling Matrix Multiplication

```
double X[N][N], Y[N][N], Z[N][N];  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        for (k=0; k<N; k++)  
            X[i][j] += Y[i][k] * Z[k][j];
```

Let us unroll the k loop once

Matmul: k loop unrolled

```
double X[N][N], Y[N][N], Z[N][N];  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        for (k=0; k<N; k+=2) /* k loop unrolled once  
            X[i][j] += Y[i][k] * Z[k][j] + Y[i][k+1] * Z[k+1][j];
```

Now, let us also unroll the j loop once

Matmul: k and j loops unrolled

```
double X[N][N], Y[N][N], Z[N][N];
```

```
for (i=0; i<N; i++)
```


```
    for (j=0; j<N; j+=2)
```

```
        for (k=0; k<N; k+=2){ /* j and k loops unrolled once
```

```
            X[i][j] += Y[i][k] * Z[k][j] + Y[i][k+1] * Z[k+1][j];
```

```
            X[i][j+1] += Y[i][k] * Z[k][j+1] + Y[i][k+1] * Z[k+1][j+1];
```



```
}
```




Exploits spatial locality for arrays Y, Z


Exploits temporal locality for array Y

Provides a programming idea for enhancing locality



$$Y \times Z = X$$



 “Block” of array elements
Block size, B

 Row sliver

Matmul: “Blocking” or “Tiling”

```
double X[N][N], Y[N][N], Z[N][N];  
for (jj=0; jj<N; jj+=B)  
    for (kk=0; kk<N; kk+=B)  
        for (i=0; i<N, i++) {  
            for (j=jj; j < min(jj+B, N); j++) {  
                sum = 0.0;  
                for (k=kk; k<min(kk+B, N), k++)  
                    sum += Y[i][k] * Z[k][j];  
                X[i][j] += sum;  
            } /* for j */  
        } /* for i */
```

Vector Operations

Example: Vector Sum

```
double A[2048], B[2048], C[2048];
```

```
for (i=0; i<2048, i++) C[i] = A[i] + B[i];
```

- What if a CPU has 4 adders?
- It can be designed to support an instruction to do 4 iterations of the Vector Sum loop at a time
 - VADD v1_A[0:3], v2_B[0:3], v3_C[0:3]
- Called a vector instruction

Multimedia Extensions

- Hardware support for operations on “short vectors” is provided in existing microprocessors
- Example: 256 bit registers, each split into 4x64b (or 8x32b)
 - Maximum vector length
- Example: Intel “x86” processors
 - SSE (Streaming SIMD Extension)
 - AVX (Advanced Vector Extension)

Vectorization of Loops

We will use a generic notation

Instead of

VADD C[0:3], A[0:3], B[0:3]

$C[0:3] = A[0:3] + B[0:3]$

An example of vectorization

- Given maximum vector length, VL

```
for (i=0; i < N; i++)
```

```
    A[i] = A[i] + B[i];
```

```
for (i=0; i < N; i+=VL)
```

```
    A[i:i+VL-1] = A[i:i+VL-1] + B[i:i+VL-1];
```

What if N is not divisible by VL?

```
for (i=0; i < (N - N%VL); i+=VL)
```

```
    A[i:i+VL-1] = A[i:i+VL-1] + B[i:i+VL-1];
```

```
for (; i<N; i++) A[i] = A[i] + B[i];
```

- This technique is called Stripmining

Possible complications


- Dependences between statements within the loop

Example 1

```
for (i=0; i < N; i++) {  
    A[i] = B[i] + C[i];  
    D[i] = (A[i] + A[i+1])/2;  
}  
  
for (i=0; i < (N - N%VL); i+=VL){  
    A[i:i+VL-1] = B[i:i+VL-1] + C[i:i+VL-1];  
    D[i: ... will get wrong value of A[i+1], etc
```

Example 1

```
for (i=0; i < N; i++) {  
    A[i] = B[i] + C[i];  
    D[i] = (A[i] + A[i+1])/2;  
}  
  
for (i=0; i < N; i++) {  
    temp[i] = A[i+1];  
    A[i] = B[i] + C[i];  
    D[i] = (A[i] + temp[i])/2;  
}
```




- This loop transformation, through copying of data, is called Node Splitting

Example 2

```
for (i=0; i < N; i++) {  
    X = A[i] + 1;  
    B[i] = X + C[i];  
}
```


```
for (i=0; i < N; i++) {  
    temp[i] = A[i] + 1;  
    B[i] = temp[i] + C[i];  
}
```

- Scalar expansion



Example 3

```
for (i=0; i < N; i++) {  
    A[i] = B[i];  
    C[i] = C[i-1] + 1;  
}
```



```
for (i=0; i < N; i++) A[i] = B[i];  
for (i=0; i<N; i++) C[i] = C[i-1] + 1;
```




- Loop fission

Example 4

```
for (j=1; i < N; j++)  
    for (i=2; i < N; i++)  
        A[i,j] = A[i-1, j] + B[i];
```

```
for (i=2; i < N; i++)  
    for (j=1; j<N; j++)  
        A[i,j] = A[i-1,j] + B[i];
```

- Loop interchange



Data Representation

Integer Data

- Signed vs Unsigned integer
- Representing a signed integer
 - 2s complement representation

The n bit quantity

$$x_{n-1}x_{n-2}\dots x_2x_1x_0$$

least significant bit

represents the signed integer value

$$-x_{n-1}2^{n-1} + \sum_{i=0}^{n-2} x_i 2^i$$

Example: 2s complement

- The signed integer -14_{10} (decimal) is represented as
 - 10010 in 5 bits (i.e., $-16 + 2$)
 - 110010 in 6 bits (i.e., $-32 + 16 + 2$)
 - 111...1110010 in 32 bits

Aside: Hexadecimal (base 16)

- Digits 0 1 2 3 4 5 6 7 8 9 A B C D E F
0000 0001 0010 ... 1101 1110 1111
- Binary sequences can be written more compactly in hexadecimal

1001 9

1010 A

1011 B

1100 C

1101 D

1110 E

1111 F

Example: 2s complement

- The signed integer -14_{10} (decimal) is represented as

10010 in 5 bits ($-16 + 2$)

110010 in 6 bits ($-32 + 16 + 2$)

$111\dots1110010$ in 32 bits

$1111\ 1111\ 1111\dots0010$

FFFFFF2

Usually written as $0x\text{FFFFFF2}$

Real Data

- How to represent real data?
- Fixed point representation
 - (sign bit) (Integer part – 23 bits)(Fraction – 8 bits)
- Disadvantage?
- The range of numbers not adequate for many practical problems
- Hence, floating point representation

Normalized floating point numbers

- Consists of two parts
- Mantissa with sign
- Exponent with sign
- Floating point represented as
- $(\text{sign}) \times \text{mantissa} \times 2^{+\text{-exponent}}$
- Mantissa (23 bits) – a binary fraction with non-zero leading bit
- Exponent (8 bits) – 1 bit used for sign of the exponent, other 7 bits used for the exponent magnitude
- Range of exponent?: -127 to +127
- Disadvantage: Two representation for 0 exponent: -0 and +0

Excess representation or bias format

- Exponent has no sign bit
- 8 bits of exponent divided as
 - (0-127) and (128-255)
 - (0-126): negative
 - 127: represents 0
 - (128-255): positive
 - Called as bias 127 for exponent
- Given an exponent, exp, value of exponent will be $(exp - 127)$
- Largest and smallest floating point numbers that can be represented?

Example

Representing 52.21875 in 32-bit floating point format.

Step 1 – Represent using binary: 110100.00111

Step 2 – Normalized representation: 1.1010000111×2^5

Exponent of 5 represented as $(127+5=132) = 10000100$

Floating point represented as (sign)(exponent-8 bits)(mantissa – 23 bits)

0 10000100 101000011000000000000000

Note that the leading 1 in the normalized representation is ignored in the mantissa.

IEEE 754 Floating Point Standard

The scheme that is just described is IEEE 754 floating point standard

Mantissa is called *significand* as per the standard

The standard uses a normalized significand – most significant bit is always 1

Thus significand is 24 bits long – 1 is implied + 23 explicit

Floating point number represented by:
 $(-1)s \times (1.f)^2 \times 2^{(\text{exponent}-127)}$

Special Cases (B&O 2.4.2)

- Representation of 0: All 31 (exponent and mantissa/significand/fraction) bits are 0's
 - +0: 0 for the sign bit
 - -0: 1 for the sign bit
 - All 0's for the exponents is not allowed to be used for any other number
- Infinity: All 1's in the exponent and all 0's in the mantissa
 - +infinity: 0 for the sign bit
 - -infinity: 1 for the sign bit

Largest and Smallest Positive Numbers?

Rounding (B&O 2.4.4)

- When mathematical operations are performed with two floating point numbers, the significand of the result may exceed 23 bits after the adjustment of the exponent.
- Rounding:
 - Rounding upwards:
 - e.g., significand: 0.110...01,
 - overflow: 1,
 - significand after rounding: 0.110....10,
i.e., add 1 to LSB.
 - Rounding downwards: extra bits ignored

Summary (from notes by Prof. Rajaraman)

Value	Sign	Exponent (8 bits)	Significand (23 bits)	
+0	0	00000000	00 ... 00	(all 23 bits 0)
- 0	1	00000000	00 ... 00	(all 23 bits 0)
$+1.f \times 2^{(e-b)}$ e exponent, b bias	0	00000001 to 11111110	$a a \dots a a$	($a = 0$ or 1)
$-1.f \times 2^{(e-b)}$	1	00000001 to 11111110	$a a \dots a a$	($a = 0$ or 1)
$+\infty$	0	11111111	000 ... 00	(all 23 bits 0)
$-\infty$	1	11111111	000 ... 00	(all 23 bits 0)
SNaN	0 or 1	11111111	000 ... 01 to 011... 11	leading bit 0 (at least one 1 in the rest)
QNaN	0 or 1	11111111	1000 ... 10	leading bit 1
Positive subnormal $0.f \times 2^{x+1-b}$ (x is the number of leading 0s in significand)	0	00000000	000 ... 01 to 111... 11	(at least one 1)

Summary (from notes by Prof. Rajaraman)

Operation	Result	Operation	Result
$n / \pm\infty$	0	$\pm 0 / \pm 0$	NaN
$\pm\infty / \pm\infty$	$\pm\infty$	$\infty - \infty$	NaN
$\pm n / 0$	$\pm\infty$	$\pm\infty / \pm\infty$	NaN
$\infty + \infty$	∞	$\pm\infty \times 0$	NaN

IEEE-754 Standard for 64-bit floating point numbers

- s: 1 bit
- e: 11 bits
- f: 52 bits

Reading

- Read sections in Bryant and O'Hallaron on the topics we have discussed in class
 - Bryant, O'Hallaron. Computer Systems – A Programmer's Perspective, Pearson Education Limited 2016, 3rd Global Edition
- Try to solve some of the problems

IEEE Standard for Floating Point Numbers

V Rajaraman

Floating point numbers are an important data type in computation which is used extensively. Yet, many users do not know the standard which is used in almost all computer hardware to store and process these. In this article, we explain the standards evolved by The Institute of Electrical and Electronic Engineers in 1985 and augmented in 2008 to represent floating point numbers and process them. This standard is now used by all computer manufacturers while designing floating point arithmetic units so that programs are portable among computers.

Introduction

There are two types of arithmetic which are performed in computers: integer arithmetic and real arithmetic. Integer arithmetic is simple. A decimal number is converted to its binary equivalent and arithmetic is performed using binary arithmetic operations. The largest positive integer that may be stored in an 8-bit byte is +127, if 1 bit is used for sign. If 16 bits are used, the largest positive integer is +32767 and with 32 bits, it is +2147483647, quite large! Integers are used mostly for counting. Most scientific computations are however performed using real numbers, that is, numbers with a fractional part. In order to represent real numbers in computers, we have to ask two questions. The first is to decide how many bits are needed to encode real numbers and the second is to decide how to represent real numbers using these bits. Normally, in numerical computation in science and engineering, one would need at least 7 to 8 significant digits precision. Thus, the number of bits needed to encode 8 decimal digits is approximately 26, as $\log_2 10 = 3.32$ bits are needed on the average to encode a digit. In computers, numbers are stored as a sequence of 8-bit bytes. Thus 32 bits (4 bytes) which is bigger than 26 bits is a logical size to use for real numbers. Given 32 bits to encode real




V Rajaraman is at the Indian Institute of Science, Bengaluru. Several generations of scientists and engineers in India have learnt computer science using his lucidly written textbooks on programming and computer fundamentals. His current research interests are parallel computing and history of computing.

Keywords

Floating point numbers, rounding, decimal floating point numbers, IEEE 754-2008 Standard.



Figure 1. Fixed point representation of real numbers in binary using 32 bits.



numbers, the next question is how to break it up into an integer part and a fractional part. One method is to divide the 32 bits into 2 parts, one part to represent an integer part of the number and the other the fractional part as shown in *Figure 1*.

In this figure, we have arbitrarily fixed the (virtual) binary point between bits 7 and 8. With this representation, called *fixed point representation*, the largest and the smallest positive binary numbers that may be represented using 32 bits are

$$\text{Largest: } +111 \dots 1.1111111 = 1677215.998046875$$

 23 bits 8 bits

$$\text{Smallest: } +000 \dots 0.00000001 = 0.00390625$$

 23 bits 8 bits

Binary Floating Point Numbers

This range of real numbers, when fixed point representation is used, is not sufficient in many practical problems. Therefore, another representation called *normalized floating point representation* is used for real numbers in computers. In this representation, 32 bits are divided into two parts: a part called the *mantissa* with its sign and the other called the *exponent* with its sign. The mantissa represents fractions with a non-zero leading bit and the exponent the power of 2 by which the mantissa is multiplied. This method increases the range of numbers that may be represented using 32 bits. In this method, a binary floating point number is represented by

$$(\text{sign}) \times \text{mantissa} \times 2^{\pm \text{exponent}}$$


As the range of real numbers representable with fixed point is not sufficient, normalized floating point is used to represent real numbers.

where the sign is one bit, the mantissa is a binary fraction with a non-zero leading bit, and the exponent is a binary integer. If 32

bits are available to store floating point numbers, we have to decide the following:

1. How many bits are to be used to represent the mantissa (1 bit is reserved for the sign of the number).
2. How many bits are to be used for the exponent.
3. How to represent the sign of the exponent.

The number of bits to be used for the mantissa is determined by the number of significant decimal digits required in computation, which is at least seven. Based on experience in numerical computation, it is found that at least seven significant decimal digits are needed if results are expected without too much error. The number of bits required to represent seven decimal digits is approximately 23. The remaining 8 bits are allocated for the exponent. The exponent may be represented in sign magnitude form. In this case, 1 bit is used for sign and 7 bits for the magnitude. The exponent will range from -127 to $+127$. The only disadvantage to this method is that there are two representations for 0 exponent: $+0$ and -0 . To avoid this, one may use an *excess representation* or *biased format* to represent the exponent. The exponent has no sign in this format. The range 0 to 255 of an 8-bit number is divided into two parts 0 to 127 and 128 to 255 as shown below:



All bit strings from 0 to 126 are considered negative, exponent 127 represents 0, and values greater than 127 are positive. Thus the range of exponents is -127 to $+128$. Given an exponent string exp , the value of the exponent will be taken as $(exp - 127)$. The main advantage of this format is a unique representation for exponent zero. With the representation of binary floating point numbers explained above, the largest floating point number which can be represented is

The number of bits to be used for the mantissa is determined by the number of significant decimal digits required in computation, which is at least seven.

The main advantage of biased exponent format is unique representation for exponent 0.



$$\begin{aligned}
 & 0.1111\dots1111 \times 2^{11111111} \\
 & |\leftarrow 23 \text{ bits} \rightarrow| \\
 & = (1 - 2^{-23}) \times 2^{255-127} \\
 & \cong 3.4 \times 10^{38}.
 \end{aligned}$$

The smallest floating point number is

$$\begin{aligned}
 & 0.10000\dots00 \times 2^{-127} \\
 & |\leftarrow 23 \text{ bits} \rightarrow| \\
 & \cong 0.293 \times 10^{-38}.
 \end{aligned}$$

Example. Represent 52.21875 in 32-bit binary floating point format.

$$52.21875 = 110100.00111 = .11010000111 \times 2^6.$$

Normalized 23 bit mantissa = 0.11010000111000000000000.

As excess representation is being used for exponent, it is equal to $127 + 6 = 133$.

Thus the representation is

$$52.21875 = 0.11010000111 \times 2^{133} = 0.110100000111 \times 2^{10000101}.$$

The 32-bit string used to store 52.21875 in a computer will thus be 0100001011101000011100000000000.

Here, the most significant bit is the sign, the next 8 bits the exponent, and the last 23 bits the mantissa.

Institute of Electrical and Electronics

Engineers standardized how floating point binary numbers would be represented in computers, how these would be rounded, how 0 and ∞ would be represented and how to treat exception condition such as attempt to divide by 0.

IEEE Floating Point Standard 754-1985

Floating point binary numbers were beginning to be used in the mid 50s. There was no uniformity in the formats used to represent floating point numbers and programs were not portable from one manufacturer's computer to another. By the mid 1980s, with the advent of personal computers, the number of bits used to store floating point numbers was standardized as 32 bits. A Standards Committee was formed by the Institute of Electrical and Electronics Engineers to standardize how floating point binary numbers would be represented in computers. In addition, the standard specified uniformity in rounding numbers, treating exception conditions such as attempt to divide by 0, and representation of 0 and infinity (∞). This standard, called IEEE Standard 754 for



floating point numbers, was adopted in 1985 by all computer manufacturers. It allowed porting of programs from one computer to another without the answers being different. This standard defined floating point formats for 32-bit and 64-bit numbers. With improvement in computer technology it became feasible to use a larger number of bits for floating point numbers. After the standard was used for a number of years, many improvements were suggested. The standard was updated in 2008. The current standard is IEEE 754-2008 version. This version retained all the features of the 1985 standard and introduced new standards for 16 and 128-bit numbers. It also introduced standards for representing decimal floating point numbers. We will describe these standards later in this article. In this section, we will describe IEEE 754-1985 Standard.

IEEE floating point representation for binary real numbers consists of three parts. For a 32-bit (called single precision) number, they are:

1. Sign, for which 1 bit is allocated.
2. Mantissa (called *significand* in the standard) is allocated 23 bits.
3. Exponent is allocated 8 bits. As both positive and negative numbers are required for the exponent, instead of using a separate sign bit for the exponent, the standard uses a biased representation. The value of the bias is 127. Thus an exponent 0 means that -127 is stored in the exponent field. A stored value 198 means that the exponent value is $(198 - 127) = 71$. The exponents -127 (all 0s) and +128 (all 1s) are reserved for representing special numbers which we discuss later.

To increase the precision of the significand, the IEEE 754 Standard uses a normalized significand which implies that its most significant bit is always 1. As this is implied, it is assumed to be on the left of the (virtual) decimal point of the significand. Thus in the IEEE Standard, the significand is 24 bits long – 23 bits of the significand which is stored in the memory and an implied 1 as the most significant 24th bit. The extra bit increases the number

The current standard is IEEE 754-2008 version. This version retained all the features of the 1985 standard and introduced new standards for 16-bit and 128-bit numbers. It also introduced standards for representing decimal floating point numbers.



Figure 2. IEEE 754 representation of 32-bit floating point number.

b ₀	b ₁ b ₂ b ₃ b ₈	b ₉ b ₁₀ b ₁₁b ₃₀ b ₃₁
Sign	Exponent	Significand

of significant digits in the significand. Thus, a floating point number in the IEEE Standard is

$$(-1)^s \times (1.f)_2 \times 2^{\text{exponent} - 127}$$

where s is the sign bit; $s = 0$ is used for positive numbers and $s = 1$ for representing negative numbers. f represents the bits in the significand. Observe that the implied 1 as the most significant bit of the significand is explicitly shown for clarity.

For a 32-bit word machine, the allocation of bits for a floating point number are given in *Figure 2*. Observe that the exponent is placed before the significand (see *Box 1*).

Example. Represent 52.21875 in IEEE 754 – 32-bit floating point format.

$$\begin{aligned} 52.21875 &= 110100.00111 \\ &= 1.1010000111 \times 2^5. \end{aligned}$$

Normalized significand = .1010000111.

Exponent: $(e - 127) = 5$ or, $e = 132$.

The bit representation in IEEE format is

0	10000100	1010000111000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

Box 1.

Why does IEEE 754 standard use biased exponent representation and place it before the significand?

Computers normally have separate arithmetic units to compute with integer operands and floating point operands. They are called FPU (Floating Point Unit) and IU (Integer Unit). An IU works faster and is cheaper to build. In most recent computers, there are several IUs and a smaller number of FPUs. It is preferable to use IUs whenever possible. By placing the sign bit and biased exponent bits as shown in *Figure 2*, operations such as $x < r.o. > y$, where *r.o.* is a relational operator, namely, $>$, $<$, \geq , \leq , and x, y are reals, can be carried out by integer arithmetic units as the exponents are integers. Sorting real numbers may also be carried out by IUs as biased exponent are the most significant bits. This ensures lexicographical ordering of real numbers; the significand need not be considered as a fraction.

Special Values in IEEE 754-1985 Standard (32 Bits)

Representation of Zero: As the significand is assumed to have a hidden 1 as the most significant bit, all 0s in the significand part of the number will be taken as 1.00...0. Thus zero is represented in the IEEE Standard by all 0s for the exponent and all 0s for the significand. All 0s for the exponent is not allowed to be used for any other number. If the sign bit is 0 and all the other bits 0, the number is +0. If the sign bit is 1 and all the other bits 0, it is -0. Even though +0 and -0 have distinct representations they are assumed equal. Thus the representations of +0 and -0 are:

+0

0	00000000	00000000000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

-0

1	00000000	00000000000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

Representation of Infinity: All 1s in the exponent field is assumed to represent infinity (∞). A sign bit 0 represents $+\infty$ and a sign bit 1 represents $-\infty$. Thus the representations of $+\infty$ and $-\infty$ are:

$+\infty$

0	11111111	00000000000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

$-\infty$

1	11111111	00000000000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

Representation of Non Numbers: When an operation is performed by a computer on a pair of operands, the result may not be mathematically defined. For example, if zero is divided by zero,

As the significand is assumed to have a hidden 1 as the most significant bit, all 0s in the significand part of the number will be taken as 1.00...0. Thus zero is represented in the IEEE Standard by all 0s for the exponent and all 0s for the significand.

All 1s in the exponent field is assumed to represent infinity (∞).



When an arithmetic operation is performed on two numbers which results in an indeterminate answer, it is called NaN (Not a Number) in IEEE Standard.

the result is indeterminate. Such a result is called Not a Number (NaN) in the IEEE Standard. In fact the IEEE Standard defines two types of NaN. When the result of an operation is not defined (i.e., indeterminate) it is called a Quiet NaN (QNaN). Examples are: $0/0$, $(\infty - \infty)$, $\sqrt{-1}$. Quiet NaNs are normally carried over in the computation. The other type of NaN is called a Signalling Nan (SNaN). This is used to give an error message. When an operation leads to a floating point underflow, i.e., the result of a computation is smaller than the smallest number that can be stored as a floating point number, or the result is an overflow, i.e., it is larger than the largest number that can be stored, SNaN is used. When no valid value is stored in a variable name (i.e., it is undefined) and an attempt is made to use it in an arithmetic operation, SNaN would result. QNaN is represented by 0 or 1 as the sign bit, all 1s as exponent, and a 0 as the left-most bit of the significand and at least one 1 in the rest of the significand. SNaN is represented by 0 or 1 as the sign bit, all 1s as exponent, and a 1 as the left-most bit of the significand and any string of bits for the remaining 22 bits. We give below the representations of QNaN and SNaN.

QNaN

0 or 1	11111111	0001000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

The most significant bit of the significand is 0. There is at least one 1 in the rest of the significand.

SNaN

0 or 1	11111111	10000000000001000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

The most significant bit of the significand is 1. Any combination of bits is allowed for the other bits of the significand.

Largest and Smallest Positive Floating Point Numbers:

Largest Positive Number

0	11111110	1111111111111111111111
Sign 1 bit	Exponent 8 bits	Significand 23 bits

Significand: $1111 \dots 1 = 1 + (1 - 2^{-23}) = 2 - 2^{-23}$.

Exponent: $(254 - 127) = 127$.

Largest Number $= (2 - 2^{-23}) \times 2^{127} \cong 3.403 \times 10^{38}$.

If the result of a computation exceeds the largest number that can be stored in the computer, then it is called an *overflow*.

Smallest Positive Number

0	00000001	00000000000000000000000000000000
Sign 1 bit	Exponent 8 bits	Significand 23 bits

Significand = 1.0.

Exponent = $1 - 127 = -126$.

The smallest normalized number is $= 2^{-126} \cong 1.1755 \times 10^{-38}$.

Subnormal Numbers: When all the exponent bits are 0 and the leading hidden bit of the significand is 0, then the floating point number is called a *subnormal number*. Thus, one logical representation of a subnormal number is

$$(-1)^s \times 0.f \times 2^{-127} \text{ (all 0s for the exponent)},$$

where f has at least one 1 (otherwise the number will be taken as 0). However, the standard uses -126, i.e., bias +1 for the exponent rather than -127 which is the bias for some not so obvious reason, possibly because by using -126 instead of -127, the gap between the largest subnormal number and the smallest normalized number is smaller.

The largest subnormal number is $0.999999988 \times 2^{-126}$. It is close to the smallest normalized number 2^{-126} .

When all the exponent bits are 0 and the leading hidden bit of the significand is 0, then the floating point number is called a *subnormal number*.



By using subnormal numbers, underflow which may occur in some calculations are gradual.

The smallest positive subnormal number is

0	00000000	000000000000000000000001
Sign 1 bit	Exponent 8 bits	Significand 23 bits

the value of which is $2^{-23} \times 2^{-126} = 2^{-149}$.

A result that is smaller than the smallest number that can be stored in a computer is called an *underflow*.

You may wonder why subnormal numbers are allowed in the IEEE Standard. By using subnormal numbers, underflow which may occur in some calculations are gradual. Also, the smallest number that could be represented in a machine is closer to zero. (See Box 2.)

Box 2. Machine Epsilon and Dwarf

In any of the formats for representing floating point numbers, *machine epsilon* is defined as the difference between 1 and the next larger number that can be stored in that format. For example, in 32-bit IEEE format with a 23-bit significand, the machine epsilon is $2^{-23} = 1.19 \times 10^{-7}$. This essentially tells us that the precision of decimal numbers stored in this format is 7 digits. The term precision and accuracy are not the same. Accuracy implies correctness whereas precision does not. For 64-bit representation of IEEE floating point numbers the significand length is 52 bits. Thus, machine epsilon is $2^{-52} = 2.22 \times 10^{-16}$. Therefore, decimal calculations with 64 bits give 16 digit precision. This is a conservative definition used by industry. In MATLAB, machine epsilon is as defined above. However, academics define machine epsilon as the upper bound of the relative error when numbers are rounded. Thus, for 32-bit floating point numbers, the machine epsilon is $2^{-23}/2 \cong 5.96 \times 10^{-8}$.

The machine epsilon is useful in iterative computation. When two successive iterates differ by less than $|\text{epsilon}|$ one may assume that the iteration has converged. The IEEE Standard does not define machine epsilon.

Tiny or dwarf is the minimum subnormal number that can be represented using the specified floating point format. Thus, for IEEE 32-bit format, it is

0	00000000	000000000000000000000001
Sign 1 bit	Exponent 8 bits	Significand 23 bits

whose value is: $2^{-126-23} = 2^{-149} \cong 1.4 \times 10^{-45}$. Any number less than this will signal underflow. The advantage of using subnormal numbers is that underflow is gradual as was stated earlier.



When a mathematical operation such as add, subtract, multiply, or divide is performed with two floating point numbers, the significand of the result may exceed 23 bits after the adjustment of the exponent. In such a case, there are two alternatives. One is to truncate the result, i.e., ignore all bits beyond the 32nd bit. The other is to round the result to the nearest significand. For example, if a significand is 0.110 .. 01 and the overflow bit is 1, a rounded value would be 0.111 ... 10. In other words, if the first bit which is truncated is 1, add 1 to the least significant bit, else ignore the bits. This is called *rounding upwards*. If the extra bits are ignored, it is called *rounding downwards*. The IEEE Standard suggests to hardware designers to get the best possible result while performing arithmetic that are *reproducible* across different manufacturer's computers using the standard. The Standard suggests that in the equation

$$c = a \text{ } <\text{op}> b,$$

where a and b are operands and $<\text{op}>$ an arithmetic operation, the result c should be as if it was computed exactly and then rounded. This is called *correct rounding*.

In Tables 1 and 2, we summarise the discussions in this section.

Table 1. IEEE 754-85 floating point standard. We use f to represent the significand, e to represent the exponent, b the bias of the exponent, and \pm for the sign.

Value	Sign	Exponent (8 bits)	Significand (23 bits)	
+0	0	00000000	00 ... 00	(all 23 bits 0)
- 0	1	00000000	00 ... 00	(all 23 bits 0)
$+ 1.f \times 2^{(e - b)}$ e exponent, b bias	0	00000001 to 11111110	$a a \dots a a$	($a = 0$ or 1)
$- 1.f \times 2^{(e - b)}$	1	00000001 to 11111110	$a a \dots a a$	($a = 0$ or 1)
$+ \infty$	0	11111111	000 ... 00	(all 23 bits 0)
$- \infty$	1	11111111	000 ... 00	(all 23 bits 0)
SNaN	0 or 1	11111111	000 ... 01 to 011... 11	leading bit 0 (at least one 1 in the rest)
QNaN	0 or 1	11111111	1000 ... 10	leading bit 1
Positive subnormal $0.f \times 2^{x+1-b}$ (x is the number of leading 0s in significand)	0	00000000	000 ... 01 to 111... 11	(at least one 1)



Table 2. Operations on special numbers. All NaNs are Quiet NaNs.

Operation	Result	Operation	Result
$n / \pm \infty$	0	$\pm 0 / \pm 0$	NaN
$\pm \infty / \pm \infty$	$\pm \infty$	$\infty - \infty$	NaN
$\pm n / 0$	$\pm \infty$	$\pm \infty / \pm \infty$	NaN
$\infty + \infty$	∞	$\pm \infty \times 0$	NaN

With 64 bits available to store floating point numbers, there is more freedom to increase the significant digits in the significand and increase the range by allocating more bits to the exponent.

The standard allocates 1 bit for sign, 11 bits for the exponent, and 52 bits for the significand.

IEEE 754 Floating Point 64-Bit Standard – 1985

The IEEE 754 floating point standard for 64-bit (called double precision) numbers is very similar to the 32-bit standard. The main difference is the allocation bits for the exponent and the significand. With 64 bits available to store floating point numbers, there is more freedom to increase the significant digits in the significand and increase the range by allocating more bits to the exponent. The standard allocates 1 bit for sign, 11 bits for the exponent, and 52 bits for the significand. The exponent uses a biased representation with a bias of 1023. The representation is shown below:

Sign	Exponent	Significand
1 bit	11 bits	52 bits

A number in this standard is thus

$$(-1)^s \times (1.f) \times 2^{(\text{exponent} - 1023)},$$

where $s = \pm 1$, and f is the value of the significand.

The largest positive number which can be represented in this standard is

0	11111111110	1111.....1111
Sign 1 bit	Exponent 11 bits	Significand 52 bits

which is $= (2 - 2^{-52}) \times 2^{(2046 - 1023)} = (2 - 2^{-52}) \times 2^{1023} \cong 10^{3083}$. The smallest positive number that can be represented using 64



bits using this standard is

0	000000000001	0000.....00000
Sign 1 bit	Exponent 11 bits	Significand 52 bits

whose value is $2^{1-1023} = 2^{-1022}$.

The definitions of ± 0 , $\pm \infty$, QNaN, SNaN remain the same except that the number of bits in the exponent and significand are now 11 and 52 respectively. Subnormal numbers have a similar definition as in 32-bit standard.

IEEE 754 Floating Point Standard – 2008

This standard was introduced to enhance the scope of IEEE 754 floating point standard of 1985. The enhancements were introduced to meet the demands of three sets of professionals:

1. Those working in the graphics area.
2. Those who use high performance computers for numeric intensive computations.
3. Professionals carrying out financial transactions who require exact decimal computation.

This standard has not changed the format of 32- and 64-bit numbers. It has introduced floating point numbers which are 16 and 128 bits long. These are respectively called half and quadruple precision. Besides these, it has also introduced standards for floating point decimal numbers.

The 16-Bit Standard

The 16-bit format for real numbers was introduced for pixel storage in graphics and not for computation. (Some graphics processing units use the 16-bit format for computation.) The 16-bit format is shown in *Figure 3*.

b ₀	b ₁ b ₂ b ₃ b ₄ b ₅	b ₆ b ₇ b ₈ b ₉ b ₁₀ ... b ₁₄ b ₁₅
Sign 1 bit	Exponent 5 bits	Significand 10 bits

IEEE 2008 Standard has introduced standards for 16-bit and 128-bit floating point numbers. It has also introduced standards for floating point decimal numbers.

The 16-bit format for real numbers was introduced for pixel storage in graphics and not for computation.

Figure 3. Representation of 16-bit floating point numbers.



With 16-bit representation, 5 bits are used for the exponent and 10 bits for the significand. Biased exponent is used with a bias of 15. Thus the exponent range is: -14 to $+15$. Remember that all 0s and all 1s for the exponent are reserved to represent 0 and ∞ , respectively.

The maximum positive integer that can be stored is

$$\begin{aligned} +1.111..1 \times 2^{15} &= 1 + (1 - 2^{-11}) \times 2^{15} = (2 - 2^{-11}) \times 2^{15} \\ &\approx 65504. \end{aligned}$$

The minimum positive number is

$$+1.000\dots0 \times 2^{-14} = 2^{-14} = 0.61 \times 10^{-4}.$$

The minimum subnormal 16-bit floating point number is

$$2^{-24} \approx 5.96 \times 10^{-8}.$$

With improvements in computer technology, using 128 bits to represent floating point numbers has become feasible. This is sometimes used in numeric-intensive computation, where large rounding errors may occur.

The definitions of ± 0 , $\pm \infty$, NaN, and SNaN follow the same ideas as in 32-bit format.

The 128-Bit Standard

With improvements in computer technology, using 128 bits to represent floating point numbers has become feasible. This is sometimes used in numeric-intensive computation, where large rounding errors may occur. In this standard, besides a sign bit, 14 bits are used for the exponent, and 113 bits are used for the significand. The representation is shown in *Figure 4*.

A number in this standard is

$$(-1)^s \times (1.f) \times 2^{(\text{exponent} - 16384)}.$$

Figure 4. Representation of 128-bit floating point number.

Sign	Exponent	Significand
1 bit	14 bits	113 bits



The largest positive number which can be represented in this format is:

0	11111111111110	111.....111
Sign 1 bit	Exponent 14 bits	Significand 113 bits

which equals $(1 - 2^{-113}) \times 2^{16383} = \approx 10^{4932}$.

The smallest normalized positive number which can be represented is

$$2^{1-16383} = 2^{-16382} \approx 10^{-4931}.$$

A subnormal number is represented by

$$(-1)^s \times 0.f \times 2^{-16382}.$$

The smallest positive subnormal number is

$$2^{-16382-113} = 2^{-16485}.$$

The definitions of ± 0 , $\pm \infty$, and QNaN, and SNaN are the same as in the 1985 standard except for the increase of bits in the exponent and significand. There are other minor changes which we will not discuss.

The Decimal Standard

The main motivation for introducing a decimal standard is the fact that a terminating decimal fraction need not give a terminating binary fraction. For example, $(0.1)_{10} = (0.00011\text{ (0011)}_{\text{recurring}})_2$. If one computes 100000000×0.1 using binary arithmetic and rounding the non-terminating binary fraction up to the next larger number, the answer is: 10000001.490116 instead of 10000000. This is unacceptable in many situations, particularly in financial transactions. There was a demand from professionals carrying out financial transactions to introduce a standard for representing decimal floating point numbers. Decimal floating point numbers are not new. They were available in COBOL. There was, however, no standard and this resulted in non-portable programs. IEEE 754-2008 has thus introduced a standard for

The main motivation for introducing a decimal standard is the fact that a terminating decimal fraction need not give a terminating binary fraction.

There was a demand from professionals carrying out financial transactions to introduce a standard for representing decimal floating point numbers.



decimal floating point numbers to facilitate portability of programs.

Decimal Representation

The idea of encoding decimal numbers using bits is simple. For example, to represent 0.1 encoded (not converted) to binary, it is written as 0.1×10^0 . The exponent instead of being interpreted as a power of 2 is now interpreted as a power of 10. Both the significand and the exponent are now integers and their binary representations have a finite number of bits. One method of representing 0.1×10^0 using this idea in a 32-bit word is shown in *Figure 5*.

We have assumed an 8-bit binary exponent in biased form. In this representation, there is no hidden 1 in the significand. The significand is an integer. (Fractions are represented by using an appropriate power of 10 in the exponent field.) For example, $.000753658$ will be written as 0.753658×10^{-3} and 753658, an integer, will be converted to binary and stored as the significand. The power of 10, namely -3 , will be an integer in the exponent part.

The IEEE Standard does not use this simple representation. This is due to the fact that the largest decimal significand that may be represented using this method is 0.8388607.

It is preferable to obtain 7 significant digits, i.e., significand up to 0.9999999. The IEEE Standard achieves this by borrowing some bits from the exponent field and reducing the range of the exponent. In other words, instead of maximum exponent of 128, it is reduced to 96 and the bits saved are used to increase the significant digits in the significand. The coding of bits is as shown in *Figure 6*.

Figure 5. Representation of decimal floating point numbers.

Figure 6. Representation of 32-bit floating point numbers in IEEE 2008 Standard.

0	01111110	00000000000000000000000000000001
Sign	Exponent	Significand
1 bit	8 bits	23 bits

Sign	Combination Bits	Exponent Continuation	Coefficient
1 bit	5 bits	6 bits	20 bits



Observe the difference in terminology for exponent and significand fields. In this coding, one part of the five combination bits is interpreted as exponent and the other part as an extension of significand bits using an ingenious method. This will be explained later in this section.

IEEE 2008 Standard defines two formats for representing decimal floating point numbers. One of them *converts* the binary significand field to a decimal integer between 0 and 10^{p-1} where p is the number of significant bits in the significand. The other, called densely packed decimal form, *encodes* the bits in the significand field directly to decimal. Both methods represent decimal floating point numbers as shown below.

Decimal floating point number = $(-1)^s \times f \times 10^{\exp - bias}$, where s is the sign bit, f is a decimal fraction (significand), \exp and $bias$ are decimal numbers. The fraction need not be normalized but is usually normalized with a non-zero most significant digit. There is no hidden bit as in the binary standard. The standard defines decimal floating point formats for 32-bit, 64-bit and 128-bit numbers.

Densely Packed Decimal Significand Representation

Normally, when decimal numbers are coded to binary, 4 bits are required to represent each digit, as 3 bits give only 8 combinations that are not sufficient to represent 10 digits. Four bits give 16 combinations out of which we need only 10 combinations to represent decimal numbers from 0 to 9. Thus, we waste 6 combinations out of 16. For a 32-bit number, if 1 bit is used for sign, 8 bits for the exponent and 23 bits for the significand, the maximum positive value of the significand will be +0.799999. The exponent range will be 00 to 99. With a bias of 50, the maximum positive decimal number will be 0.799999×10^{49} .

The number of significand digits in the above coding is not sufficient. A clever coding scheme for decimal encoding of binary numbers was discovered by Cowlinshaw. This method encodes 10 bits as 3 decimal digits (remember that $2^{10} = 1024$ and

IEEE 2008 Standard defines two formats for representing decimal floating point numbers. One of them *converts* the binary significand field to a decimal integer between 0 and 10^{p-1} where p is the number of significant bits in the significand. The other, called densely packed decimal form, *encodes* the bits in the significand field directly to decimal.



one can encode 000 to 999 using the available 1024 combinations of 10 bits). IEEE 754 Standard uses this coding scheme. We will now describe the standard using this coding method for 32-bit numbers.

The standard defines the following:

1. A *sign bit* 0 for + and 1 for - .
2. A *combination field* of 5 bits. It is called a combination field as one part of it is used for the exponent and another part for the significand.
3. An *exponent continuation field* that is used for the exponent.
4. A *coefficient field* that encodes strings of 10 bits as 3 digits. Thus, 20 bits are encoded as 6 digits. This is part of the significand field.

For 32-bit IEEE word, the distribution of bits was given in *Figure 6* which is reproduced below for ready reference.

Sign	Combination field	Exponent continuation field	Coefficient field
1 bit	5 bits	6 bits	20 bits

The 5 bits of the combination field are used to increase the coefficient field by 1 digit. It uses the first ten 4-bit combinations of 16 possible 4-bit combinations to represent this. As there are 32 combinations of 5 bits, out of which only 10 are needed to extend the coefficient digits, the remaining 22-bit combinations are available to increase the exponent range and also encode ∞ and NaN. This is done by an ingenious method explained in *Table 3*.

Combination bits	Type	Most significant bits of exponent	Most significant digit of coefficient
$b_1 \ b_2 \ b_3 \ b_4 \ b_5$			
x y a b c	Finite	x y	0 a b c
1 1 a b c	Finite	a b	1 0 0 c
1 1 1 1 0	∞	—	—
1 1 1 1 1	NaN	—	—

Table 3. Use of combination bits in IEEE Decimal Standard.



Sign bit	1	1
Combination field (bits)	5	5
Exponent continuation field (bits)	8	12
Coefficient field (bits)	50	110
Number of bits in number	64	128
Significant digits	16	34
Exponent range (decimal)	768	12288
Exponent bias (decimal)	384	6144

Table 4. Decimal representation for 64-bit and 128-bit numbers.

Thus, the most significant 2 bits of the exponent are constrained to take on values 00, 01, 10. Using the 6 bits of exponent combination bits, we get an exponent range of $3 \times 2^6 = 192$. The exponent bias is 96. (Remember that the exponent is an integer that can be exactly represented in binary.) With the addition of 1 significant digit to the 6 digits of the coefficient bits, the decimal representation of 32-bit numbers has 7 significant digits. The largest number that can be stored is 0.9999999×10^{96} . A similar method is used to represent 64- and 128-bit floating point numbers. This is shown in *Table 4*.

In the IEEE 2008 Decimal Standard, the combination bits along with exponent bits are used to represent NaN, signaling NaN and $\pm \infty$ as shown in *Table 5*.

Quantity	Sign bit	Combination bits
	b_0	$b_1 b_2 b_3 b_4 b_5 b_6$
SNaN	0 or 1	1 1 1 1 1 1
QNaN	0 or 1	1 1 1 1 1 0
$+\infty$	0	1 1 1 1 0 x
$-\infty$	1	1 1 1 1 0 x
$+0$	0	0 0 0 0 0 x
-0	1	0 0 0 0 0 x (x is 0 or 1)

Table 5. Representation of NaNs, $\pm \infty$, and ± 0 .



Binary Integer Representation of Significand

In this method, as we stated earlier, if there are p bits in the extended coefficient field constituting the significand, they are converted as a decimal integer. This representation also uses the bits of the combination field to increase the number of significant digits. The idea is similar and we will not repeat it. The representation of NaNs and $\pm\infty$ is the same as in the densely packed format. The major advantage of this representation is the simplicity of performing arithmetic using hardware arithmetic units designed to perform arithmetic using binary numbers. The disadvantage is the time taken for conversion from integer to decimal. The densely packed decimal format allows simple encoding of binary to decimal but needs a specially designed decimal arithmetic unit.

Acknowledgements

I thank Prof. P C P Bhatt, Prof. C R Muthukrishnan, and Prof. S K Nandy for reading the first draft of this article and suggesting improvements.

Suggested Reading

- [1] David Goldberg, What every computer scientist should know about floating point arithmetic, *ACM Computing Surveys*, Vol.28, No.1, pp5–48, March 1991.
- [2] V Rajaraman, *Computer Oriented Numerical Methods*, 3rd Edition, PHI Learning, New Delhi, pp.15–32, 2013.
- [3] Marc Cowlinshaw, Densely packed decimal encoding, *IEEE Proceedings – Computer and Digital Techniques*, Vol.149, No.3, pp.102-104, May 2002.
- [4] Steve Hollasch, IEEE Standard 754 for floating point numbers, steve.hollasch.net/cgindex/coding/ieeefloat.html
- [5] Decimal Floating Point, Wikipedia, en.wikipedia.org/wiki/Decimal_floating_point

Address for Correspondence

V Rajaraman
Supercomputer Education
and Research Centre
Indian Institute of Science
Bengaluru 560 012, India.
Email: rajaram@serc.iisc.in