




DS221: Introduction to Scalable Systems

Topic: Algorithms and Data Structures





L2: Complexity Analysis & Performance Evaluation



Algorithm Analysis

- Algorithms can be evaluated on two performance measures
- Time taken to run an algorithm
- Memory space required to run an algorithm
- Later, I/O and Communication complexity
...for a given input size

- *Why are these important?*



Space Complexity

- *Estimate of the amount of peak memory required for an algorithm to run to completion, for a given input size*
 - Core dumps/OOMEx: Memory required is larger than the memory available on a given system
 - Algorithm design problem OR “memory leaks” in implementation
- In some applications, we may want load all data in memory for performance




Space Complexity

- **Fixed part:** The size required to store certain data/variables, that is independent of the size of the problem:
 - e.g., for all valid words, given a set of letters
 - e.g., etymology for each word in a dictionary
- **Variable part:** Space needed by variables, whose size is dependent on the size of the problem:
 - e.g., number of letters in a scrabble game
 - e.g., text of Shakespeare's plays

Try yourself!

For some program with variable input sizes, find the space taken by the fixed part and the variable part, using **top** command




Analyzing Running Time Empirically

- Write program
- Run Program
- Measure actual running time with some methods like

std::chrono::high_resolution_clock::now() in C++

- *Is that good enough as a programmer?*

Running Time



- Suppose the program includes an *if-then statement that may execute or not* → variable running time
- Typically algorithms are measured by their **worst case**

Try yourself!

Run some program for the same input size (but different inputs), and see how the run time changes for each input



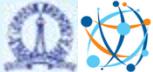
General Methodology for Analysis

- Uses High Level Description (pseudo-code) instead of implementation
- Takes into account all variations of inputs of some size “n”
- Allows one to evaluate the efficiency independent of hardware/software environment



Pseudo-Code

- Mix of natural language and high level programming concepts that describes the main idea behind algorithm
 - More detail than algo, less than implementation
- Control flow
 - If ... then ...else
 - While-loop
 - for-loop
- Simple data structures
 - Array : A[i]; A[l,j]
- Methods
 - Calls: methodName(args)
 - Returns: return value



Pseudo-Code

- Mix of natural language and high level programming concepts that describes the main idea behind algorithm
 - More detail than algo, less than implementation
- Control flow
 - If ... then ...else
 - While-loop
 - for-loop
- Simple data structures
 - Array : A[i]; A[l,j]
- Methods
 - Calls: methodName(args)
 - Returns: return value

```
int arrayMax(int[] A, int n)
    Max=A[0]
    for i=1 to n-1 do
        if Max < A[i]
            then Max = A[i]
    return Max
```



Pseudo-Code

- Mix of natural language and high level programming concepts that describes the main idea behind algorithm
 - More detail than algo, less than implementation
- Control flow
 - If ... then ...else
 - While-loop
 - for-loop
- Simple data structures
 - Array : A[i]; A[l,j]
- Methods
 - Calls: methodName(args)
 - Returns: return value



```
int arrayMax(int[] A, int n)
    Max=A.front();
    for i=1 to n-1 do
        if Max < A[i]
            then Max = A[i]
    return Max
```



Pseudo-Code

- Mix of natural language and high level programming concepts that describes the main idea behind algorithm
 - More detail than algo, less than implementation
- Control flow
 - If ... then ...else
 - While-loop
 - for-loop
- Simple data structures
 - Array : A[i]; A[l,j]
- Methods
 - Calls: methodName(args)
 - Returns: return value



```
int arrayMax(int[] A, int n)
    Max=A[0]
    for i in range(0,n) do
        if Max < A[i]
            then Max = A[i]
    return Max
```



Pseudo-Code

- Mix of natural language and high level programming concepts that describes the main idea behind algorithm
 - More detail than algo, less than implementation
- Control flow
 - If ... then ...else
 - While-loop
 - for-loop
- Simple data structures
 - Array : A[i]; A[l,j]
- Methods
 - Calls: methodName(args)
 - Returns: return value




```
int arrayMax(int[] A, int n)
    Max=A[0]
    for i=1 to n-1 do
        if Max < A[i]
            then Max = A[i]
    return Max
```



Analysis of Algorithms

- Analyze time taken by **Primitive Operations**
- Low level operations independent of programming language
 - Data movement (assign..)
 - Control (branch, subroutine call, return...)
 - Arithmetic/logical operations (add, compare..)
- By inspecting the pseudo-code, we can count the number of primitive operations executed by an algorithm



Example: Array Transpose

```
function Transpose(A[][][], n)
    for i = 0 to n-1 do
        for j = i+1 to n-1 do
            tmp = A[i][j]
            A[i][j] = A[j][i]
            A[j][i] = tmp
    end
end
end
```

		$j=0$	$j=3$	
		$i=0$	$0,1$	$0,2$
		$1,0$	$1,1$	$1,2$
		$2,0$	$2,1$	$2,2$
		$3,0$	$3,1$	$3,2$
				$3,3$

Example: Array Transpose

```
function Transpose(A[][][], n)
    for i = 0 to n-1 do
        for j = i+1 to n-1 do
            tmp = A[i][j]
            A[i][j] = A[j][i]
            A[j][i] = tmp
    end
end
end
```

		j=0				j=3			
		i=0	0,0	0,1	0,2	0,3			
		i=1	1,0	1,1	1,2	1,3			
		i=2	2,0	2,1	2,2	2,3			
		i=3	3,0	3,1	3,2	3,3			

Swap

Outer
Loop

Inner
Loop

Estimated time for $A[n][n] = (n(n-1)/2).(3+2) + 2.n$
Is this constant for a given 'n'?



Asymptotic Analysis

- **Goal:** Simplify analysis of running time by getting rid of ‘details’ which may be affected by specific implementation and hardware
 - Like ‘rounding’: $1001 = 1000$
 - $3n^2 = n^2$
- How does the running time of an algorithm increase with the size of input in the limit?
 - Asymptotically more efficient algorithms are best for all but small inputs


Asymptotic Notation: “Big O”

Definition Let $p(n)$ and $q(n)$ be two nonnegative functions.

$p(n)$ is **asymptotically bigger** ($p(n)$ asymptotically dominates $q(n)$) than the function $q(n)$ iff

$$\lim_{n \rightarrow \infty} \frac{q(n)}{p(n)} = 0$$

- **O Notation**
 - Asymptotic **upper bound**
 - $q(n) = O(p(n))$, if there exists constants c and n_0 , s.t.
 - $q(n) \leq c.p(n)$ for $n \geq n_0$
 - $q(n)$ and $p(n)$ are functions over non negative integers
- Used for **worst-case analysis**
 - $p(n)$ is the **asymptotic upper bound** of actual time taken





Asymptotic Notation

- **Simple Rule:** Drop lower order terms and constant factors
 - $(n(n-1)/2).(3+2) + 2.n$ is $O(n^2)$
 - $23.n.\log(n)$ is $O(n.\log(n))$
 - $9n-6$ is $O(n)$
 - $6n^2.\log(n) + 3n^2 + n$ is $O(n^2.\log(n))$



Asymptotic Notation

- **Simple Rule:** Drop lower order terms and constant factors
 - $(n(n-1)/2).(3+2) + 2.n$ is $O(n^2)$
 - $23.n.\log(n)$ is $O(n.\log(n))$
 - $9n-6$ is $O(n)$
 - $6n^2.\log(n) + 3n^2 + n$ is $O(n^2.\log(n))$
- **Note:** It is expected that the approximation should be as tight an order as possible (e.g., avoid giving a loose $O(n^3)$ upper bound for function $3n^2 + 7n + 42$)

Try yourself!

Plot the observed and asymptotic expected curves for a program



Asymptotic Analysis of Running Time

- Use O notation to express number of primitive operations executed as a function of input size.
- Hierarchy of functions

$$1 < \log n < n < n^2 < n^3 < 2^n$$

$\xleftarrow{\text{Better}}$

- **Warning!** Beware of large constants (say 1M).
 - This may have lower performance than one running in time $2n^2$, which is $O(n^2)$, for even modest input sizes



Example of Asymptotic Analysis

- Input: An array $X[n]$ of numbers.
- Output: An array $A[n]$ of numbers s.t $A[k]=\text{mean}(X[0]+X[1]+\dots+X[k-1])$


```
for i=0 to (n-1) do
    a=0
    for j=0 to i do
        a = a + X[j]
    end
    A[i] = a/(i+1)
end
return A
```

Analysis: running time is $O(n^2)$

- A naïve algorithm! *What's its complexity?*

A Better Algorithm?

```
s=0  
for i=0 to n do  
    s = s + X[i]  
    A[i] = s/(i+1)  
end  
return A
```



Analysis: running time is $O(n)$

Time taken by previous naive algorithm ($O(n^2)$) and current algorithm ($O(n)$). Orange line in secondary Y axis indicates n^2 line. The blue line matches the orange line. The green line should match $O(n)$ but since time is small, there may be measurement errors.




Comparison

Logarithmic	Linear	Iterated logarithmic	Quadratic	Polynomial (n^α) α is a const. > 1	Exponential (α^n) α is a const. > 1
log n	n	$n \log n$	n^2	n^3	2^n
0	1	0	1	1	2
1	2	2	4	8	4
2	4	8	16	64	16
3	8	24	64	512	256
4	16	64	256	4096	65536
5	32	160	1024	32768	4294967296

(solvable in polynomial time)

Comparison






Tasks

- Self study (Sahni Textbook)
 - ▶ Chapter 3 & 4 “Asymptotic Notation” & “Performance Measurement”
- Try the mean calculation code in C++ yourselves

Asymptotic Notation: Lower Bound


- The “big-Omega” Ω notation
 - asymptotic *lower* bound
 - $q(n) = \Omega(p(n))$ if there exists const. c and n_0 s.t.
 - $c.p(n) \leq q(n)$ for $n \geq n_0$
- Tells you: the algorithm will take at least this much time for large n (up to constant factors)
- Can also be useful to give a lower bound for the entire class of algorithms

(e.g., searching in an unsorted array is $\Omega(n)$, that is, no algorithm can offer faster worst-case time complexity)



Asymptotic Notation: Tight Bound

- The “big-Theta” Θ -Notation
 - Asymptotically tight bound
 - $q(n) = \Theta(p(n))$ if there exists constants c_1, c_2 and n_0 s.t.
 $c_1 p(n) \leq q(n) \leq c_2 p(n)$ for $n \geq n_0$
- $q(n) = \Theta(p(n))$ if and only if $q(n)=O(p(n))$ and $q(n)=\Omega(p(n))$





Asymptotic Notation

- Analogy with real numbers

- $q(n) = O(p(n)) \rightarrow q \leq p$
- $q(n) = \Omega(p(n)) \rightarrow q \geq p$
- $q(n) = \Theta(p(n)) \rightarrow q \approx p$



Analyze the runtime (In-class exercise)

```
void mystery(int n) {  
    count = 0;  
  
    for (int i = 1; i <= n; i++) {  
        for (int j = 1; j <= i*i; j++) {  
            if (j % i == 0) {  
                count++;  
            }  
        }  
    }  
}
```

Say $q(n)$ denotes the worst-case runtime of this algorithm
 $q(n) = O(p_1(n))$. Specify $p_1(n)$
 $q(n) = \Omega(p_2(n))$. Specify $p_2(n)$
 $q(n) = \Theta(p_3(n))$. Specify $p_3(n)$




Polynomial and Intractable Algorithms

- **Polynomial Time complexity**
 - An algorithm is said to be polynomial if it is $O(n^\alpha)$ for some integer α
 - Polynomial algorithms are said to be efficient
 - They solve problems in “reasonable” times
- **Intractable Algorithms**
 - Algorithms for which there is no *known* polynomial time algorithm



Complexity: List using Arrays

- **Storage Complexity:** Amount of storage required by the data structure, relative to items stored
- List using Array: ...
- **Computational Complexity:** Number of CPU cycles required to perform each data structure operation
- `size()`, `set()`, `get()`, `indexOf()`



Complexity: List using Linked List

- Storage Complexity
 - Only store as many items as you need
 - But...
- Computational Complexity
 - set(), get(), remove()
 - indexOf()
- Other Pros & Cons?
 - Memory management, mixed item types




Empirical Validation

- How do I check if complexity analysis matches reality?
 - Timing
 - Static overheads
 - Asymptotic behaviour
 - Locality effects
 - Disk Thrashing
 - Memory used
 - Adding up variable sizes, reference pointer sizes, “sizeof”
 - Deep vs. shallow size
 - Effect of padding for struct to align with word length/cache line
 - Profiling: CPU used, top, cache hits/misses, iops, context switching

Perf of Array, LinkedList


- Time to insert into array
 - ▶ Time to copy from one array to larger, on resize?
- Time to insert into linked list



Time to insert n items into a list. X axis is “ n ”.
We are doing an append to the end.
So linked list and array implementations have $O(n)$ time.
But time to allocate memory for an item is higher in LL than time to set allocated array location.

Perf of Array List

■ Array Copy Times on growth




Time to insert $100k$ items incrementally into a single array list. X axis is the number of inserted so far.

We expect each 100k insertion to take constant time.

Some spikes indicate array capacity being full and reallocation/moving of prior data.

The red dots show where the capacity may have been increased based on implementation logic.



Complexity of Matrix Ops

- Generating a 2D matrix with $n \times n$ elements

$O(n^2)$

```
function MatMult(A[][], B[][], n)
```

```
    for i = 0 to n-1 {
```

```
        for j = 0 to n-1 {
```

```
            sum=0
```

```
            for k = 0 to n-1 {
```

```
                sum = sum + A[i][k]*B[k][j]
```


```
            }
```

```
            c[i][j] = sum
```


```
        }
```

```
}
```

$O(n^3)$

Mat Generate...O(n^2)

Time to generate and set values in matrix of size $n \times n$. X axis is number of elements in matrix. Orange line is a linear trend line.

Mat Mult ... O(n^3)

Time to multiply two matrices of size $n \times n$. X axis is n^3 . Orange line is a linear trend line.



Tasks

- **Self study (Sahni Textbook)**
 - Chapter 3 & 4 “Asymptotic Notation” & “Performance Measurement”
- Try the code in C++ yourselves
- Assignment 1
 - Due on **20 September**
 - No extension will be granted
 - 20% weightage
 - Get started early to avoid any last day hiccups
 - Access to teaching cluster- soon
- 1st tutorial on Aug 29 (Friday)