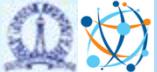




# DS221: Introduction to Scalable Systems

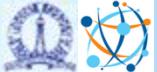
Topic: Algorithms and Data Structures





# L4: Fast Searching

Search Trees, B-Tree, Hashmap



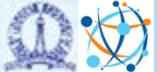
# Dictionary Abstract Data Structure

- Store `<key,value>` as a pair
- *Lookup* the value for a given key
- Goal: Lookup has to fast
- Different implementations
  - Ordered List
  - Hash table (or Hash Map)
  - Binary Search Tree



# Dictionary using List

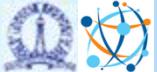
- Dictionary stored as a List of <key,value> items
  - Unsorted Linked List and Array
  - Insertion time? Searching time?
- Dictionary stored as an *Ordered* List of <key,value> elements, ordered by key
  - Linked List vs. Array
  - What's the advantage?



# Dictionary as a Sorted Array

- Idea: **Divide and Conquer**
- Narrow down the search range by half at each stage
- E.g. **find (8)**
- Start with **floor(|search space| / 2)**
- 2 5 8 **9** 11 17 20 22
- 2 **5** 8 9 11 17 20 22
- 2 5 **8** 9 11 17 20 22

*Binary search over array  
Takes  $O(\log_2(n))$  searches*



# Dictionary as a Sorted List

```
int bsearch(KVP[] list, int start, int end, int k) {  
    if (end < start) return -1 // No match!  
    i = start+(end-start)/2 // midpoint  
    if (list[i].key == k) // Found!  
        return list[i].value  
    if (list[i].key < k) // check 2nd half  
        return bsearch(list, i+1, end, k)  
    else // check 1st half  
        return bsearch(list, start, i-1, k)  
}
```

Usual problem with arrays!

- Unused capacity
- Costly to update and maintain sorted list...many shifts

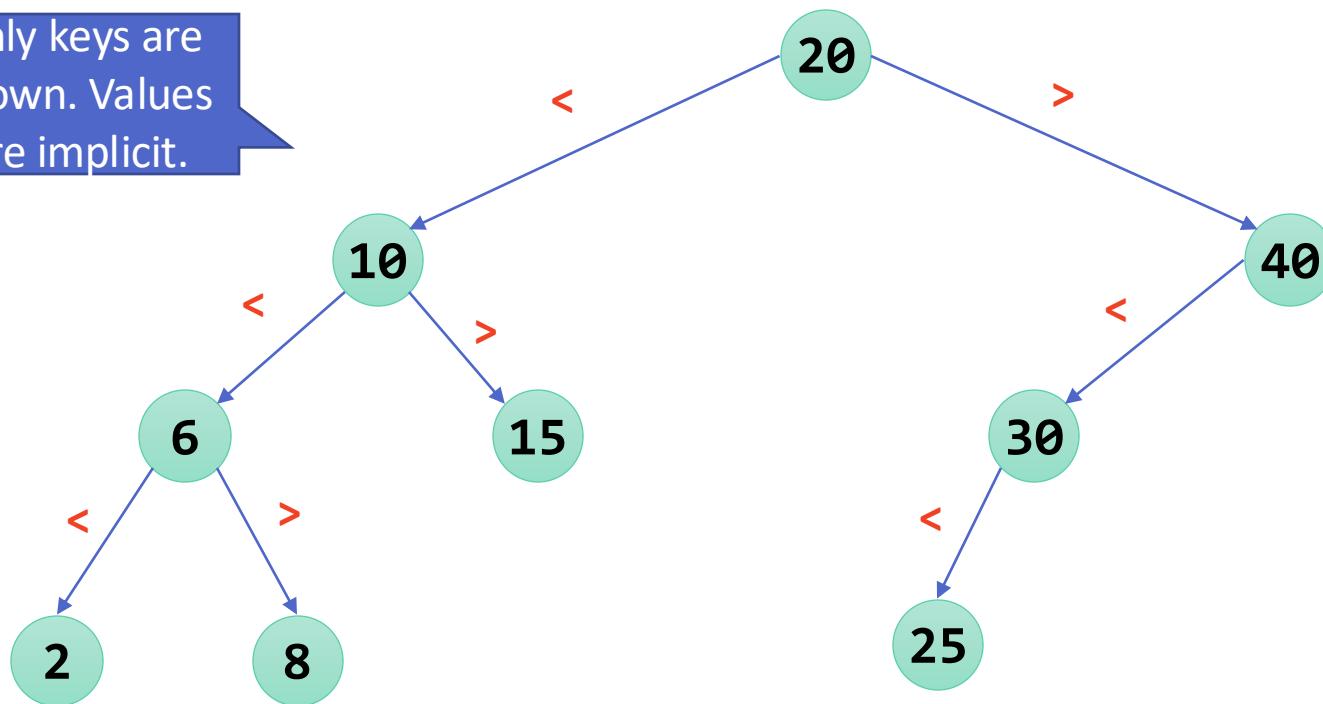


# Dictionary using Binary Search Tree (BST)

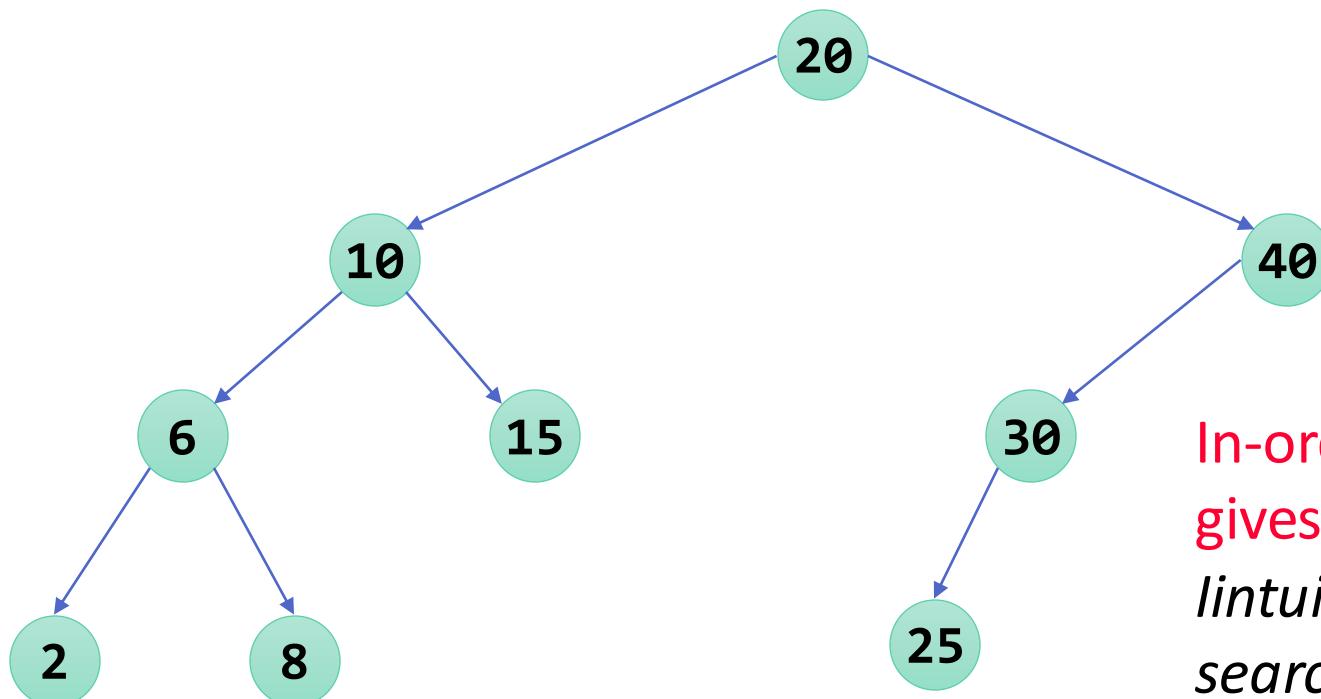
- Combining speed of binary search over array with dynamic capacity of a linked list
- A binary tree with each node having a **(key, value)** pair
- For each node  $x$ ,
  - All keys in the *left subtree* of  $x$  are *smaller* than the key of  $x$
  - All keys in the *right subtree* of  $x$  are *greater* than the key of  $x$
- Dictionary Operations
  - `find(key)`
  - `insert(key, value)`
  - `delete(key)`

# Example Binary Search Tree

Only keys are shown. Values are implicit.



# The Operation find()



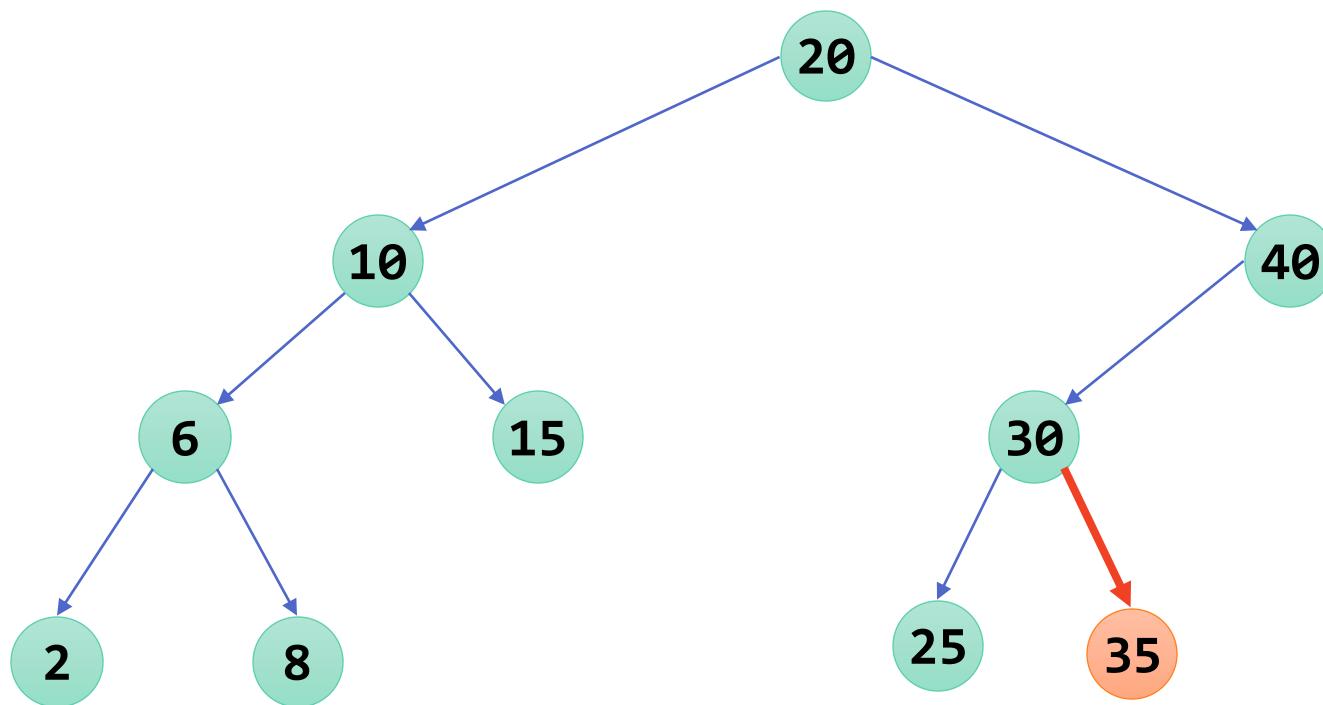
In-order traversal of BST gives a sorted array.  
*Intuition behind “binary search” by partitioning into two.*

2	6	8	10	-	15	-	20	25	30	-	40	-	-	-
---	---	---	----	---	----	---	----	----	----	---	----	---	---	---

Complexity is  $O(\text{height}) = O(n)$ , where  $n$  is the number of nodes/elements.

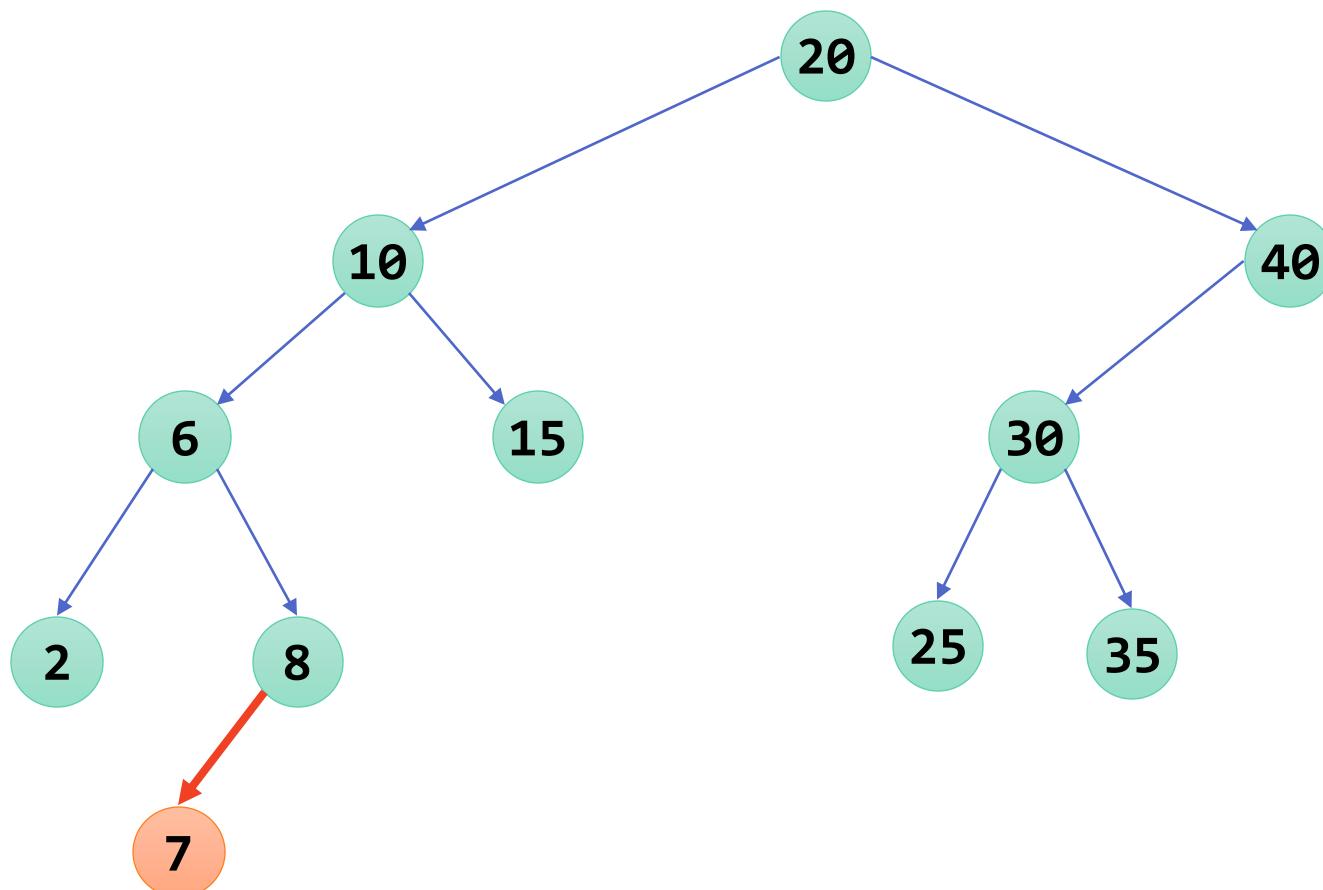


# The Operation insert()

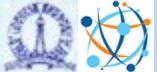


Insert a pair whose key is 35.

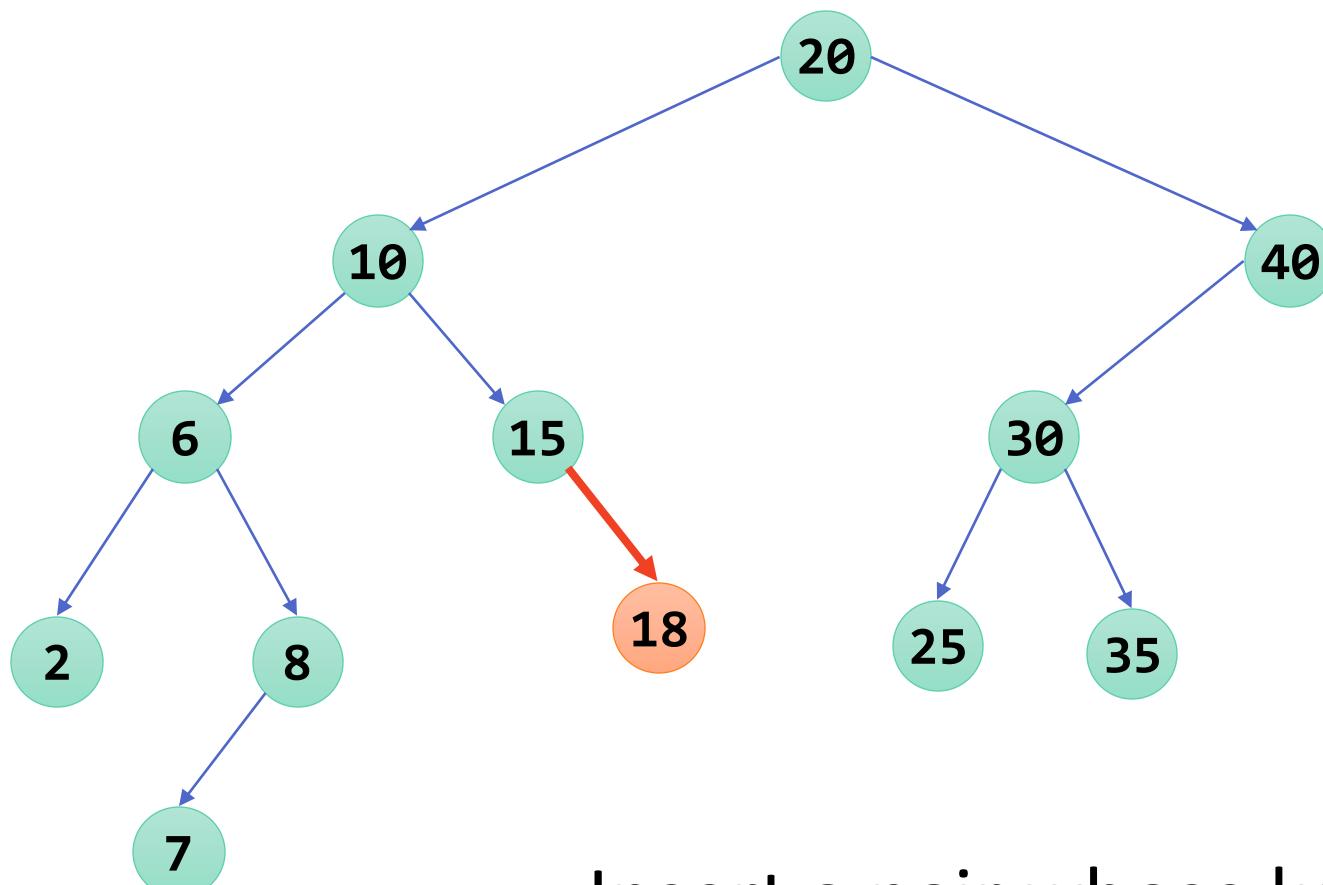
# The Operation insert()



Insert a pair whose key is 7.

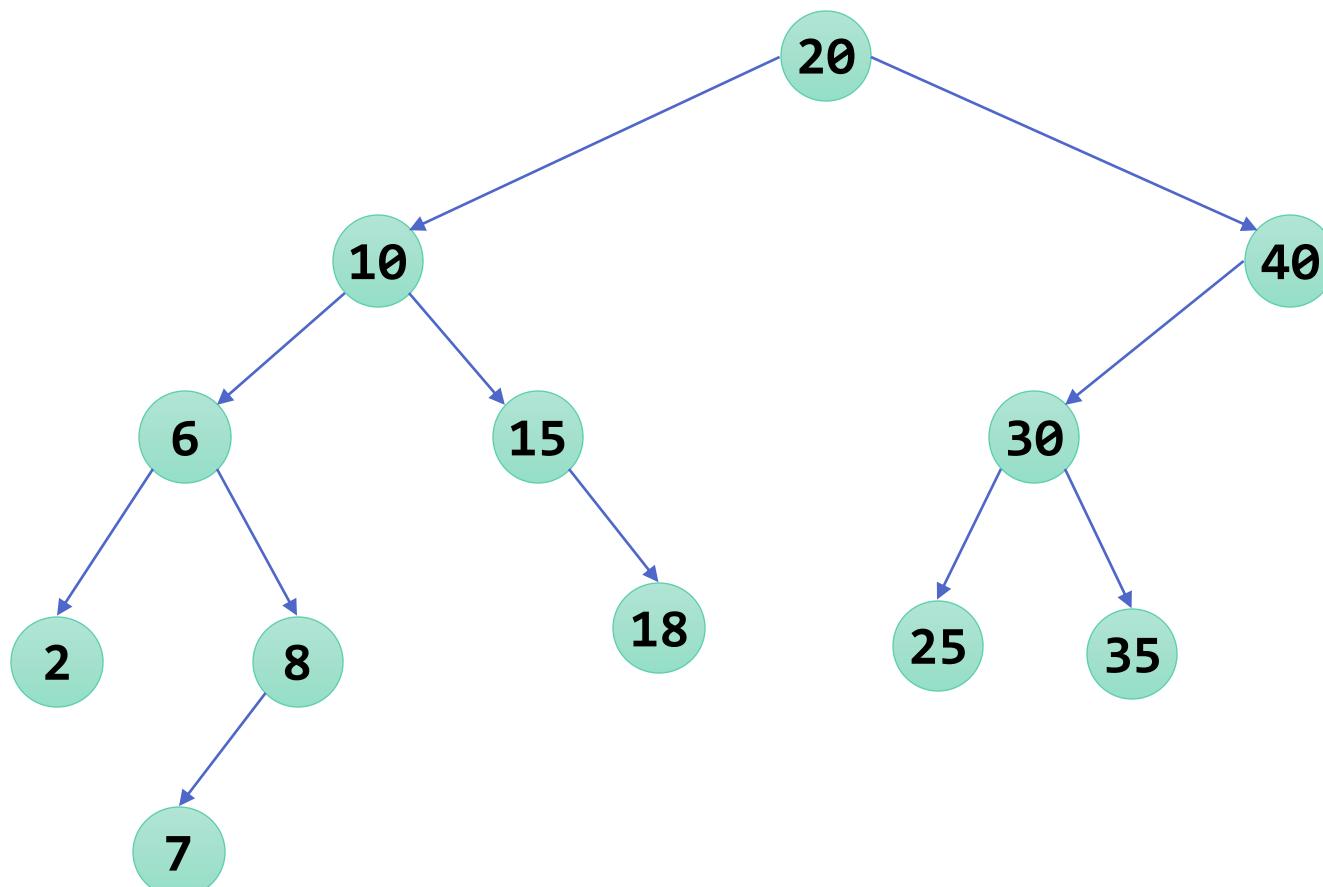


# The Operation insert()

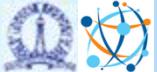


Insert a pair whose key is 18.

# The Operation insert()

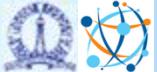


Complexity of `insert()` is  $O(\text{height})$ .

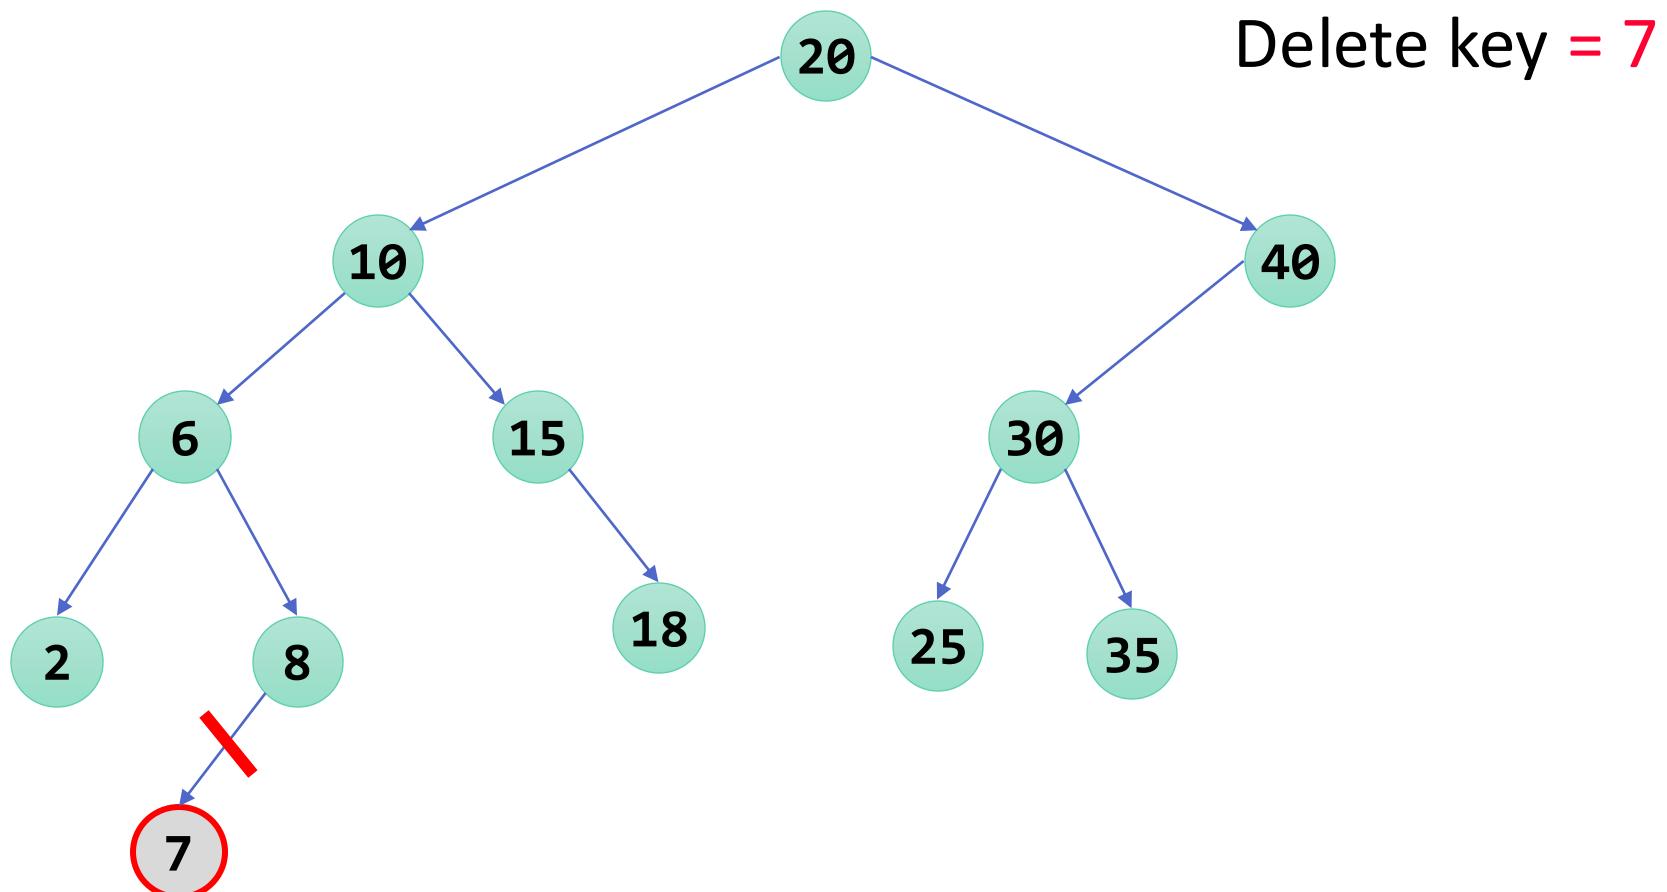


# The Operation delete()

- Three cases:
  - Element is in a leaf.
  - Element is in a degree 1 node.
  - Element is in a degree 2 node.

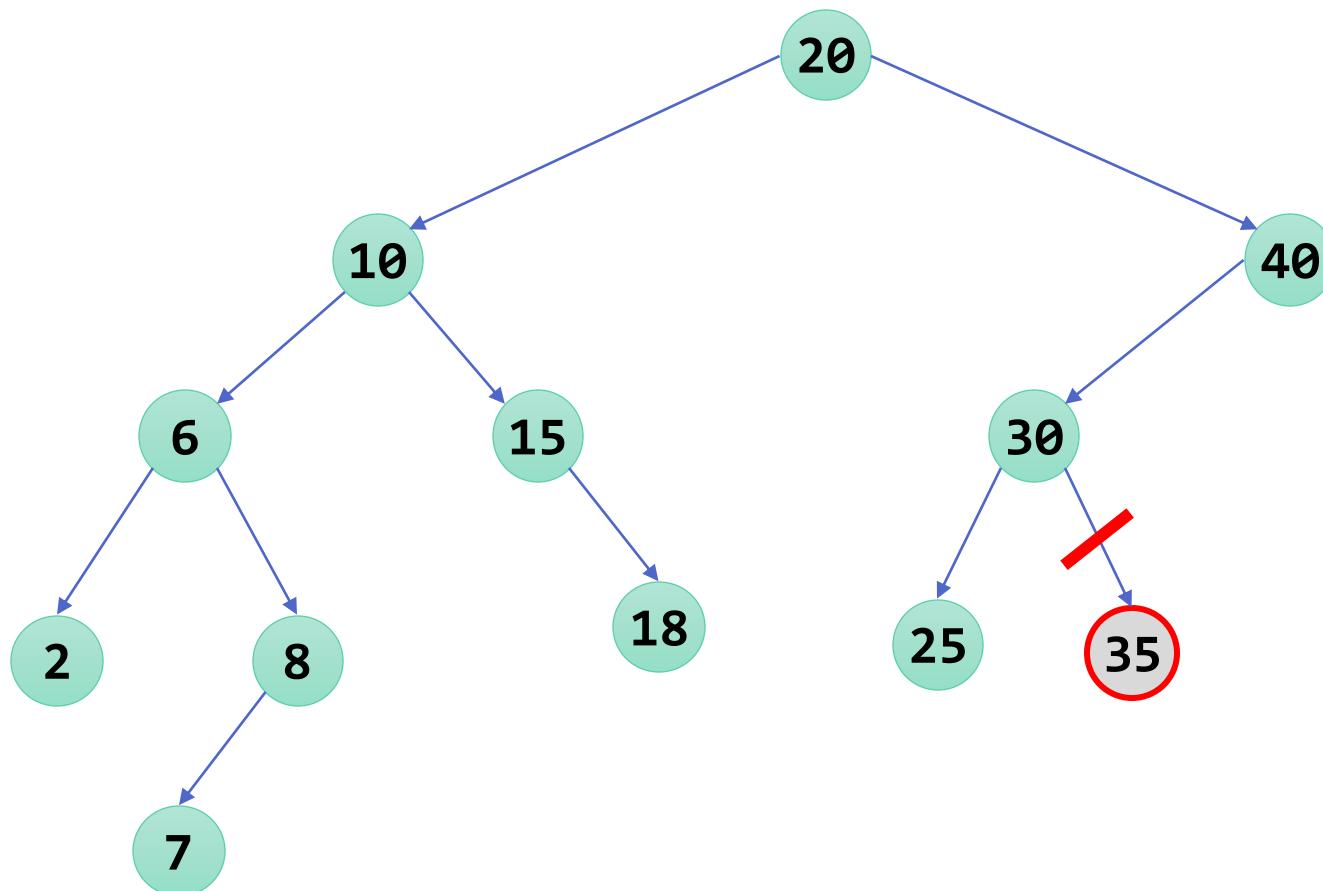


# Delete From A Leaf



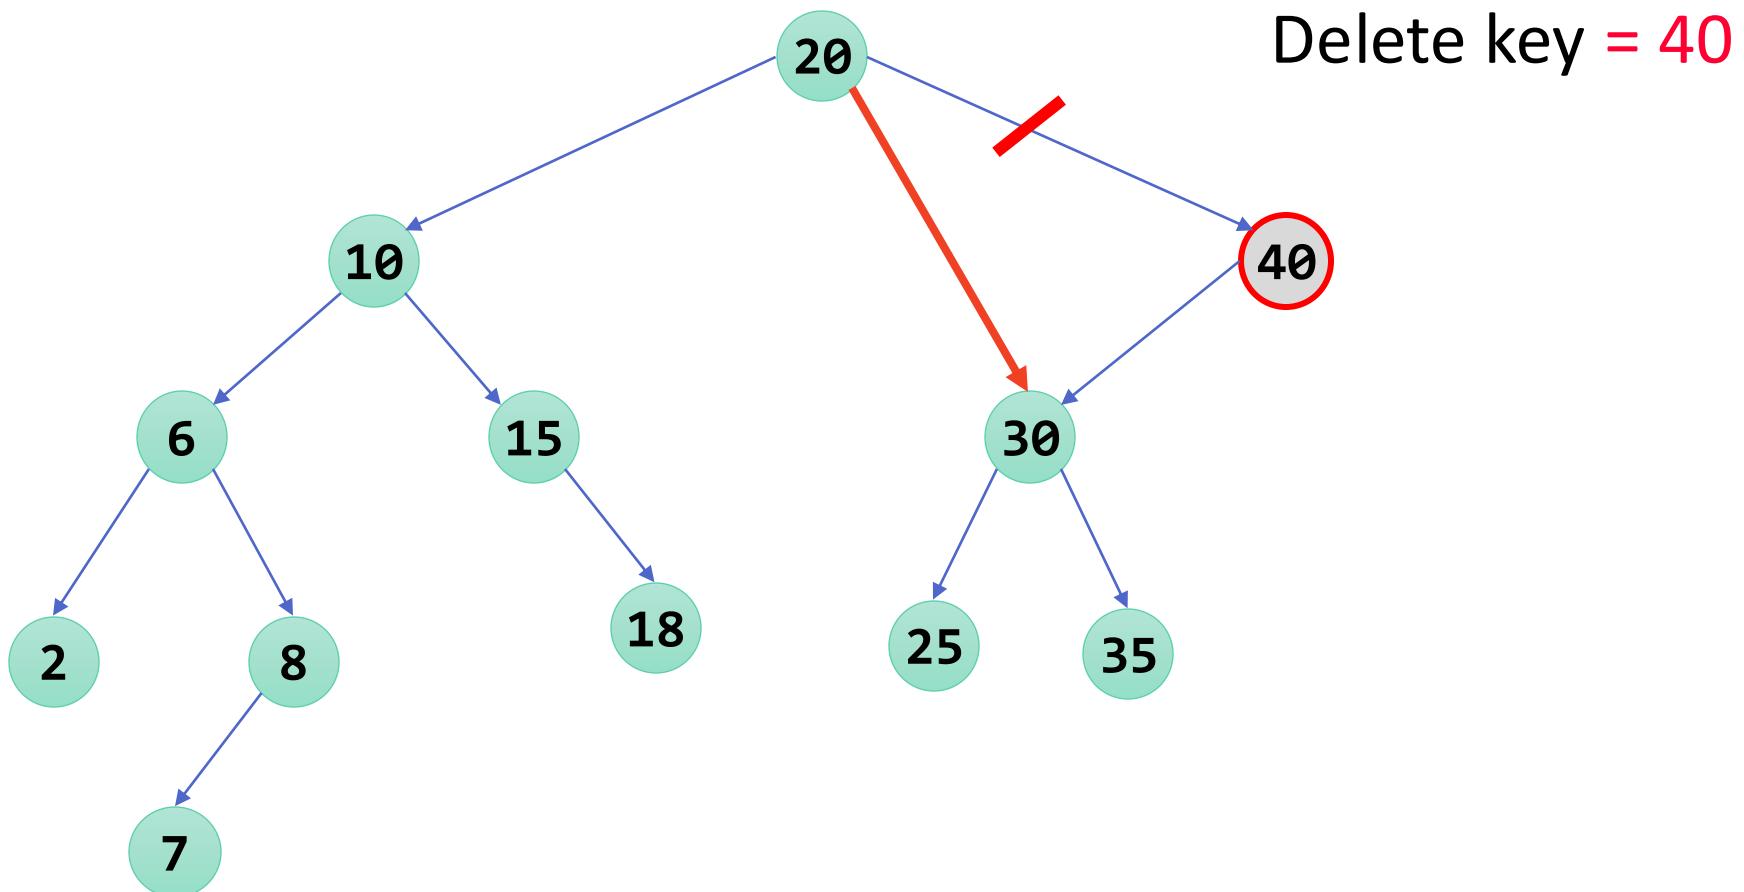
Delete a leaf element.  
*Set parent to NULL*

# Delete From A Leaf

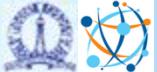


Delete a leaf element. key = 35

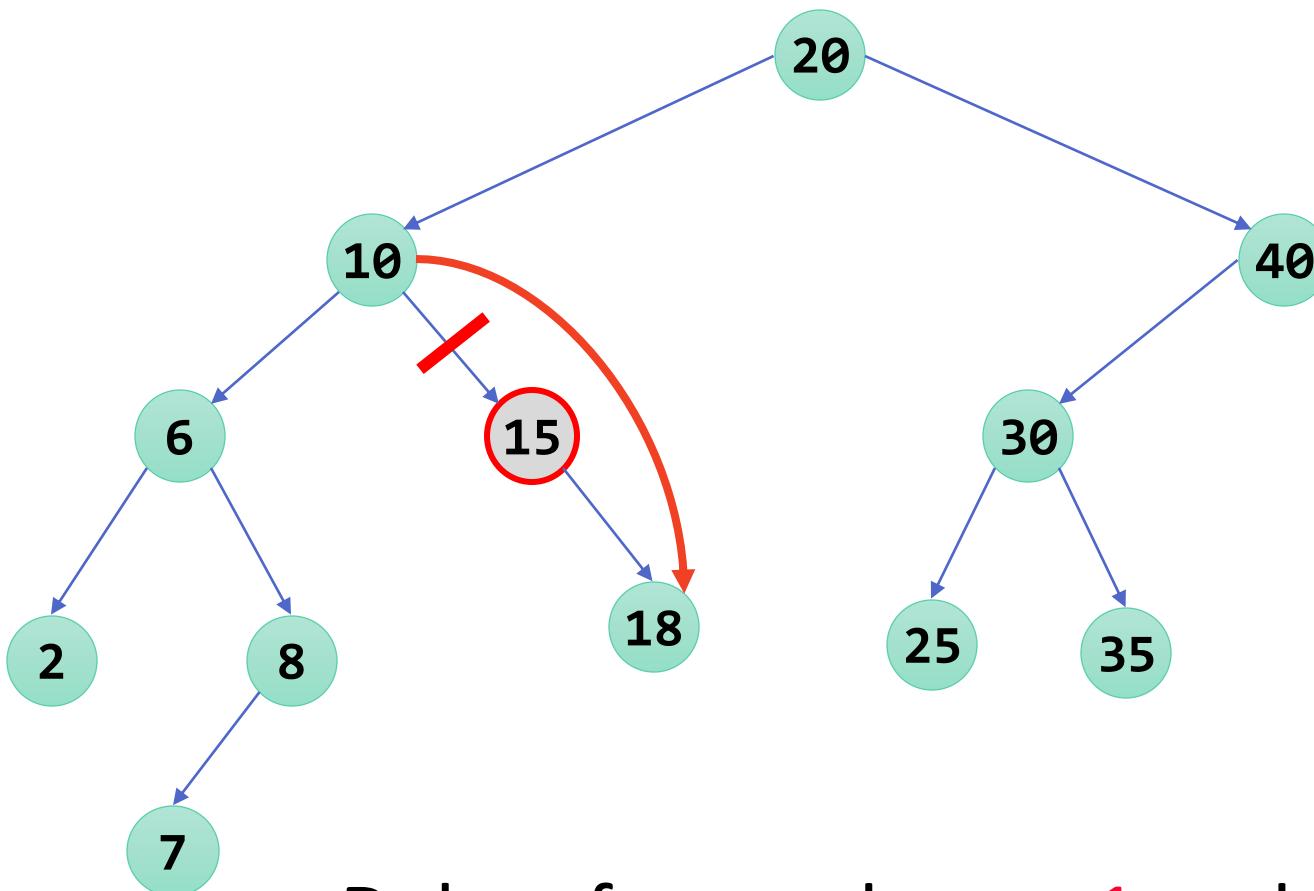
# Delete From Degree 1 Node

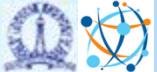


Delete from a degree 1 node.  
*Point parent to child.*

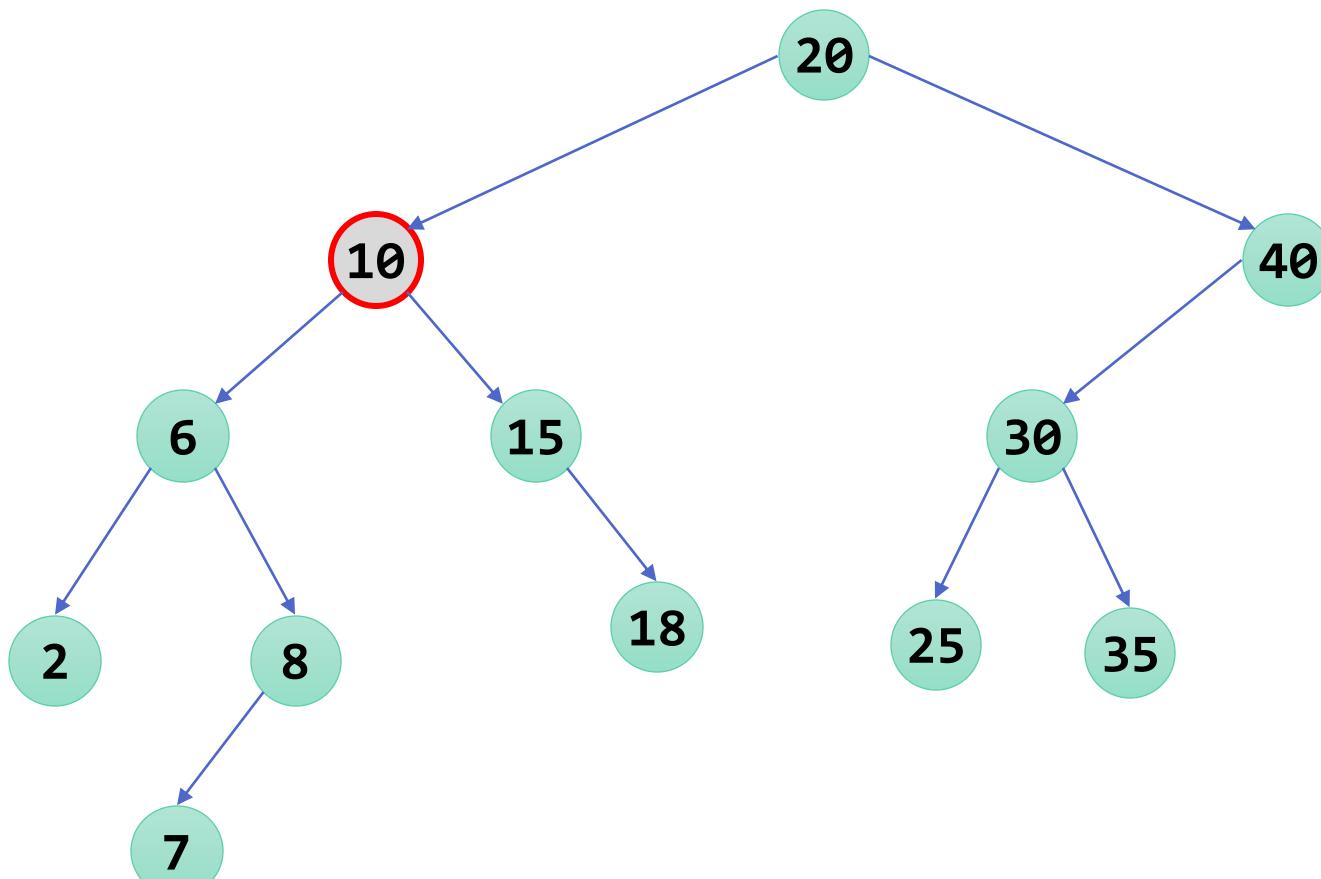


# Delete From Degree 1 Node

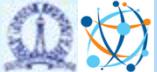




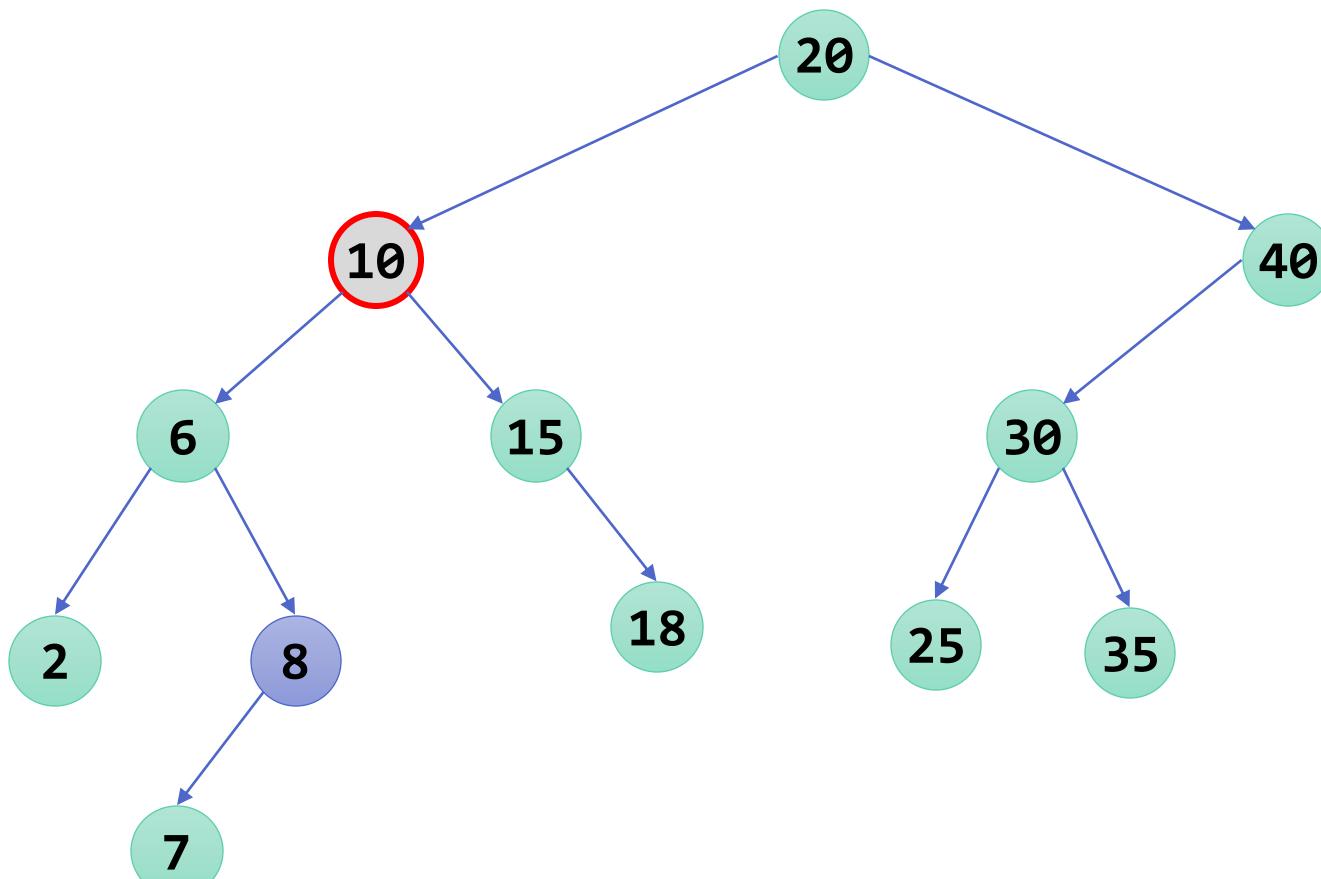
# Delete From Degree 2 Node



Delete from a degree 2 node. key = 10

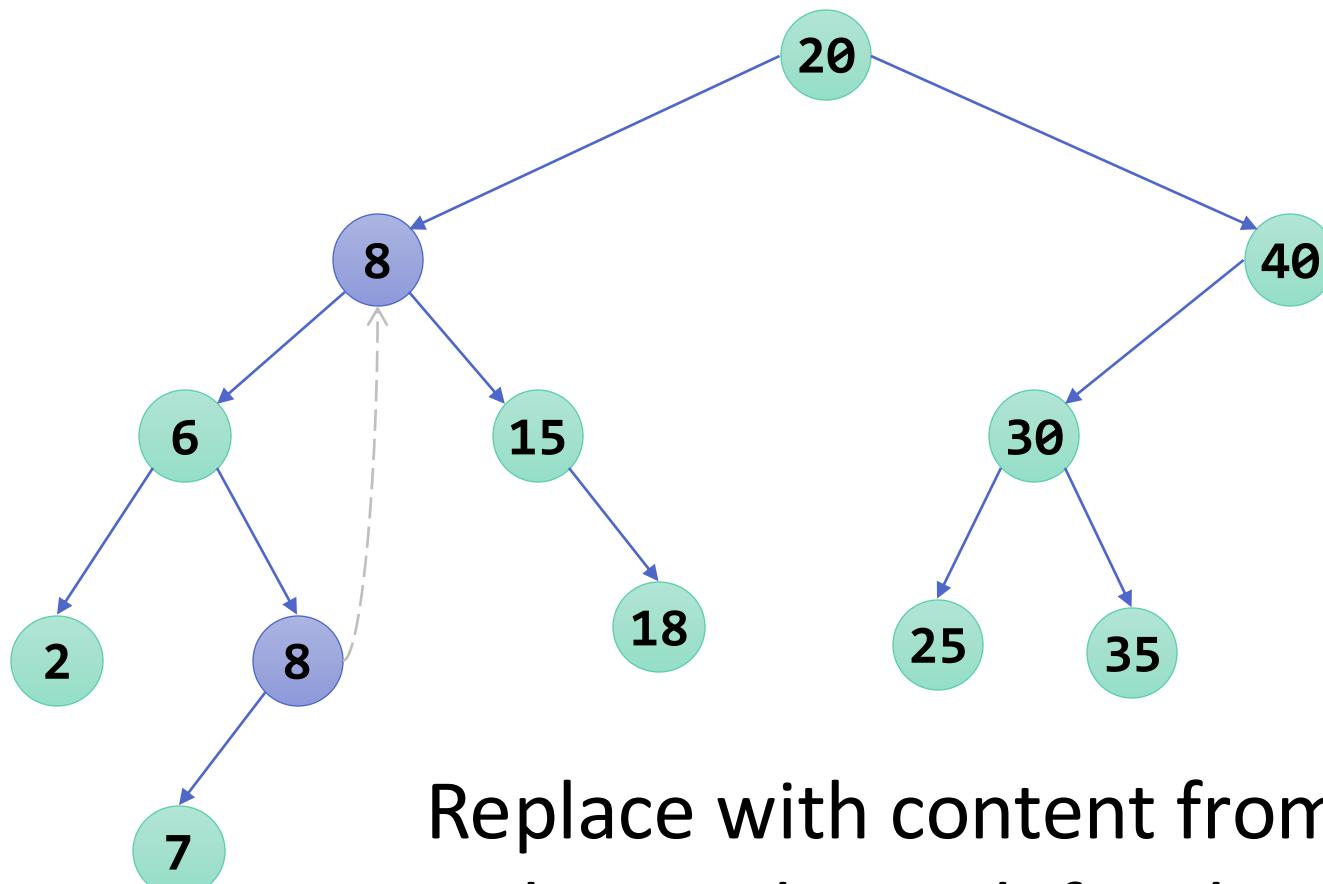


# Delete From Degree 2 Node



Replace with largest key in left subtree  
(or smallest in right subtree).

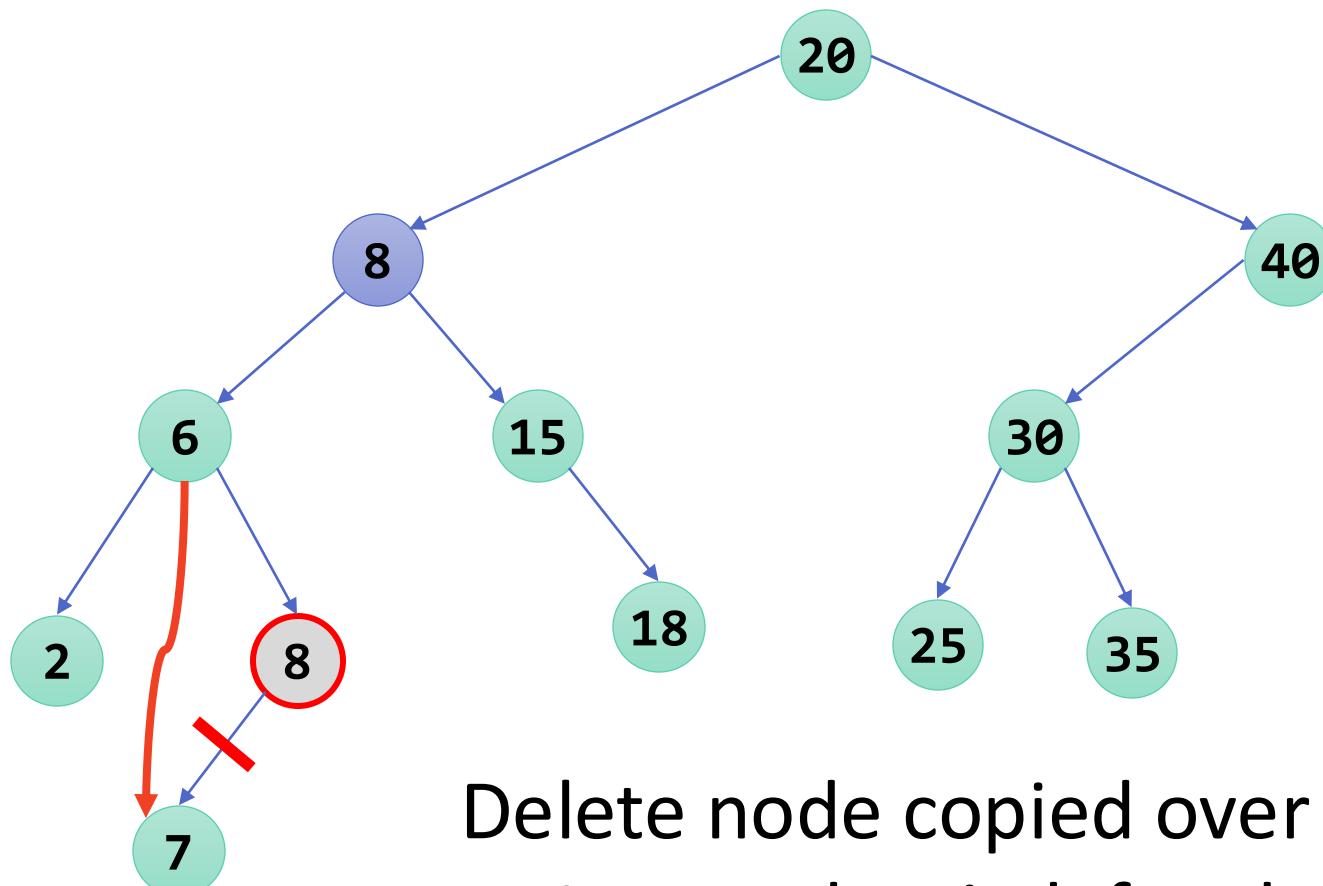
# Delete From Degree 2 Node



Replace with content from

- largest key in left subtree, or
- smallest in right subtree

# Delete From Degree 2 Node

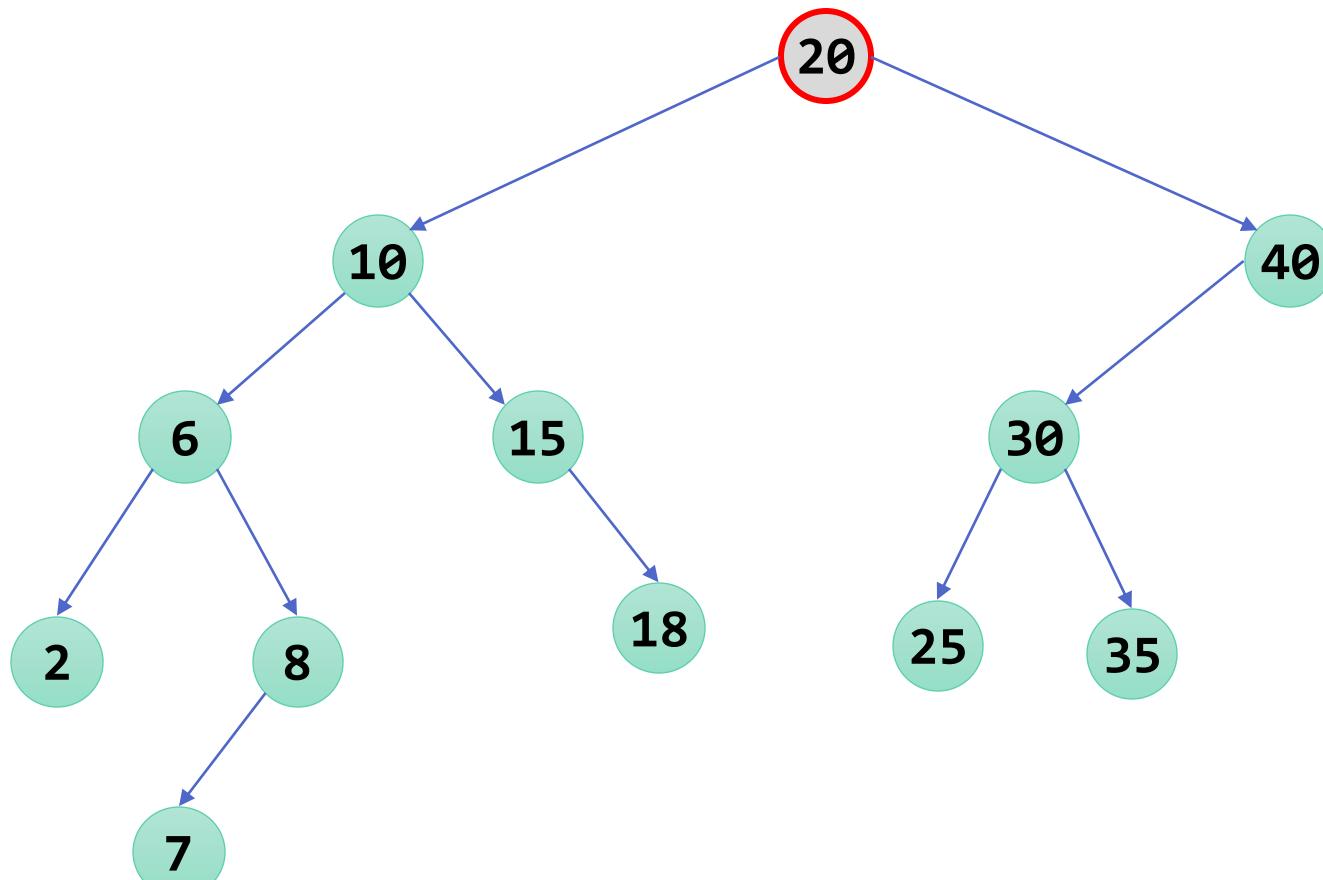


Delete node copied over

- Largest key in left subtree will be a leaf, or degree **1** node.

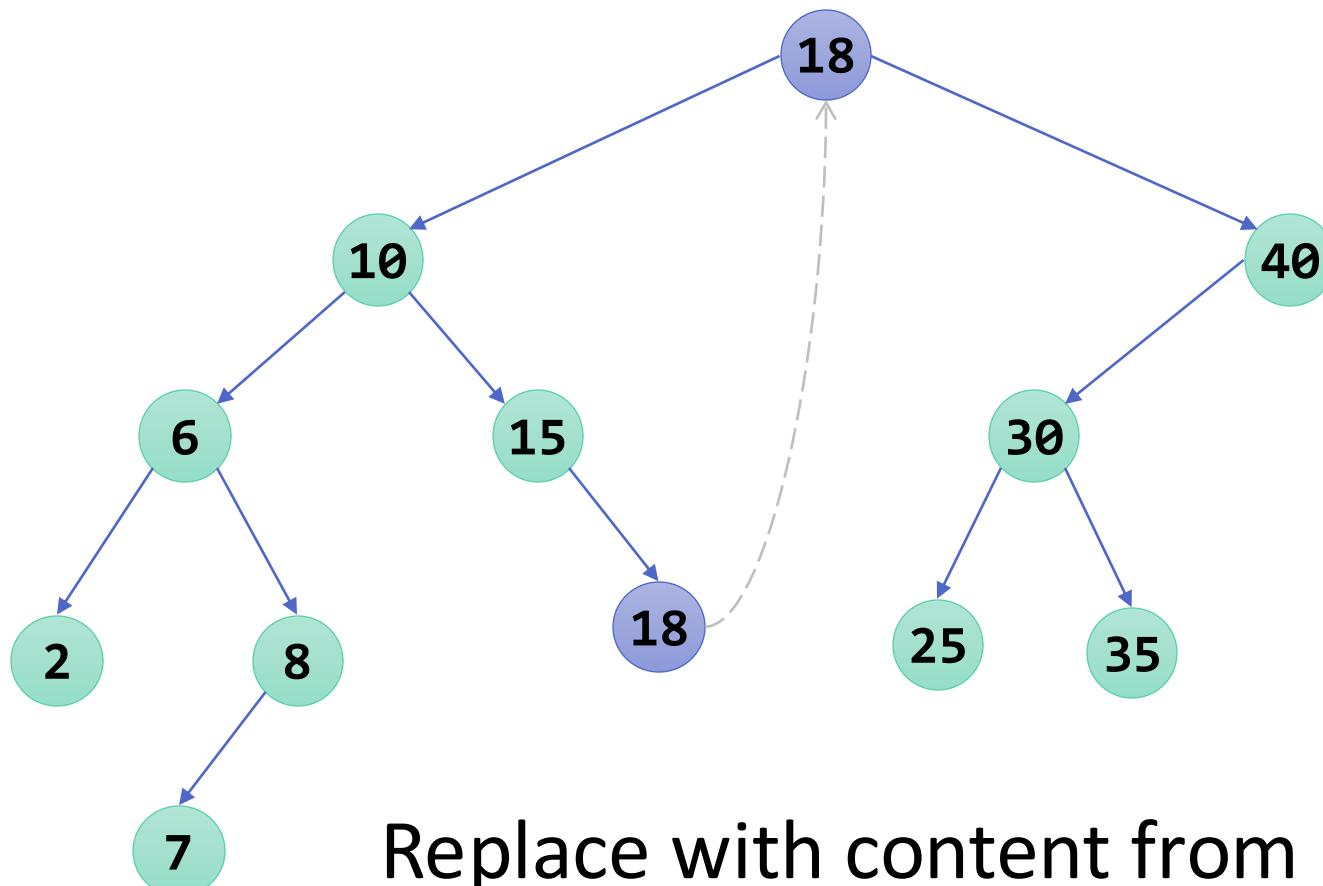


# Delete From Degree 2 Node



Delete from a degree 2 node. key = 20

# Delete From Degree 2 Node

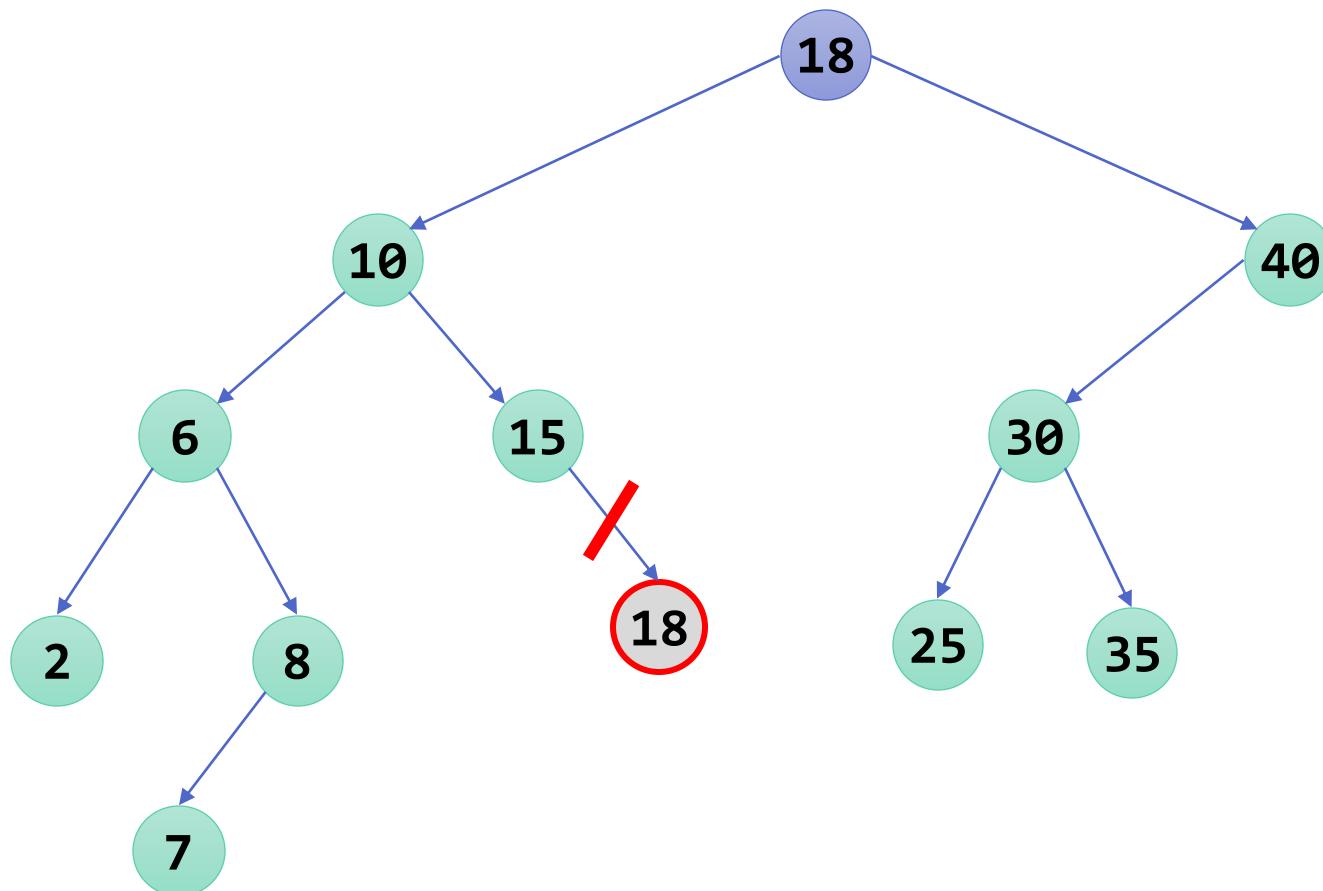


Replace with content from

- largest key in left subtree, or
- smallest in right subtree



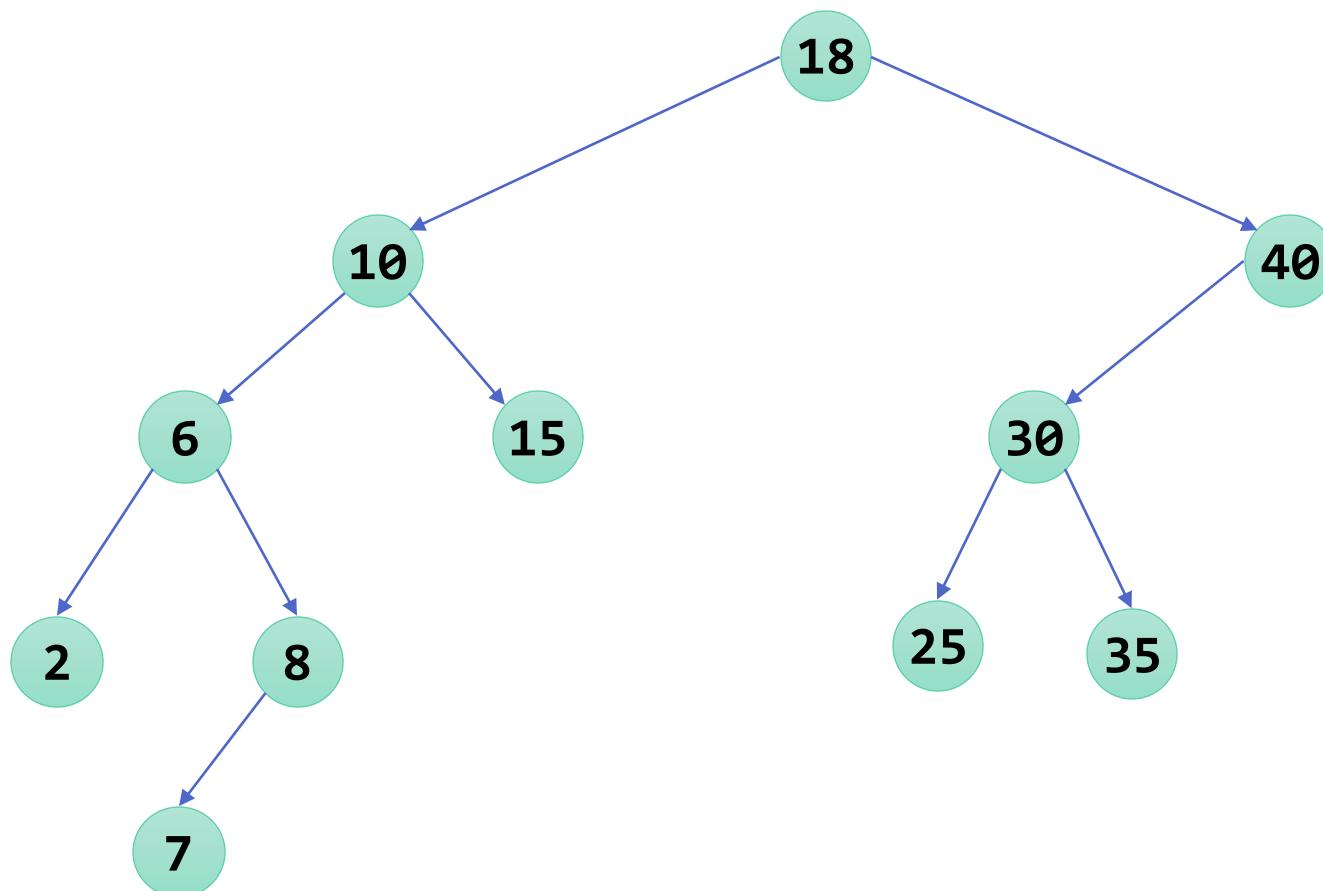
# Delete From Degree 2 Node



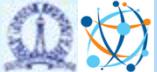
Delete node copied over



# Delete From Degree 2 Node



Complexity is  $O(\text{height})$



# Tree Imbalances

- Similarities between **BST vs. Sorted Array**
- Inserting and Deleting in specific orders can cause tree to be imbalanced
  - E.g. insert in sorted ascending/descending order
  - Height of left and right subtrees are very different, skewed
- Causes complexity to tend to  $O(n)$  rather than  $O(\log(n))$
- Periodically *rebalance* if skew greater than a threshold
  - *Balanced BST*, e.g., AVL Tree, Red-Black Tree, etc.



# Hash Table

- Uses a 1D array (or table) `table[0:b-1]`
  - Each position of this array is a **bucket**
  - Number of buckets is **b**
  - A bucket can normally hold only one dictionary pair: `<key, value>`
    - But larger capacity allowed per bucket as well
    - Bucket sizes can be unbounded as well
- Uses a hash function **h** that converts each key **k** into an index in the range `[0, b-1]`.
  - $h(k)$  is the “**home bucket**” for key **k**.
- Every dictionary pair is stored in its home bucket  
$$\text{table}[h(\text{item.key})] = \text{item}$$



# Ideal Hashing Example

- KVPs are: (22,a), (33,c), (3,d), (73,e), (85,f).
- Hash table is **table[0:7]**, b = 8.
- Hash function **h=key/11**
- Pairs are stored in table as below

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
-------	--	--------	--------	--	--	--------	--------

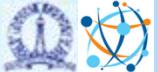
- Lookup, Insert and Delete are done similarly
  - Apply hash, find bucket, perform op.
  - Take **O(1)** time to apply hash and do array access



# What Can Go Wrong?

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
-------	--	--------	--------	--	--	--------	--------

- Where does (99,k) go?
- Hash function causes us to go beyond table size
- **Simple fix:** do a “mod” with the bucket size by default
- $\mathbf{h = (k / 11) \% 8}$



# What Can Go Wrong?

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
-------	--	--------	--------	--	--	--------	--------

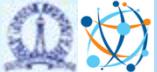
- Where does (26,g) go?
- Keys 22 and 26 have the same *home bucket*, are synonyms with respect to the hash function used
  - This is a **collision**
- The home bucket for (26,g) is already occupied
  - And capacity of bucket is only 1 item
  - This is called an **overflow**



# What Can Go Wrong?

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
-------	--	--------	--------	--	--	--------	--------

- A **collision** occurs when the home bucket for a new pair is occupied by a pair with a different key
- An **overflow** occurs when there is no space in the home bucket for a new pair
  - E.g. if each bucket has capacity to hold two values for same key, and more than 2 values for the key are inserted
- If a bucket has a capacity of 1, *collisions and overflows occur together*
  - *Can we allow buckets to hold multiple item? Unbounded items?*
  - Using a linked list for each bucket item is called “Chaining”
- Need a method to handle overflows



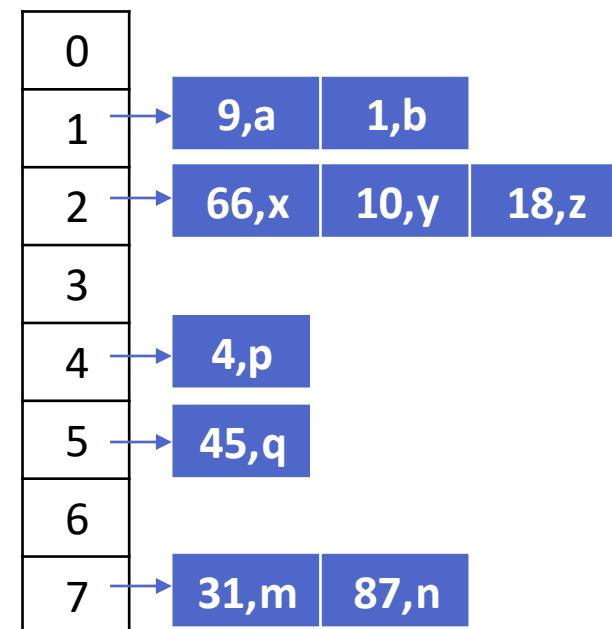
# Designing>Selecting a Hash Table

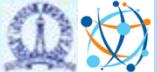
- Choice of hash function
  - Should be **fast** to compute
  - Distributes keys **uniformly** throughout the table
  - Each bucket should have an **equal probability** of receiving a key
  - E.g.  $h=k \% b$  is a *uniform hash function* for keys in the range  $[0..r]$   
... *assuming all keys have equal probability of occurrence*
  - Each buckets gets about  **$\text{ceil}(r/b)$**  or  **$\text{floor}(r/b)$**  keys
- Size (number of buckets) of hash table
  - Decides frequency of collision
- Overflow handling method



# Hash Table using Array & Linked List

- Buckets with unbounded capacity
  - Bucket as a linked list
- Hash function gives array index
- Array contains pointer to head of linked list
  - Items are <key,value> pairs
- Traverse list to lookup element
- What if key not present?
- Time complexity for Insert?  
Lookup?





# Open Addressing to handle Overflows

- All elements are stored in the hash table
  - Elements to store  $\leq$  capacity of table
- Each table entry contains either a `<key,value>` element or *null*
- While **inserting** an element **systematically probe** table slots if overflow occurs
- While **searching** for an element **systematically probe** table slots if bucket does not match key



# Open Addressing

- Modify the hash function to take the *probe number i* as second parameter
  - $h: K \times \{\theta, 1, \dots, b-1\} \rightarrow \{\theta, 1, \dots, b-1\}$
- Hash function,  $h$ , also determines the sequence of slots “probed” for a given key
- Probe sequence for a given key  $k$  is the series of buckets  $h(k, \theta), h(k, 1), \dots, h(k, b-1)$ 
  - Use  $h(k, \theta)$  as bucket if no overflow
  - Else probe each bucket from successive hash fns., i.e. a permutation of  $\langle \theta, 1, \dots, b-1 \rangle$



# Linear Probing

- If the current location is occupied, try the next location

**LPInsert(k)**

**If** (table is full) return error

probe = h(k)

**while** (table[probe] is occupied)

    probe = (probe+1) mod b

table[probe]=k



# Linear Probing – Example

- Home bucket  $h(k) = k \bmod 17$
- Insert keys: 6, 12, 34, 29, 28, 11, 23, 7, 0, 33, 30, 45

0	4	8	12	16
34	0			33

0	4	8	12	16
34	0	45		33



# Lookup in Linear Probing

- Search for a key: Go to  $(k \bmod 17)$  and continue looking at successive locations till we find  $k$  or reach empty location.
  - Longer (unsuccessful) lookup time
  - Deletion?

0	4	8	12	16
34	0	45	28	12

The table shows an array of 17 slots. Slots 0, 4, 8, 12, and 16 are explicitly labeled with their indices above them. The array contains the following values: index 0: 34, index 1: 0, index 2: 45, index 3: (empty), index 4: (empty), index 5: 6, index 6: 23, index 7: 7, index 8: (empty), index 9: (empty), index 10: 28, index 11: 12, index 12: 29, index 13: 11, index 14: 30, index 15: 33, index 16: (empty).



# Deletion

- Shift all elements to previous location?
  - Costly
- Instead, place flag at vacated location
  - `neverUsed=false`
  - Lookup continues till `neverUsed=true`
  - Insert puts element in first location with `neverUsed=true`, sets it to `false`
- Too many markers degrade performance
  - Perform Rehashing



# Complexity Of Dictionary Operations

## find(), insert()

- Given  $n$  elements in the dictionary

Data Structure	Worst Case	Expected
Hash Table (linear probing)	$O(n)$	$O(1)$ *
Binary Search Tree	$O(n)$	$O(\log n)$
<i>Balanced Binary Search Tree</i>	$O(\log n)$	$O(\log n)$

\* Assumptions: (i) Each key's hash is uniform and independent over the ' $b$ ' buckets and (ii) the *load factor* ( $= n/b$ ) is a constant strictly less than one

For a proof, see [https://en.wikipedia.org/wiki/Linear\\_probing](https://en.wikipedia.org/wiki/Linear_probing)



# Demo

- Demo: Chaining vs Linear probing vs BST

[https://github.com/cjain7/DS221-Chirag-LiveDemos/tree/main/Lecture\\_5](https://github.com/cjain7/DS221-Chirag-LiveDemos/tree/main/Lecture_5)



# B Tree

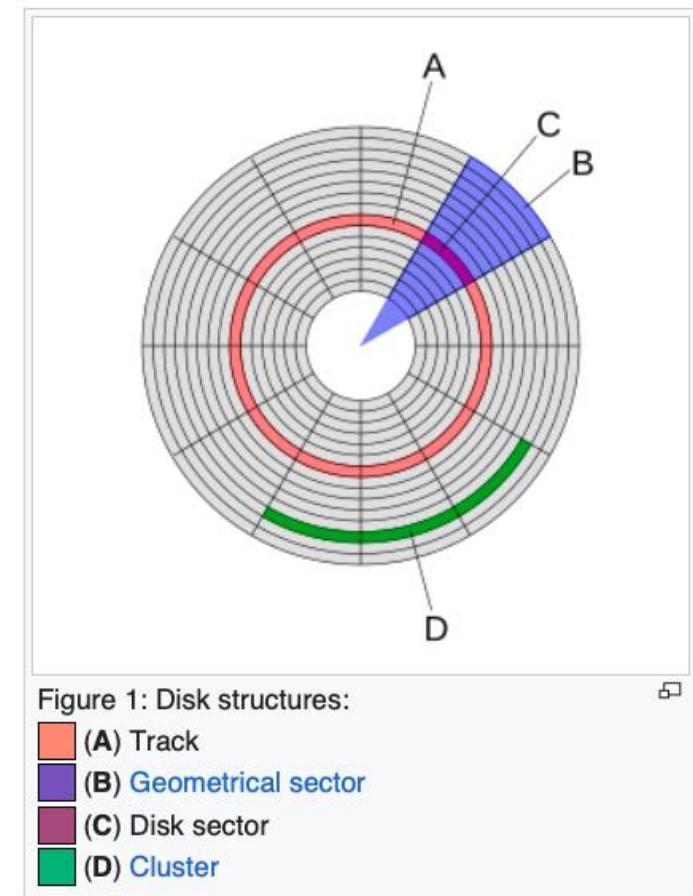


# B-Tree: Searching External Storage

- Main memory (RAM) is fast, but has limited capacity
- Different considerations for in-memory vs. on-disk data structures for search
- Problem: Database too big to fit memory
  - Disk reads are slow
- Example: 1,000,000 records on disk
- Binary search might take 20 disk reads
  - $\log_2(1M) \approx 20$

# Searching External Storage

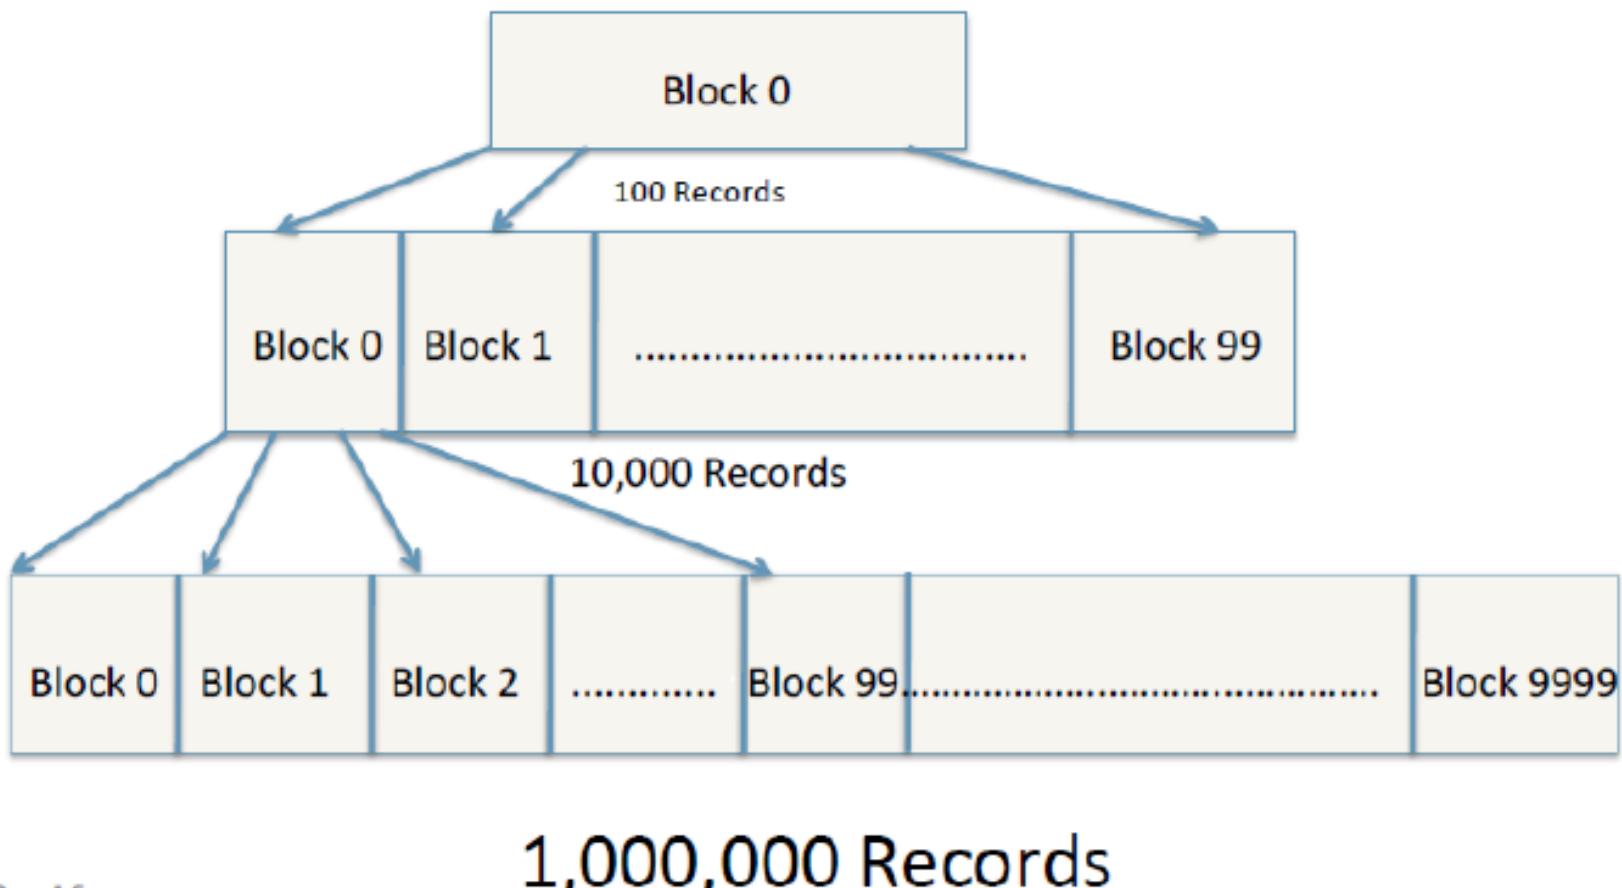
- But disks are accessed “block at a time” by OS
- Blocks are typically 1KiB–4KiB in size
  - Can span multiple “sectors” on HDD
  - Access time per block
  - 10-15 ms for HDD
  - <1ms for SSD
- Say 1KiB block, and say each block can hold 100 records
  - 10,000 blocks for 1M records



Source: Wikipedia



# Searching External Storage





# B-Trees

- Data structures for external memory, not main memory
  - Goal is to reduce number of block accesses, not number of comparisons
- Similar to *binary* search tree
  - But allow more than 1 value and 2 children per node
  - Each node is one disk block with data records plus block addresses of children
- B-Trees
  - Proposed by R. Bayer and E. M. McCreigh in 1972.
  - “Bayer”, “Balanced”, Bushy”, “Boeing” trees?
  - Different from **binary** trees
- NOTE
  - An in-memory data structure generally outperforms an on-disk one. Thus, an in-memory binary tree will be faster than an on-disk B-tree.
  - However, when working on disk, a B-tree is more efficient than a binary tree.



# B-Trees

- Like a BST, nodes contain alternating records (keys & values) and child pointers (blocks).
  - A node with  $k$  records has  $(k + 1)$  children.
- Key ordering rule:
  - Keys in a node are sorted in increasing order.
  - Every key is greater than all keys in its left child's subtree and smaller than all keys in its right child's subtree.
- Bounds on minimum and maximum number of children in a node. For an 'order m' tree:
  - Each node can have **at most** 'm' children
  - Each internal node (except root) must have **at least**  $[m/2]$  children

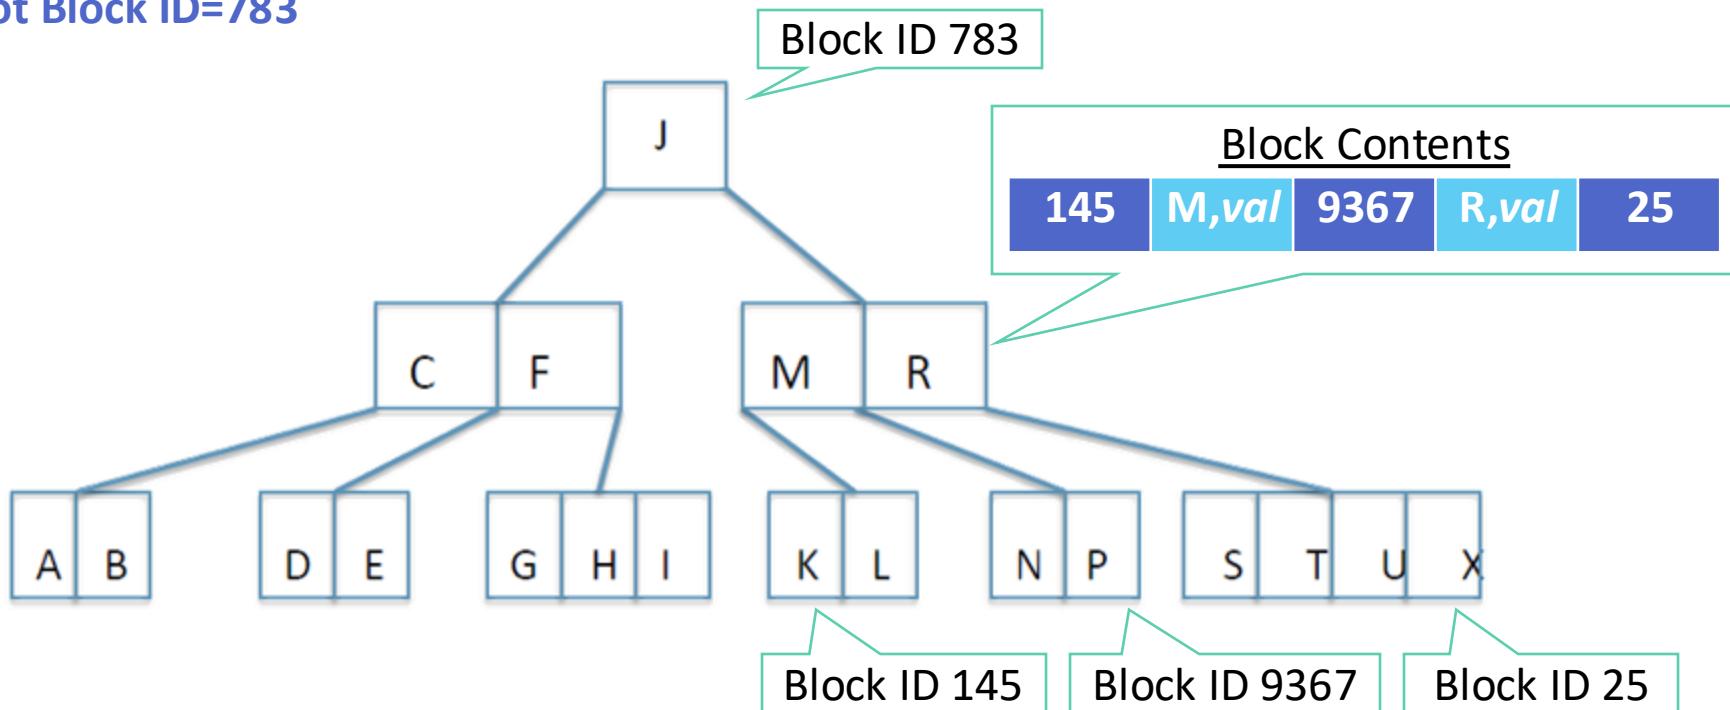
*E.g. Largest-sized node in order 5 B-Tree*



# B-Tree Search (Order 5)

A G F B K D H M J E S I R X C L N T U P

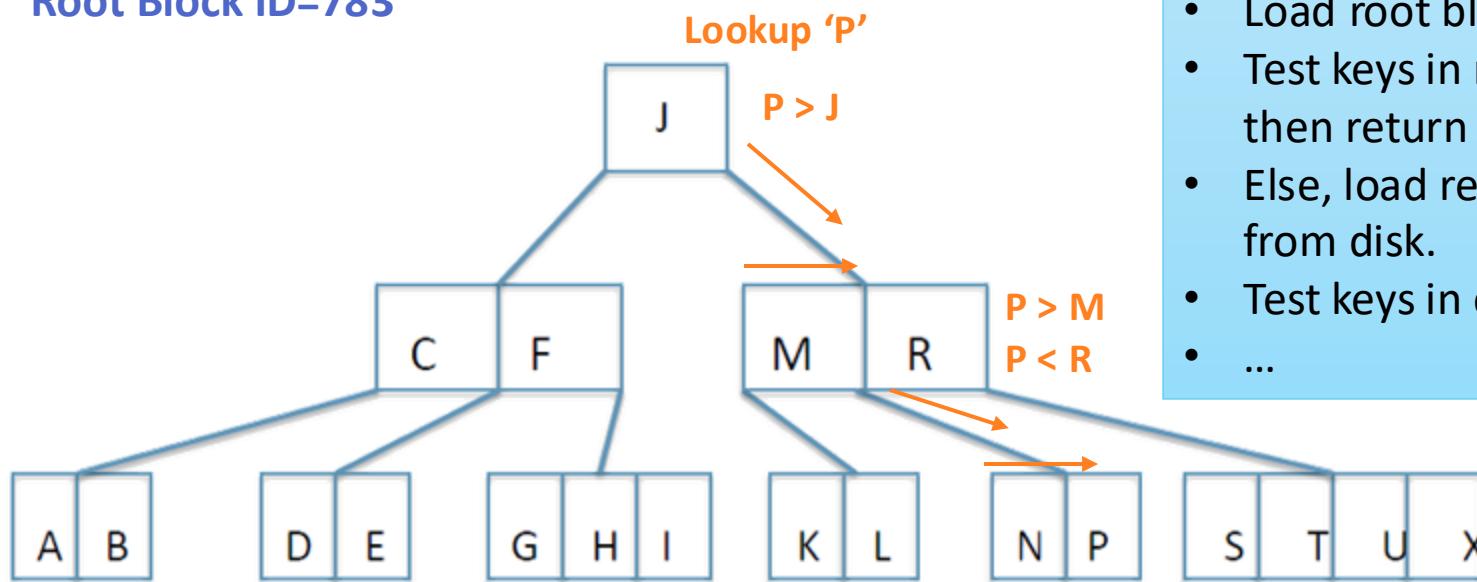
Root Block ID=783



# B-Tree Search (Order 5)

A G F B K D H M J E S I R X C L N T U P

Root Block ID=783



**Lookup is similar to BST**

- Load root block from disk
- Test keys in root block. If match, then return record.
- Else, load relevant child block from disk.
- Test keys in child block...
- ...



# B-Tree Creation

A G F B K D H M J E S I R X C L N T U P

A	B	F	G	K
---	---	---	---	---

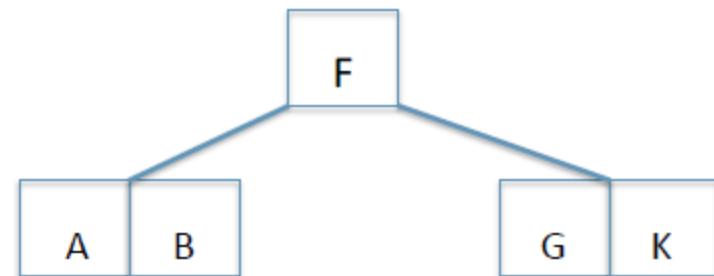
# B-Tree Creation

A G F B K D H M J E S I R X C L N T U P

A	B	F	G	K
---	---	---	---	---



Split if keys > m-1  
Add mid-point to parent.  
Create parent if root.

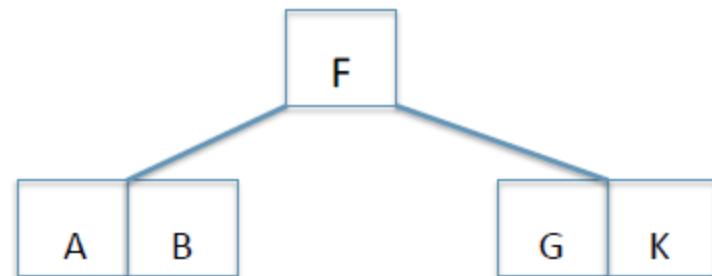




# B-Tree Creation

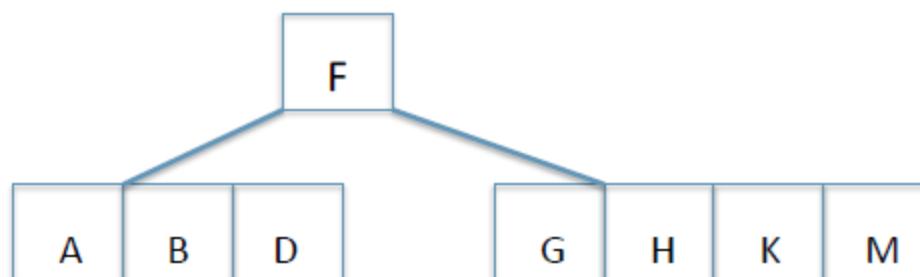
A G F B K D H M J E S I R X C L N T U P

A	B	F	G	K
---	---	---	---	---



Split if keys >  $m-1$   
Add mid-point to parent.  
Create parent if root.

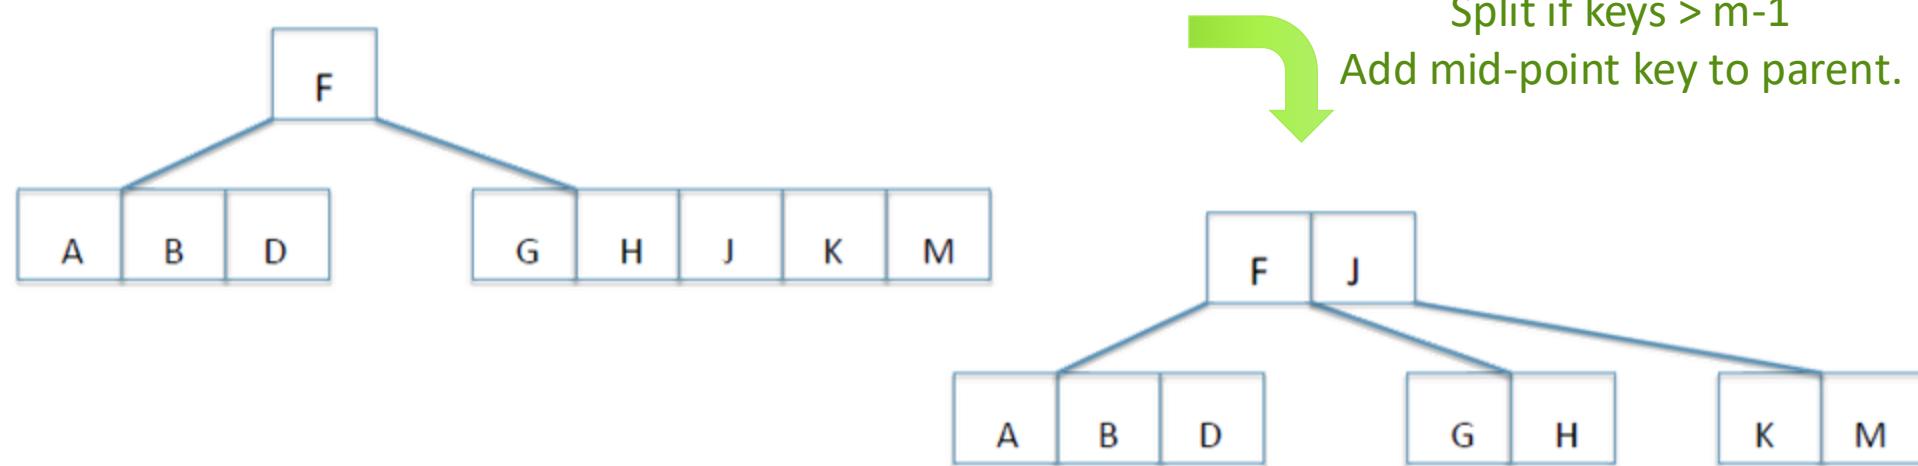
A G F B K D H M J E S I R X C L N T U P



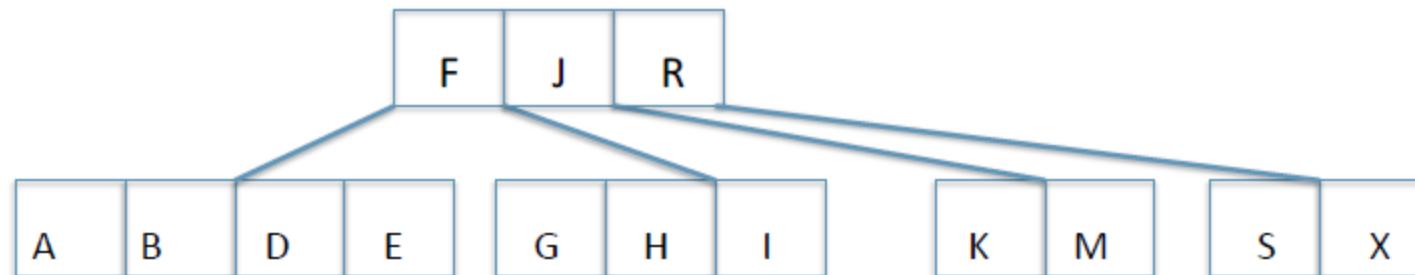


# B-Tree Creation

A G F B K D H M J E S I R X C L N T U P



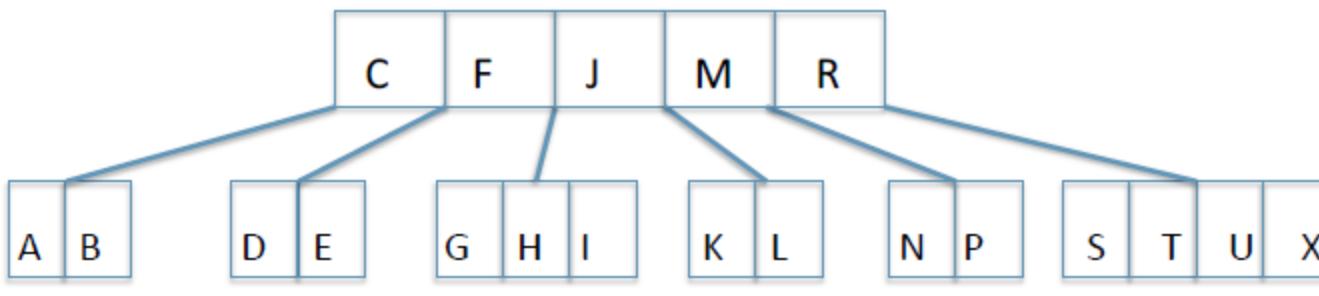
A G F B K D H M J E S I R X C L N T U P



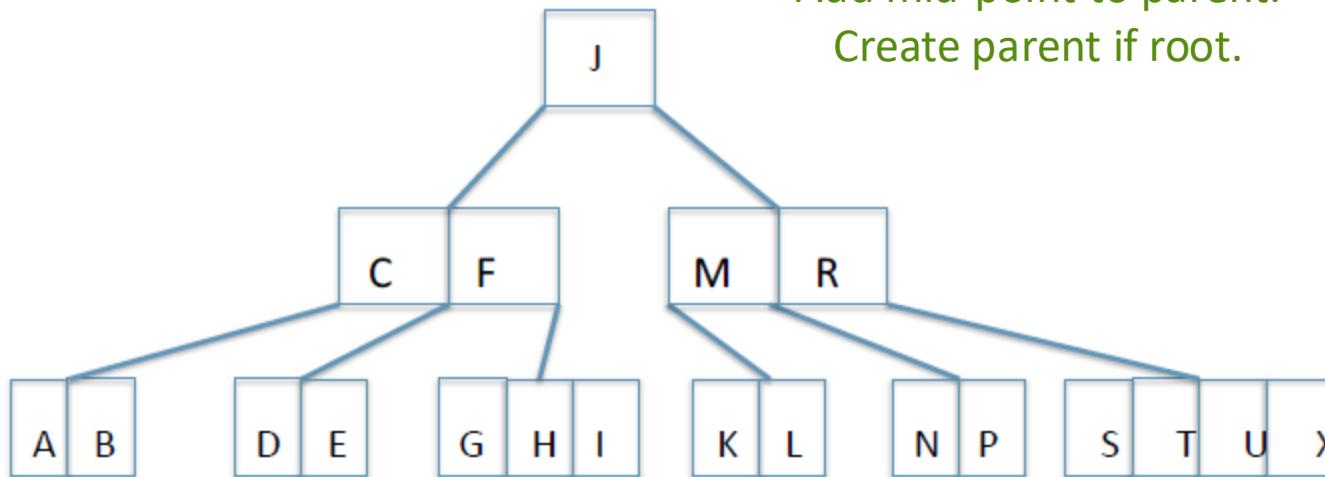


# B-Tree Creation

A G F B K D H M J E S I R X C L N T U P



Split if keys >  $m-1$   
Add mid-point to parent.  
Create parent if root.





# Efficiency of B-trees

- If a B-tree has order  $m$ , then each node (apart from the root) has at least  **$m/2$  children**
- So the depth of the tree is at most  **$\log_{m/2}(\text{size})+1$** 
  - ▶ These many blocks have to be loaded from disk
- In the worst case, we have to make  **$m-1$  comparisons** in each node
  - ▶ Linear search, but  $(m-1)$  is a *constant factor* and *in-memory scan cost* is lower



# Tasks

- Self study (Sahni Textbook)

- Chapter 10.5, Hashing from textbook
- Chapter 11.0-11.6, Trees & Binary Trees from textbook
- B Trees (online sources)