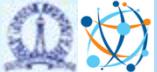




DS221: Introduction to Scalable Systems

Topic: Algorithms and Data Structures



L5: Graphs (continued)

Graph ADT, Algorithms

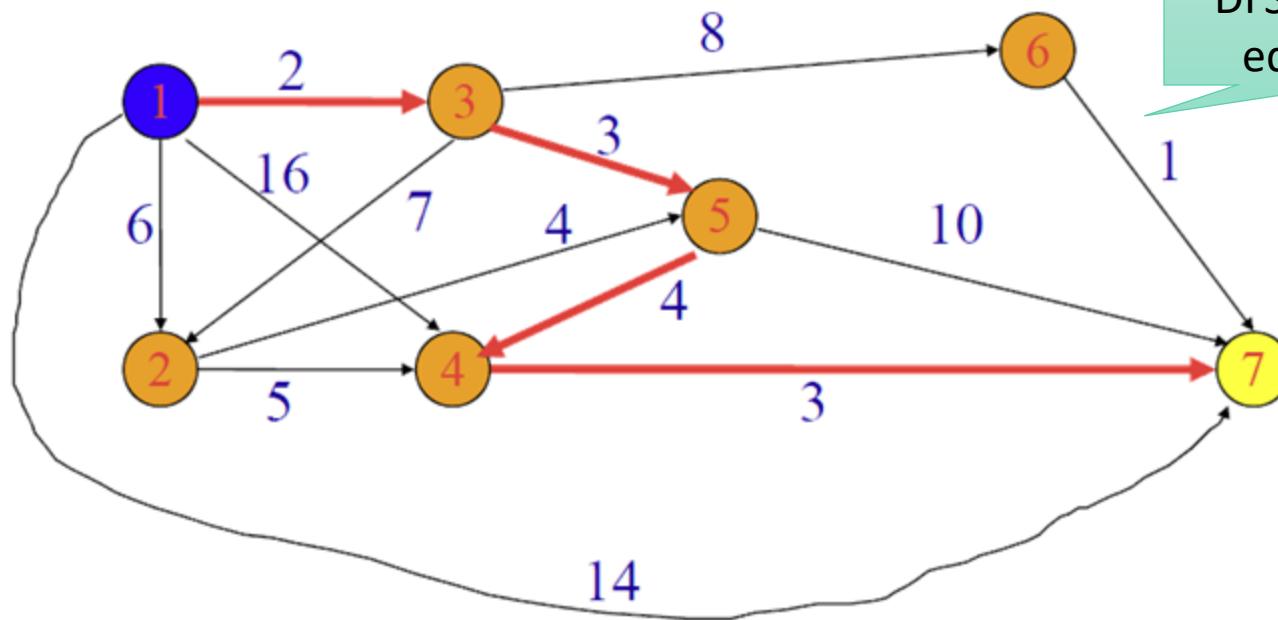
Shortest Path: Single source, single destination

■ Possible greedy algorithm

- Leave source vertex using *shortest outgoing edge*
- Leave new vertex again using shortest outgoing edge to an *unvisited vertex*
- Continue until destination is reached

This is same as the DFS with shortest edge selected!

Greedy Path
from 1 To 7



Length of 12 is not shortest path!

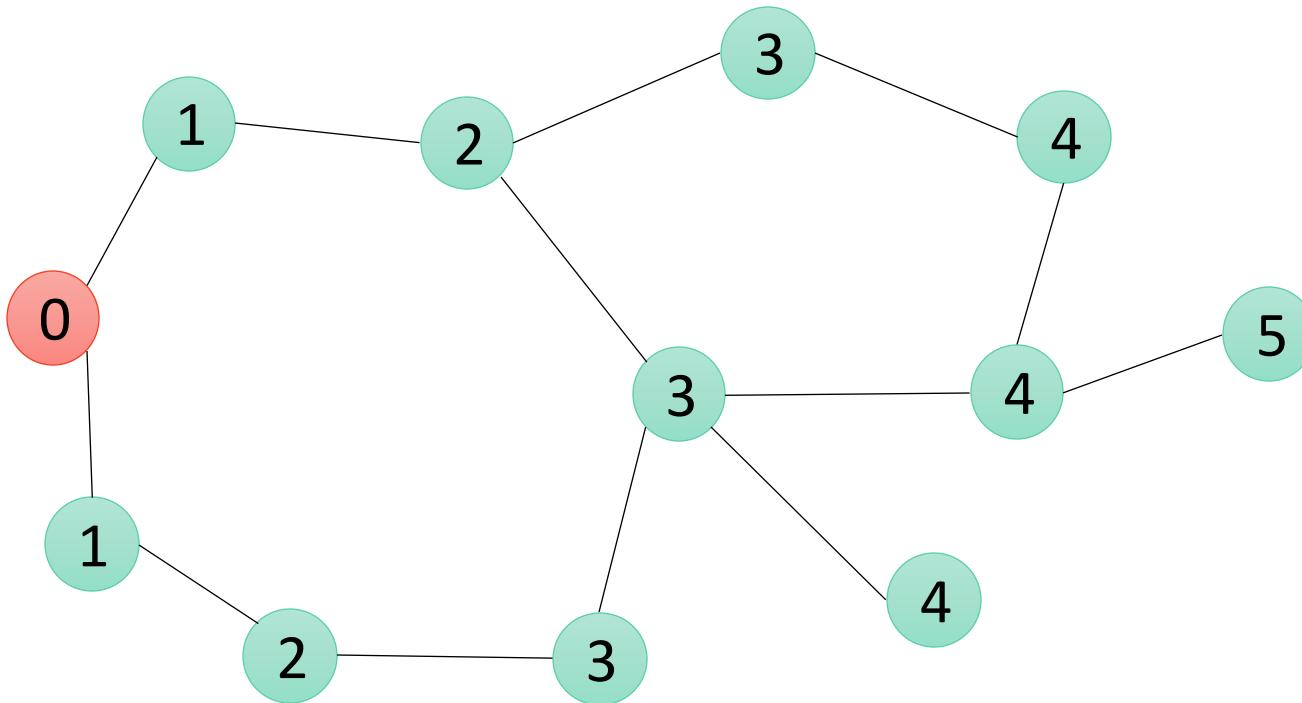


Single Source Shortest Path

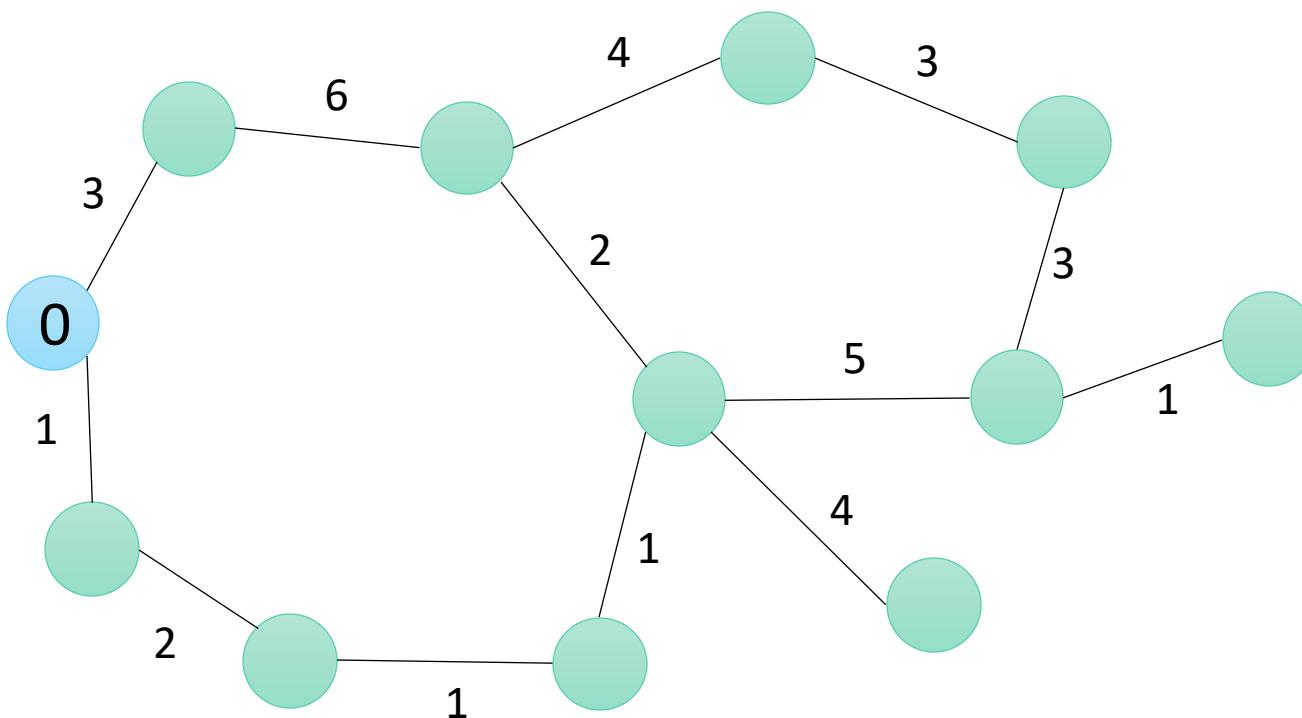
- Shortest distance from one source vertex to all destination vertices
- Is there a simple way to solve this?
- ...Say if you had an unit-weighted graph?
- Just do Breadth First Search (BFS)! 😊



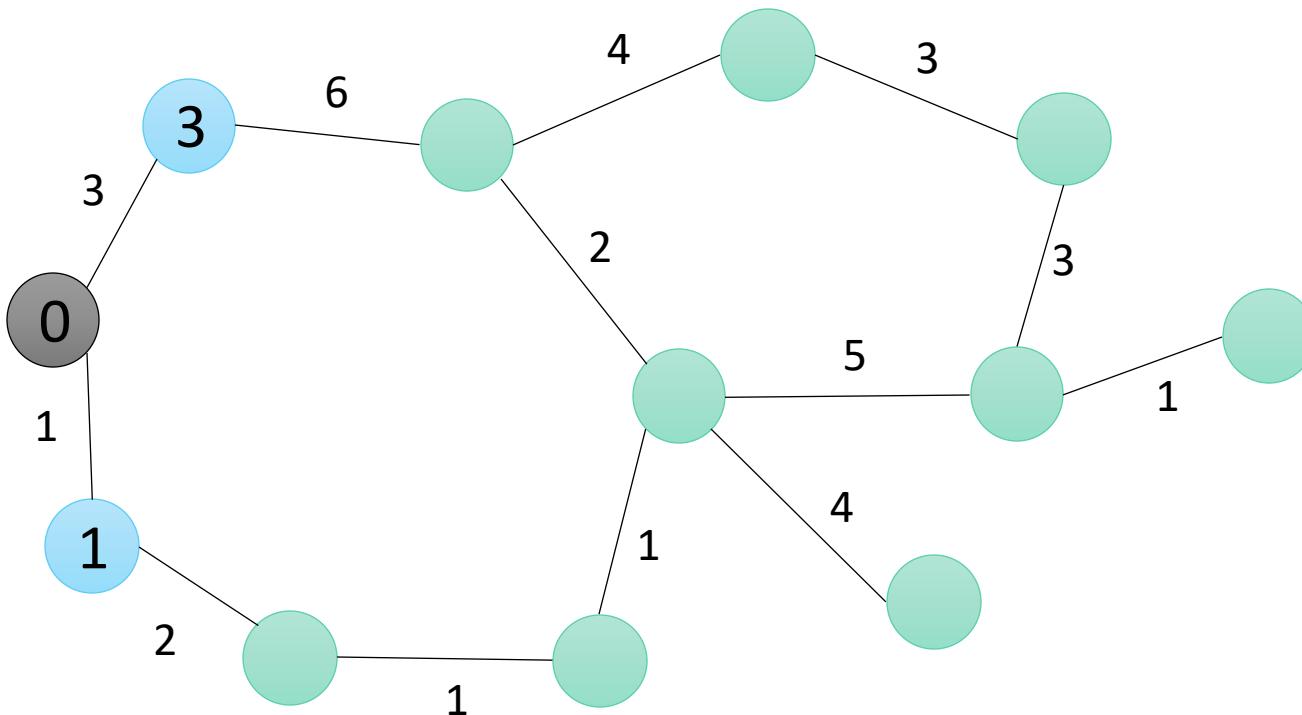
SSSP: BFS on Unweighted Graphs



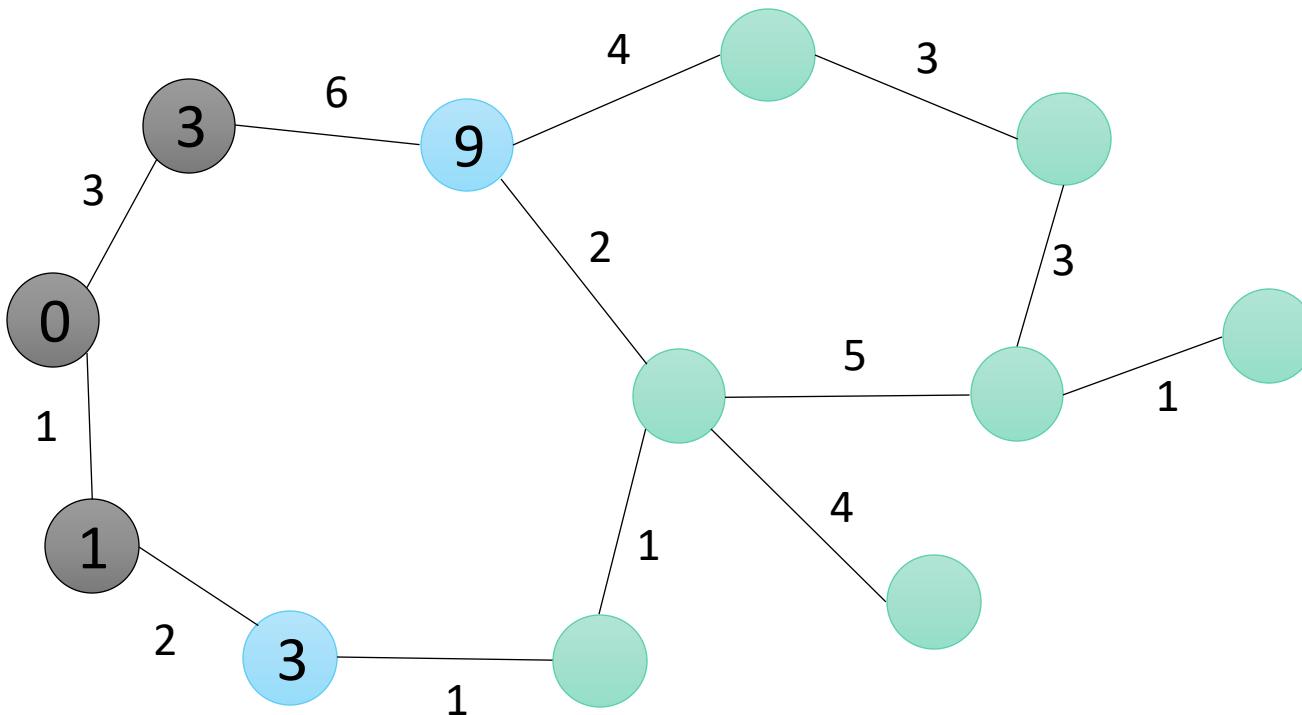
SSSP: *BFS on Weighted Graphs?*



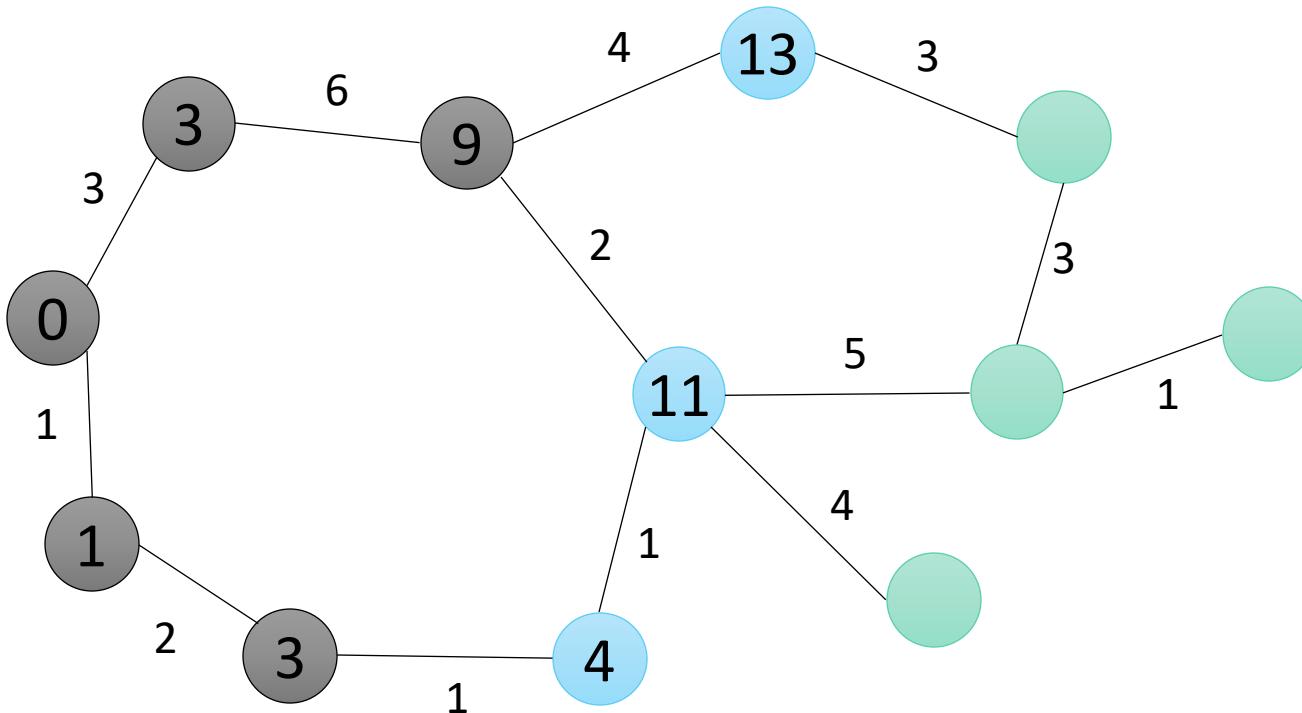
SSSP: BFS on Weighted Graphs?



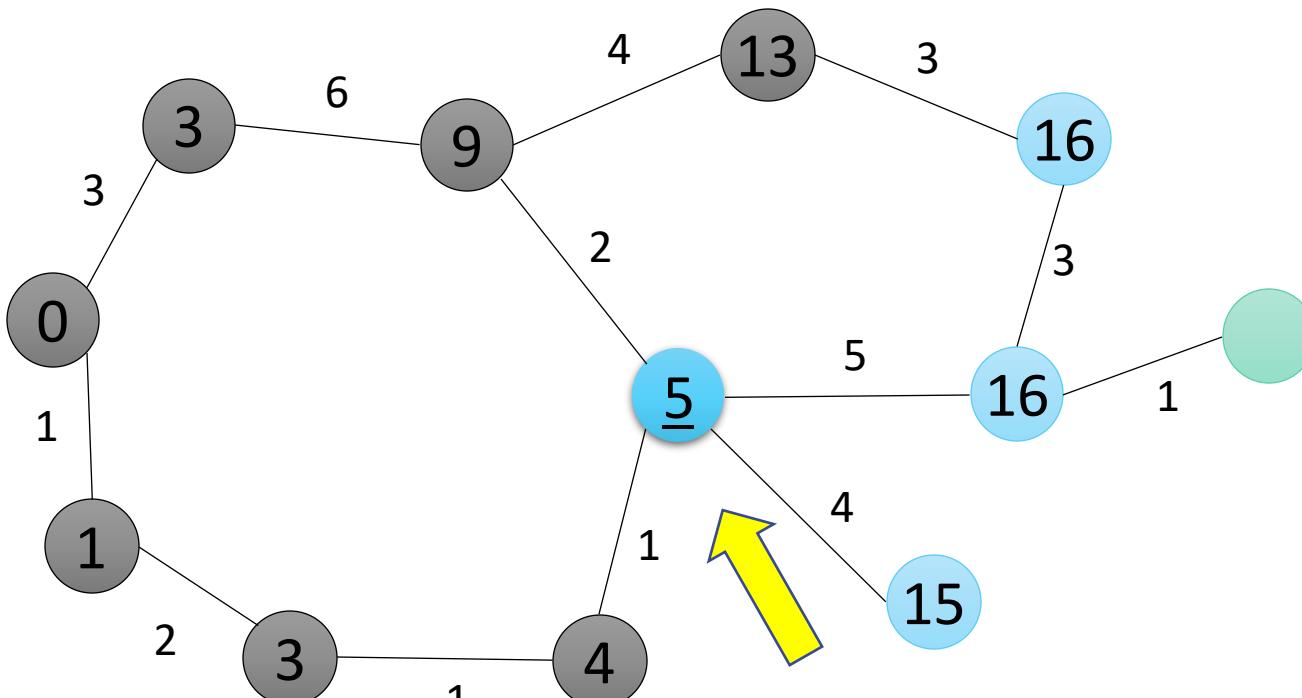
SSSP: BFS on Weighted Graphs?



SSSP: BFS on Weighted Graphs?

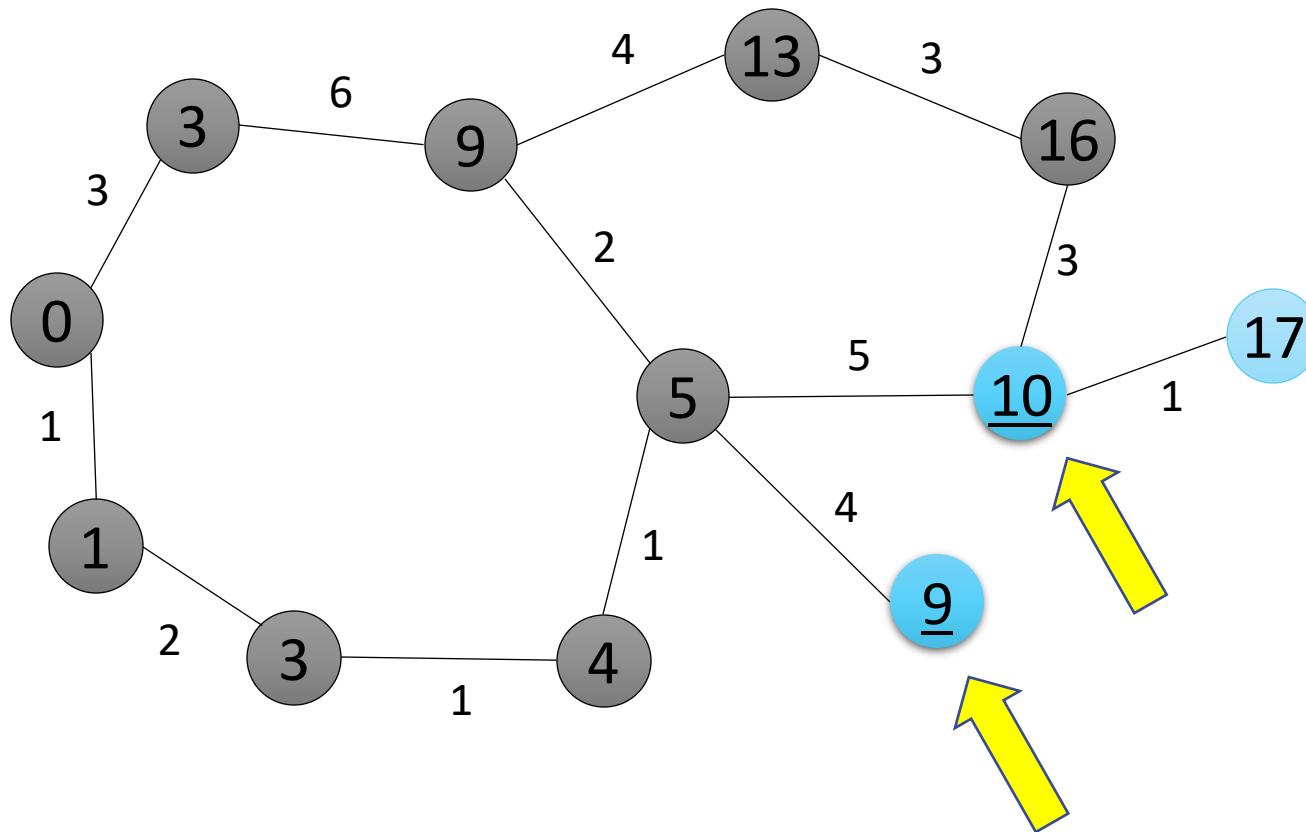


SSSP: BFS on Weighted Graphs?

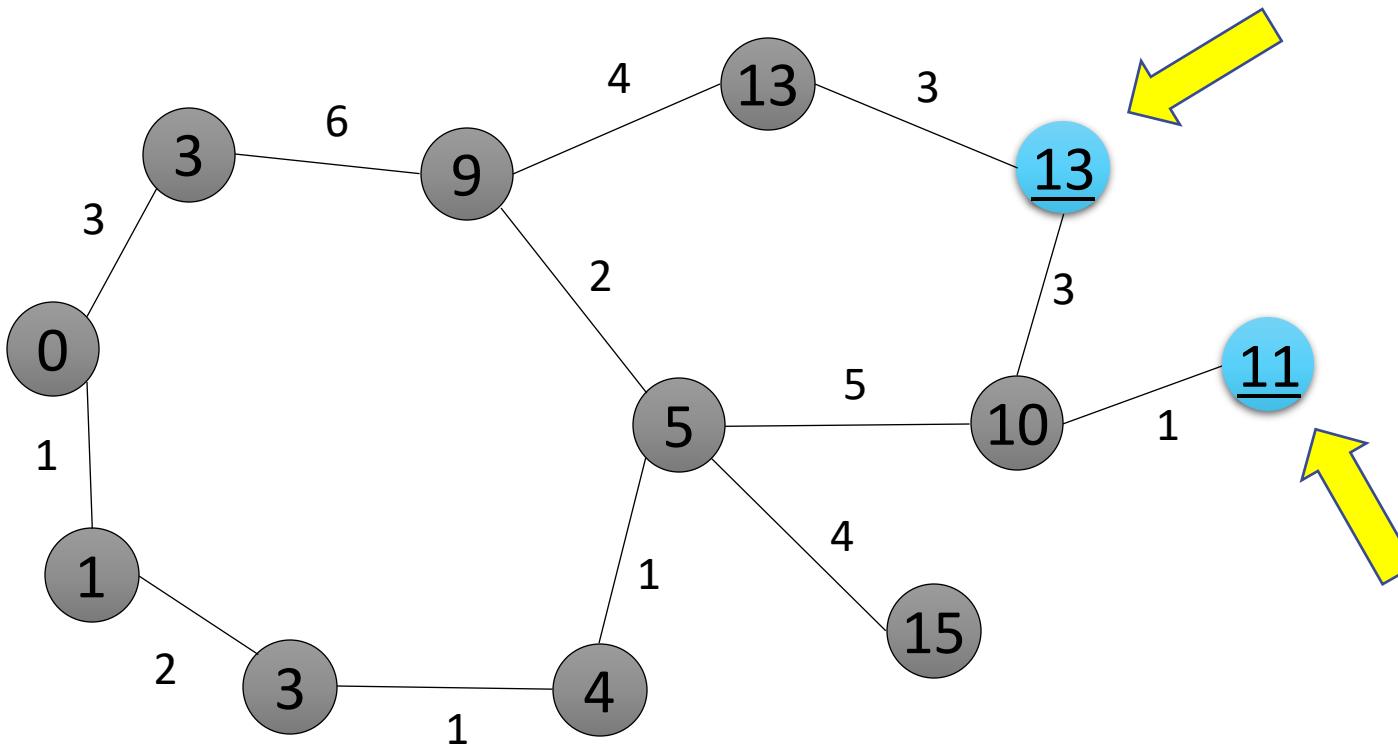


*Revisit, recalculate, re-propagate...
cascading effect*

SSSP: BFS on Weighted Graphs?



SSSP: BFS on Weighted Graphs?

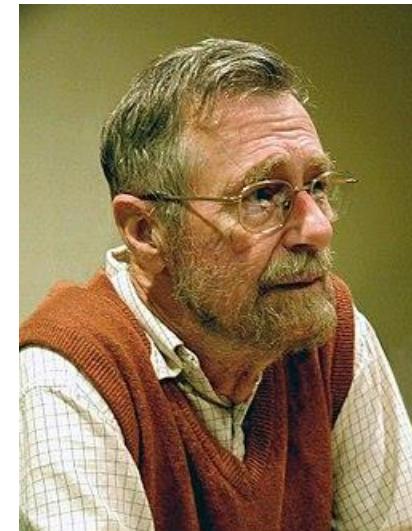


BFS with revisits is not efficient. Can we be smart about order of visits?



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- Named after its inventor Edsger Dijkstra (1930-2002)
- One of the “founders” of computer science; this algorithm is one of his many contributions





Invariant of an algorithm

- An **invariant of an algorithm** is a property (a condition or statement) that:
 - Holds true before the algorithm starts,
 - Remains true throughout every step
 - Is still true when the algorithm finishes
- Invariants are often used to **prove correctness** of algorithms.
- Example: In the binary search algorithm over a sorted array, an invariant is:

“The target value, if it exists, is always within the current search interval.”



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- Solves SSSP problem on a weighted, directed graph $G = (V, E)$
- Assumption: All edge weights are nonnegative
- s is the source vertex
- Let $\delta(s,x)$ be the shortest distance from s to x
- Algorithm will maintain the following invariants:
 - ▶ $\forall x \in V$, the algorithm maintains an **estimate** $d[x]$ of the distance of vertex x from s such that:
 - At any point in time, $d[x] \geq \delta(s,x)$
 - When x is “finished”, $d[x] = \delta(s,x)$



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

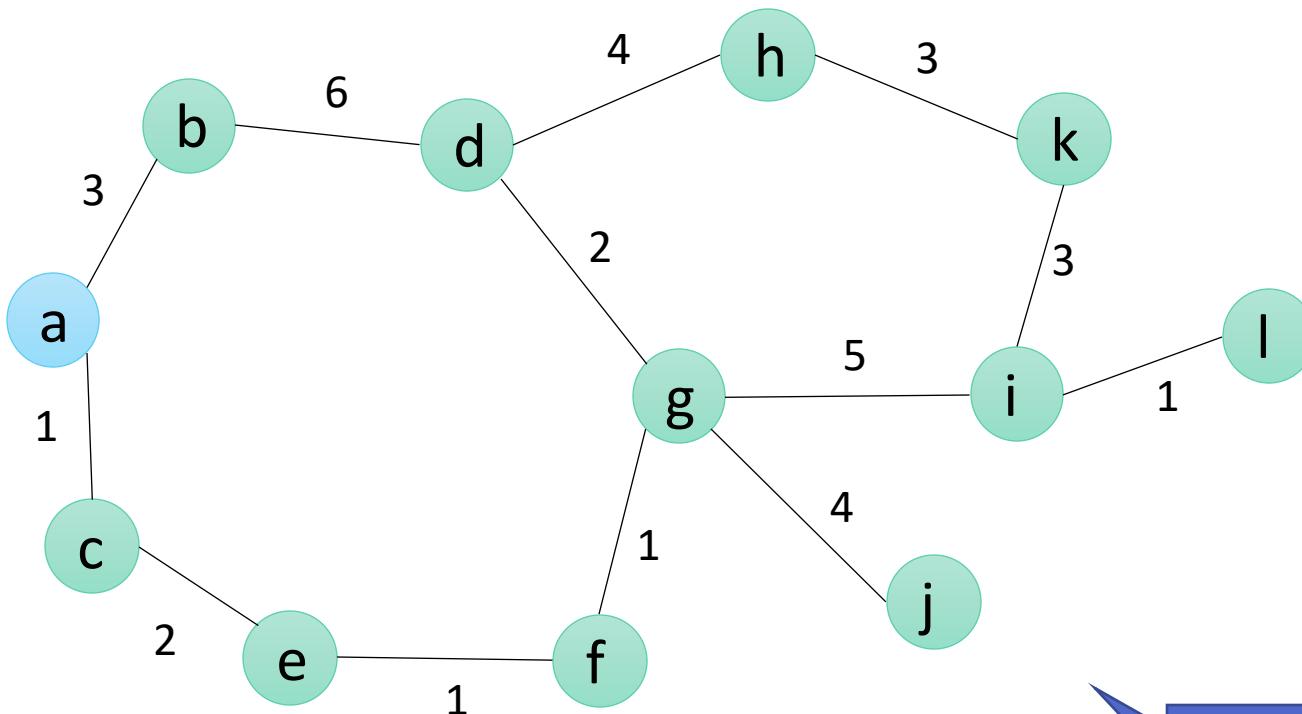
- Let $w(x,y)$ denote the weight of an edge from x to y
- Initialize distance estimates of all vertices to ∞
- Update $d[s]$ to 0
- F is the set of vertices that are yet to achieve final distance estimates. Initialize F to V
- D is the set of vertices that have achieved final distance estimates



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

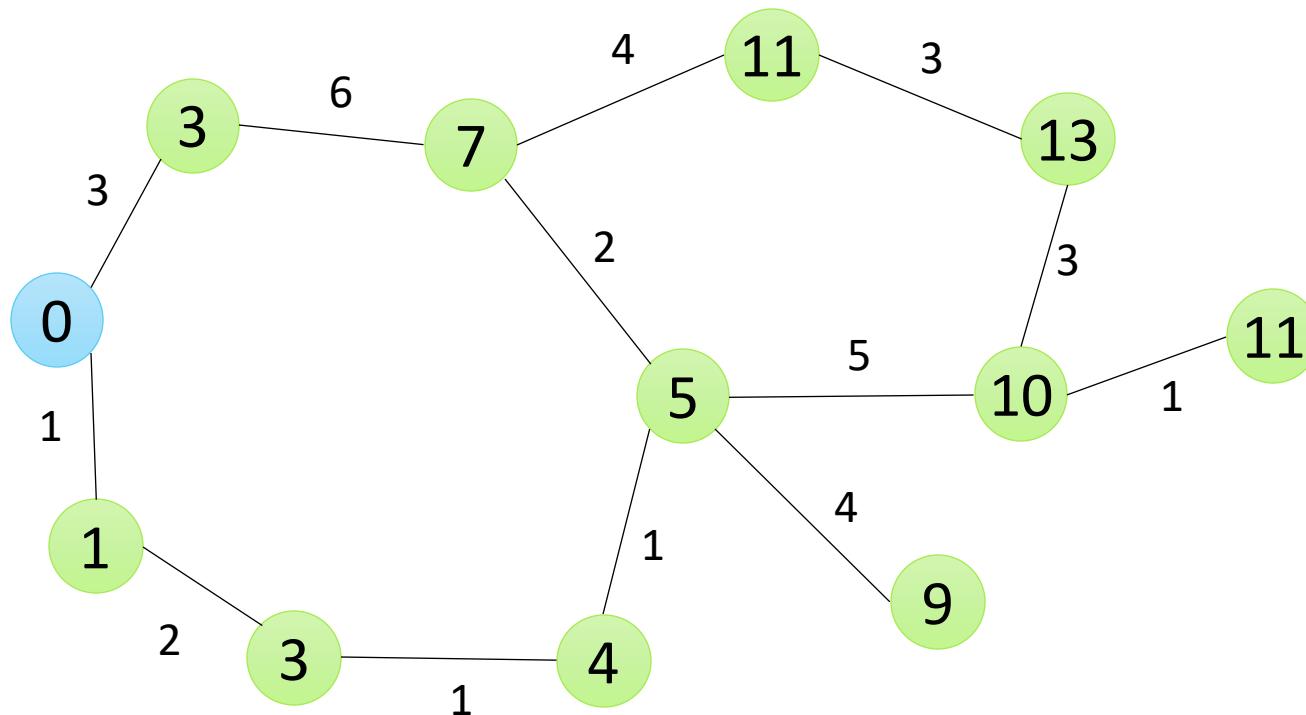
- Let $w(x,y)$ denote the weight of an edge from x to y
- Initialize distance estimates of all vertices to ∞
- Update $d[s]$ to 0
- F is the set of vertices that are yet to achieve final distance estimates. Initialize F to V
- D is the set of vertices that have achieved final distance estimates
- Algorithm:
 - while (F is not empty)
 - Let x be the vertex in F with minimum distance estimate
 - Remove x from F and add it to D
 - For all edges (x,y) from x , update distance estimate of y :
$$d[y] = \min (d[y], d[x] + w(x,y))$$

SSSP on Weighted Graphs



Work out
yourself !

SSSP on Weighted Graphs





Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- Let $w(x,y)$ denote the weight of an edge from x to y
- Initialize distance estimates of all vertices to ∞
- Update $d[s]$ to 0
- F is the set of vertices that are yet to achieve final distance estimates. Initialize F to V
- D is the set of vertices that have achieved final distance estimates
- Algorithm:

while (F is not empty)

- Let x be the vertex in F with minimum distance estimate
- Remove x from F and add it to D
- For all edges (x,y) from x , update distance estimate of y :
$$d[y] = \min(d[y], d[x] + w(x,y))$$

Does invariant $d[u] \geq \delta(s,u)$ hold for all vertices $u \in V$?



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- **Claim 1:** For every vertex u , at any point of time $d[u] \geq \delta(s,u)$
- We will prove that at any point in time, if $d[u] < \infty$, then **$d[u]$ is the weight of some path from s to u**

Proof by induction

- Base case (at initialization): $d[s] = 0 = \delta(s,s)$ and all other estimates are $+\infty$
- Induction hypothesis: Claim holds till vertex x was picked from set F
- Now if y is a out-neighbour of x , we may update $d[y]$ to $d[x]+w(x,y)$
- By the induction hypothesis, there exists a path from s to x with weight $d[x]$. We also have edge (x,y) of weight $w(x,y)$. Clearly, there exists a path from s to y of weight $d[x]+w(x,y)$



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- Let $w(x,y)$ denote the weight of an edge from x to y
- Initialize distance vector $d[]$ for all vertices to **infinity**, except for source which is set to **0**
- F is the set of vertices that are yet to have their distance estimates. Initialize F to V
- D is the set of vertices that have achieved their final distance estimates
- Algorithm:
while (F is not empty)
 - Let x be the vertex in F with minimum distance estimate
 - Remove x from F and add it to D
 - For all edges (x,y) from x , update distance estimate of y :
$$d[y] = \min (d[y], d[x] + w(x,y))$$

Now, how do we show
that $d[x] = \delta(s,x)$?



Dijkstra's Single Source Shortest Path Algorithm (SSSP)

- **Claim 2:** When vertex x is placed in D , $d[x] = \delta(s,x)$
(Proof by contradiction)
- Suppose there exist at least one vertex in V for which this property is violated.
- Assume u is the **first vertex** added to D with $d[u] \neq \delta(s,u)$
- Note that u cannot be s (Why? $d[s] = \delta(s,s) = 0$)
- There must exist a shortest path from s to u in the graph. Call this path P

Dijkstra's Single Source Shortest Path Algorithm (SSSP)

Case 1: All vertices on path P except u are in D

Let x be the last vertex on path P which is in D

When x was placed in D, $d[x] = \delta(s,x)$ and $d[u]$ must have been updated to $\delta(s,u)$

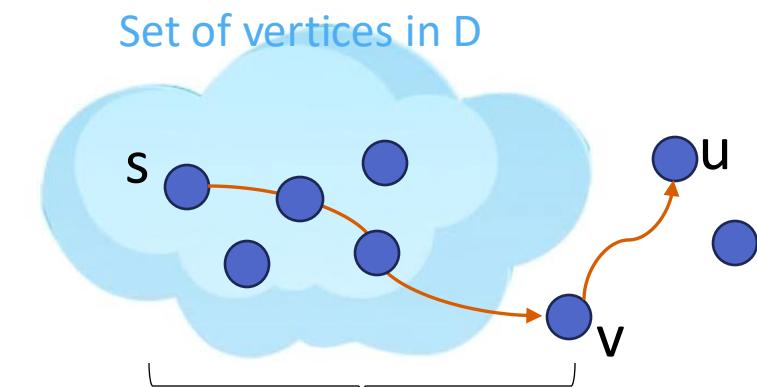
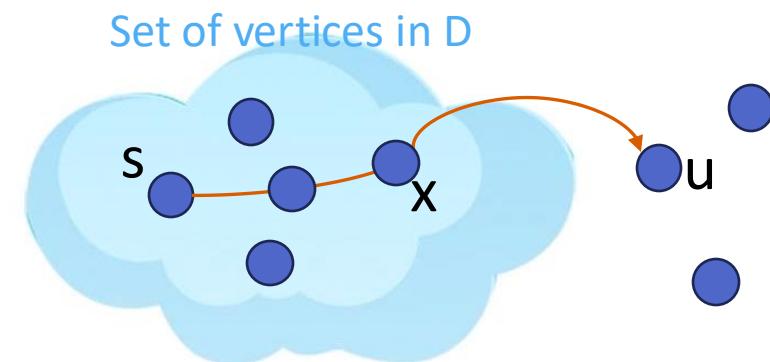
Case 2: Besides u, there is at least one vertex on path P that is not in D

Let v be the first vertex on path P which is not in D

$$d[v] = \delta(s,v) \leq \delta(s,u)$$

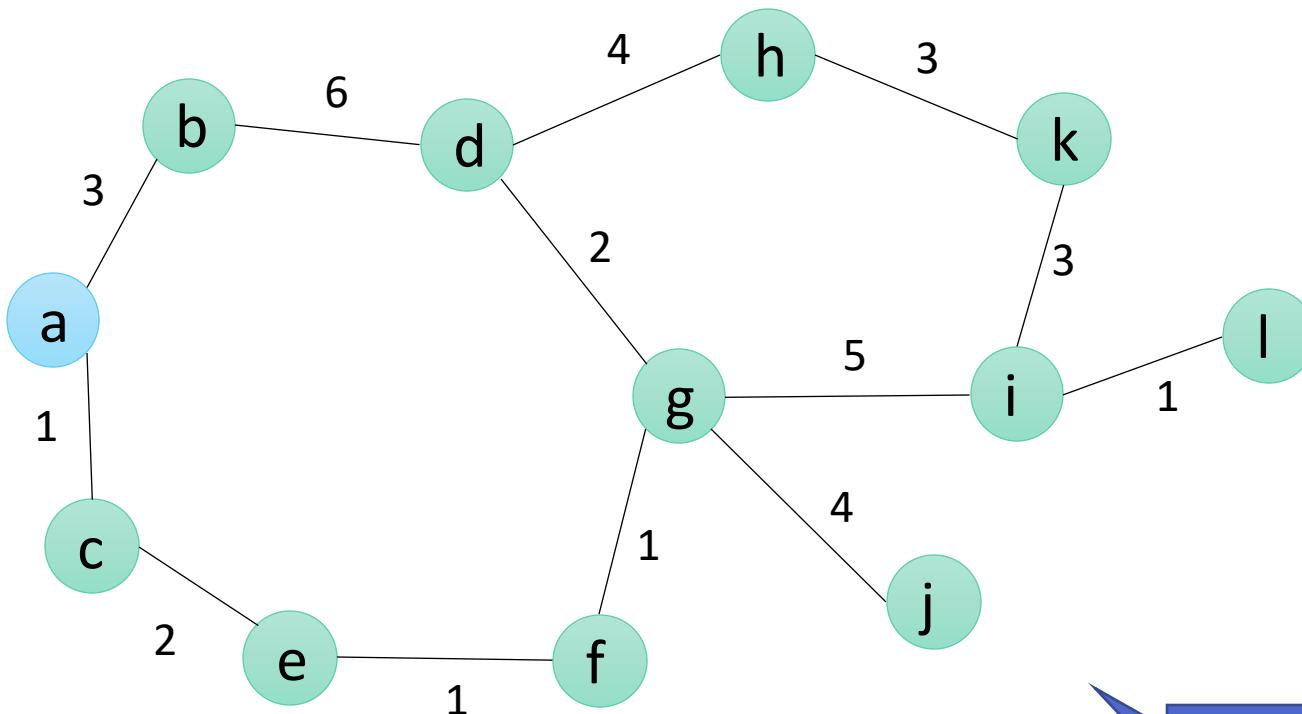
$$d[u] \neq \delta(s,u) \text{ implies } d[u] > \delta(s,u)$$

If u was picked before v, then $d[u] \leq d[v]$ (contradiction!)



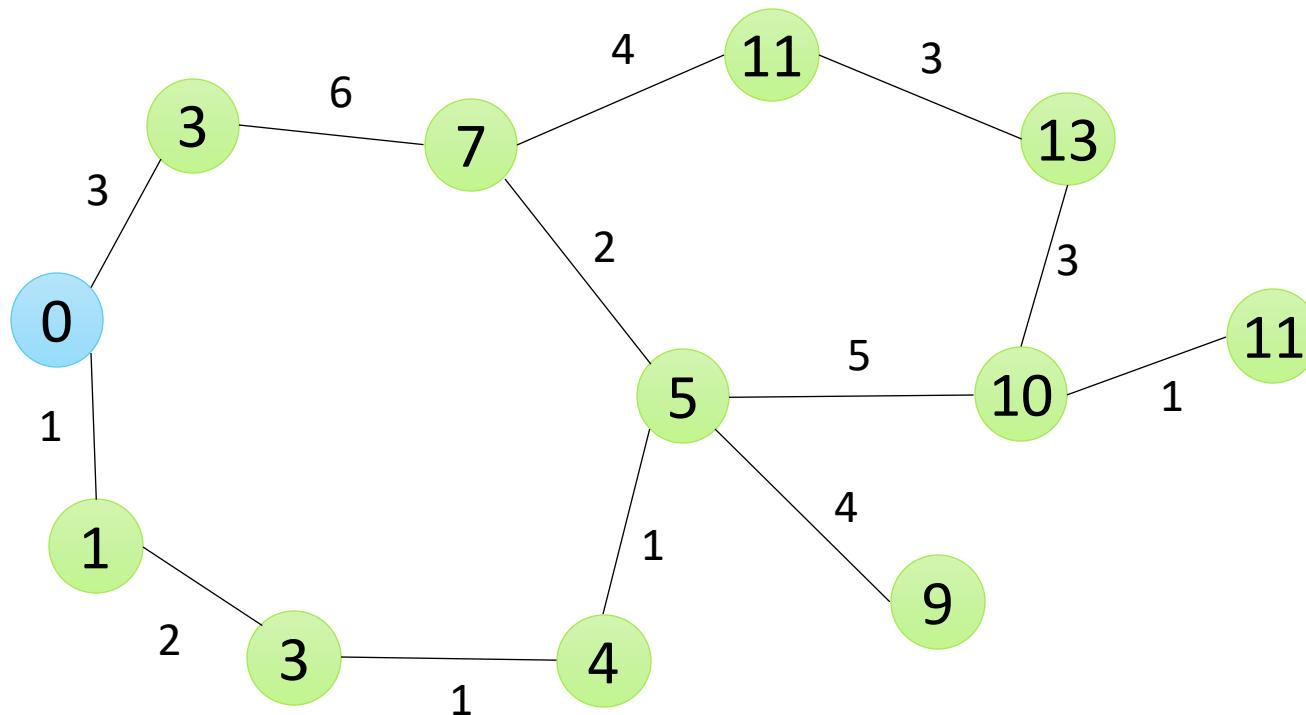
This subpath of P from s to v must be shortest path from s to v (WHY?)

SSSP on Weighted Graphs



Work out
yourself !

SSSP on Weighted Graphs





Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)



Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)

We execute EXTRACT-MIN operation $|V|$ times

We execute DECREASE-KEY operation $|E|$ times



Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)
- Option 1: Assume vertices are numbered from 1 to $|V|$. Use an **array** of size $|V|$, store $d[u]$ at the u^{th} entry, have a separate boolean array of size $|V|$ to indicate if u is in set F or not



Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)
- Option 1: Assume vertices are numbered from 1 to $|V|$. Use an **array** of size $|V|$, store $d[u]$ at the u^{th} entry, have a separate boolean array of size $|V|$ to indicate if u is in set F or not
 - EXTRACT-MIN: we linearly search the array for the smallest: $O(|V|)$
 - DECREASE-KEY: Each operation takes $O(1)$ time
 - Total time: **$O(|V|^2 + |E|)$**



Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)
- Option 1: Assume vertices are numbered from 1 to $|V|$. Use an **array** of size $|V|$, store $d[u]$ at the u^{th} entry, have a separate boolean array of size $|V|$ to indicate if u is in set F or not
 - EXTRACT-MIN: we linearly search the array for the smallest: $O(|V|)$
 - DECREASE-KEY: Each operation takes $O(1)$ time
 - Total time: **$O(|V|^2 + |E|)$**
- Option 2: When a **min heap (priority queue)** with distance as priority key, total time is **$O((|E| + |V|) \log |V|)$**
 - $O(\log |V|)$ to DECREASE-KEY
 - $O(\log |V|)$ to EXTRACT-MIN



Complexity

- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)
- Option 1: Assume vertices are numbered from 1 to $|V|$. Use an **array** of size $|V|$, store $d[u]$ at the u^{th} entry, have a separate boolean array of size $|V|$ to indicate if u is in set F or not
 - EXTRACT-MIN: we linearly search the array for the smallest: $O(|V|)$
 - DECREASE-KEY: Each operation takes $O(1)$ time
 - Total time: $O(|V|^2 + |E|)$
- Option 2: When a **min heap (priority queue)** with distance as priority key, total time is $O((|E| + |V|) \log |V|)$
 - $O(\log |V|)$ to DECREASE-KEY
 - $O(\log |V|)$ to EXTRACT-MIN
- When $|E|$ is $O(|V|^2)$ [*highly connected, dense graph*], using a min heap is worse than using a linear array



Complexity

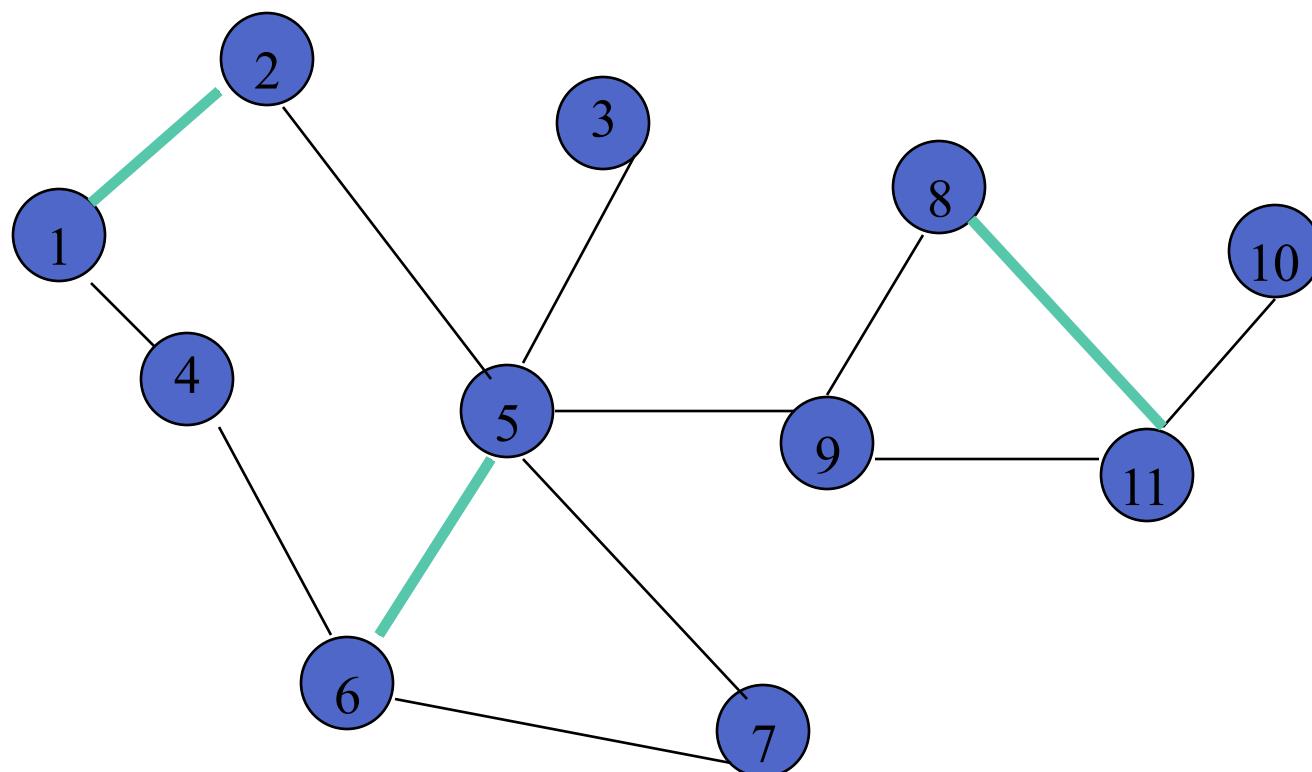
- Depends on how we store set F
- Need two operations:
 - EXTRACT-MIN (to remove the vertex with minimum distance estimate from F)
 - DECREASE-KEY (to revise a distance estimate)
- Option 1: Assume vertices are numbered from 1 to $|V|$. Use an **array** of size $|V|$, store $d[u]$ at the u^{th} entry, have a separate boolean array of size $|V|$ to indicate if u is in set F or not
 - EXTRACT-MIN: we linearly search the array for the smallest: $O(|V|)$
 - DECREASE-KEY: Each operation takes $O(1)$ time
 - Total time: $O(|V|^2 + |E|)$
- Option 2: When a **min heap (priority queue)** with distance as priority key, total time is $O((|E| + |V|) \log |V|)$
 - $O(\log |V|)$ to DECREASE-KEY
 - $O(\log |V|)$ to EXTRACT-MIN
- When $|E|$ is $O(|V|^2)$ [*highly connected, dense graph*], using a min heap is worse than using a linear array
- When a **Fibonacci heap** is used, the total time is $O(|E| + |V| \log |V|)$



Minimum Spanning Tree

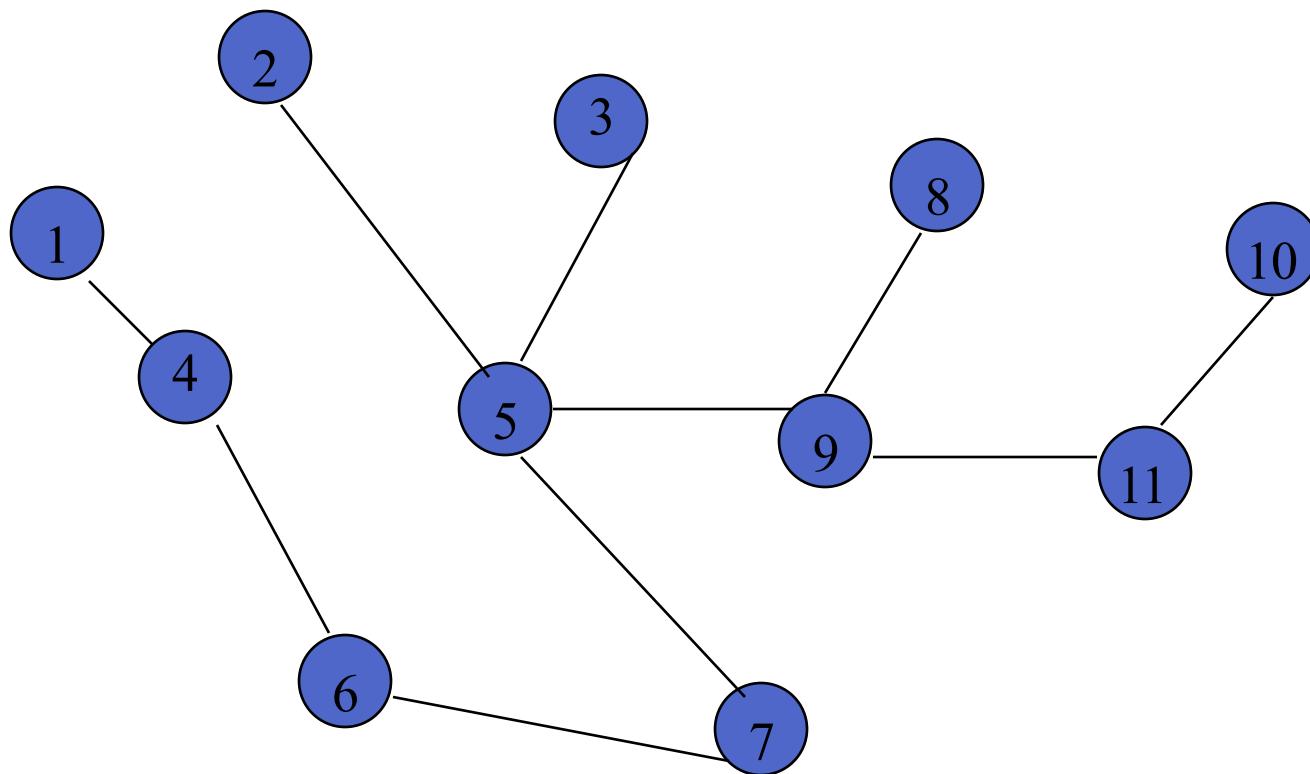
Cycles And Connectedness

Removal of an edge that is on a cycle does not affect connectedness.



Cycles And Connectedness

Connected subgraph with all vertices and minimum number of edges has no cycles.



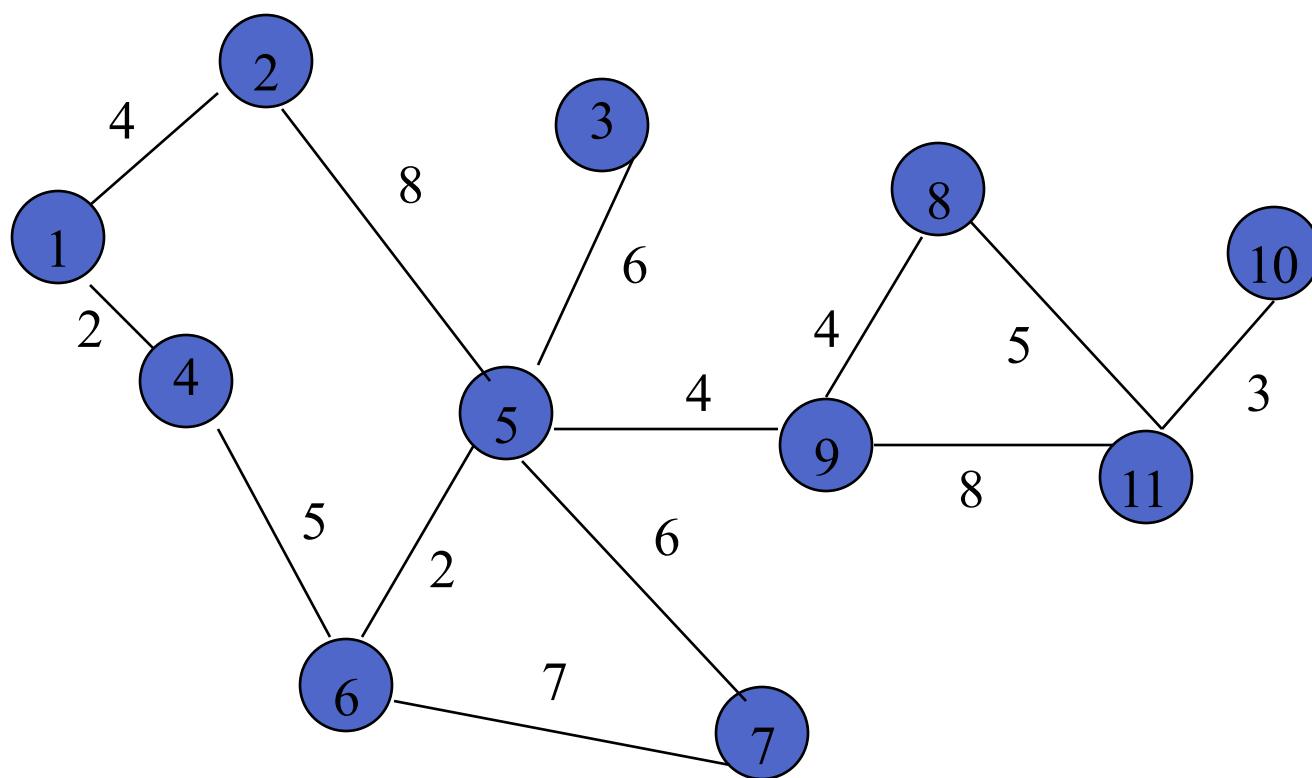


Spanning Tree

- Communication Network Problems
 - Is the network connected?
 - Can we communicate between every pair of cities?
 - Find the components.
 - Want to construct smallest number of feasible links so that resulting network is connected.
- Subgraph that includes all vertices of the original graph.
- Subgraph is a tree.
 - If original graph has n vertices, the spanning tree has n vertices and $n-1$ edges.

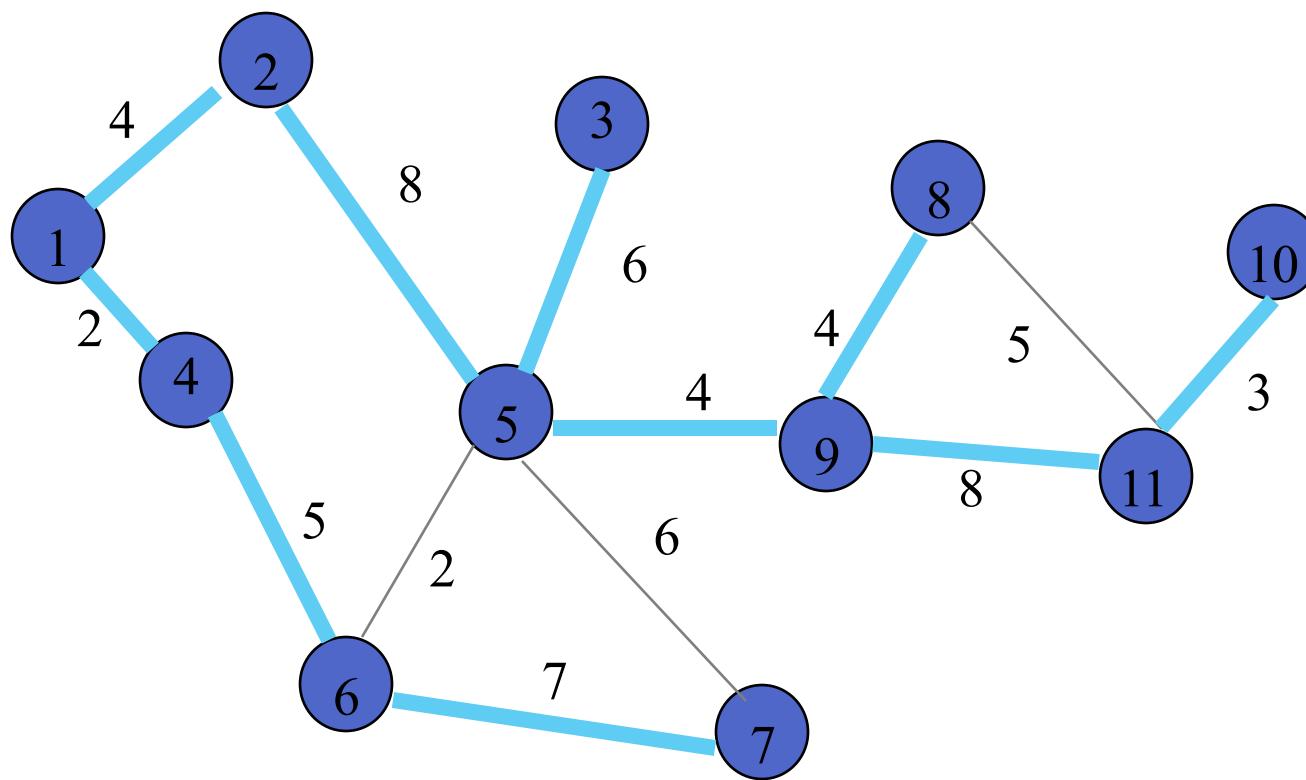
Minimum Cost Spanning Tree

- Tree cost is sum of edge weights/costs.



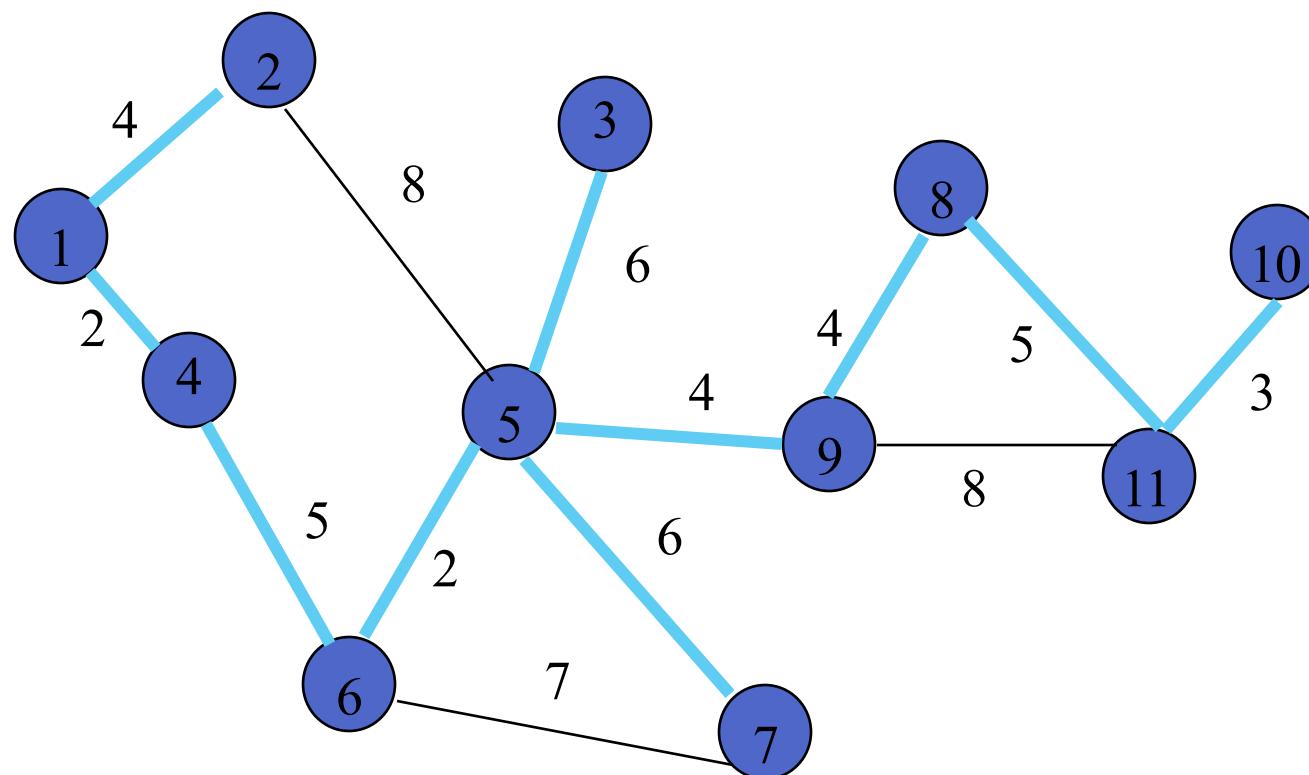
A Spanning Tree

A Spanning tree, cost = 51



Minimum Cost Spanning Tree

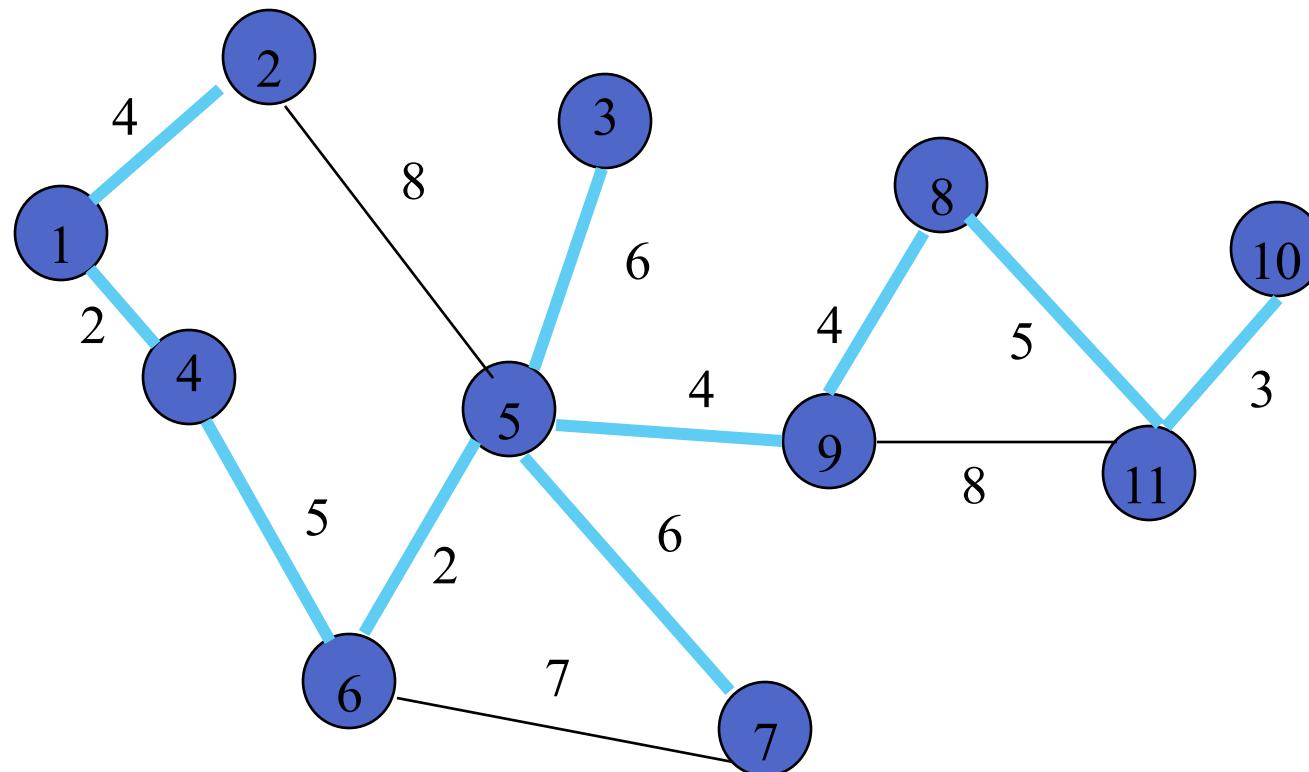
Minimum Spanning tree, cost = 41.



A Wireless Broadcast Tree

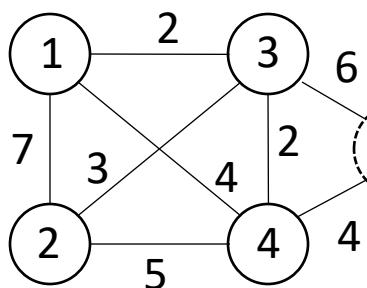
Source = 1, say edge weights = power requirement

$$\text{Cost} = 4 + 2 + 5 + 2 + 6 + 6 + 4 + 4 + 5 + 3 = 41$$



Prim's Minimum Spanning Tree (MST)

Given input graph G



Select a vertex randomly
e.g., vertex 5

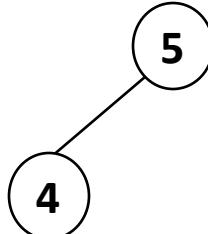


Initialize an empty graph T
with vertex 5

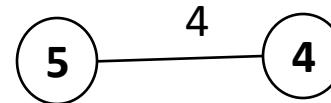


Repeat until all vertices are added to T

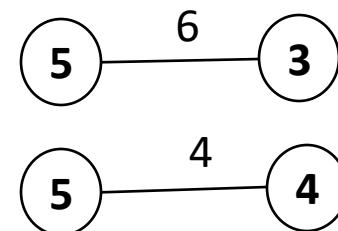
Add the
edge to T



From such edges,
select the edge with
minimum weight



Consider all those edges in
G that connect a vertex in
T to a vertex not in T





Tasks

- **Self study**

- ▶ Read the relevant sections of CLRS “Introduction to Algorithms” textbook



Questions?