

CUDA

Sathish Vadhiyar
High Performance Computing

Hierarchical Parallelism

- Parallel computations arranged as grids
- One grid executes after another
- Grid consists of blocks
- Blocks assigned to SM. A single block assigned to a single SM. Multiple blocks can be assigned to a SM.
 - Max thread blocks executed concurrently per SM = 16

Hierarchical Parallelism

- Block consists of elements
- Elements computed by threads
 - Max threads per thread block = 1024
- A thread executes on a GPU core

Hierarchical Parallelism




Figure 5. Decomposing result data into a grid of blocks partitioned into elements to be computed in parallel.

Thread Blocks

- Thread block – an array of concurrent threads that execute the same program and can cooperate to compute the result
- Has shape and dimensions (1d, 2d or 3d) for threads
- A thread ID has corresponding 1,2 or 3d indices
- Threads of a thread block share memory

CUDA Programming Model




A kernel is executed by a grid of thread blocks

- A thread block is a batch of threads that can cooperate with each other by:

- Sharing data through shared memory
- Synchronizing their execution


- Threads from different blocks cannot cooperate



CUDA Memory Spaces

- Each thread can:
 - Read/write per-thread **registers**
 - Read/write per-thread **local memory**
 - Read/write per-block **shared memory**
 - Read/write per-grid **global memory**
 - Read only per-grid **constant memory**
 - Read only per-grid **texture memory**

- The host can read/write **global, constant, and texture memory** (stored in DRAM)



CUDA Memory Spaces

- Global and Shared Memory introduced before
 - Most important, commonly used
- Local, Constant, and Texture for convenience/performance
 - Local: automatic array variables allocated there by compiler
 - Constant: useful for uniformly-accessed read-only data
 - Cached (see programming guide)
 - Texture: useful for spatially coherent random-access read-only data
 - Cached (see programming guide)
 - Provides address clamping and wrapping

Memory	Location	Cached	Access	Scope (“Who?”)
Local	Off-chip	No	Read/write	One thread
Shared	On-chip	N/A	Read/write	All threads in a block
Global	Off-chip	No	Read/write	All threads + host
Constant	Off-chip	Yes	Read	All threads + host
Texture	Off-chip	Yes	Read	All threads + host

CUDA Programming Language

- Programming language for threaded parallelism for GPUs
- Minimal extension of C
- A serial program that calls parallel kernels
- Serial code executes on CPU
- Parallel kernels executed across a set of parallel threads on the GPU
- Programmer organizes threads into a hierarchy of thread blocks and grids

CUDA Kernels and Threads



- Parallel portions of an application are executed on the device as **kernels**
 - One **kernel** is executed at a time
 - Many threads execute each **kernel**
- Differences between CUDA and CPU threads
 - CUDA threads are extremely lightweight
 - Very little creation overhead
 - Instant switching
 - CUDA uses 1000s of threads to achieve efficiency
 - Multi-core CPUs can use only a few

Definitions:

Device = GPU; *Host* = CPU

Kernel = function that runs on the device

CUDA C

- Built-in variables:
 - `threadIdx.{x,y,z}` – thread ID within a block
 - `blockIdx.{x,y,z}` – block ID within a grid
 - `blockDim.{x,y,z}` – number of threads within a block
 - `gridDim.{x,y,z}` – number of blocks within a grid
- `kernel<<<nBlocks,nThreads>>>(args)`
 - Invokes a parallel kernel function on a grid of nBlocks where each block instantiates nThreads concurrent threads

Example: Summing Up

```
void addMatrix
    (float *a, float *b, float *c, int N)
{
    int i, j, idx;
    for (i = 0; i < N; i++) {
        for (j = 0; j < N; j++) {
            idx = i + j*N;
            c[idx] = a[idx] + b[idx];
        }
    }
}
void main()
{
    ...
    addMatrix(a, b, c, N);
}
```

(a)

```
kernel function
__global__ void addMatrixG
    (float *a, float *b, float *c, int N)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int j = blockIdx.y*blockDim.y + threadIdx.y;
    int idx = i + j*N;
    if (i < N && j < N)
        c[idx] = a[idx] + b[idx];
}

grid of kernels
void main()
{
    dim3 dimBlock (blocksize, blocksize);
    dim3 dimGrid (N/dimBlock.x, N/dimBlock.y);
    addMatrixG<<<dimGrid, dimBlock>>>(a, b, c, N);
}
```

(b)

Figure 8. Serial C (a) and CUDA C (b) examples of programs that add arrays.

Variable Qualifiers (GPU code)

- **__device__**
 - stored in device memory (large, high latency, no cache)
 - Allocated with `cudaMalloc` (**__device__** qualifier implied)
 - accessible by all threads
 - lifetime: application
- **__constant__**
 - same as **__device__**, but cached and read-only by GPU
 - written by CPU via `cudaMemcpyToSymbol`(...) call
 - lifetime: application
- **__shared__**
 - stored in on-chip shared memory (very low latency)
 - accessible by all threads in the same thread block
 - lifetime: kernel launch
- **Unqualified variables:**
 - scalars and built-in vector types are stored in registers
 - arrays of more than 4 elements stored in device memory

General CUDA Steps

1. Copy data from CPU to GPU
 2. Compute on GPU
 3. Copy data back from GPU to CPU
- By default, execution on host doesn't wait for kernel to finish
 - General rules:
 - Minimize data transfer between CPU & GPU
 - Maximize number of threads on GPU

CUDA Elements

- `cudaMalloc` – for allocating memory in device
- `cudaMemcpy` – for copying data to allocated memory from host to device, and from device to host
- `cudaFree` – freeing allocated memory
- `void syncthreads__()` – synchronizing all threads in a block like barrier

EXAMPLE 1: MATRIX VECTOR MULTIPLICATION

Kernel

```
1  __global__ void matvec_mul(int m, int n, double *A, double *x,
2    double *y)
3  {
4    int row, col;
5    double sum;
6
7    row = blockIdx.x*blockDim.x+threadIdx.x;
8
9    sum=0;
10   if (row < m){
11     for (col=0; col<n; col++){
12       sum += A[row*n+col]*x[col];
13     }
14   }
15   y[row]=sum;
16 }
17 }
```

Host Program

```
18 int main(int argc, char** argv){  
19     ...  
20     size_t size_A, size_x, size_y;  
21     double *A, *x, *y;  
22     double *dA, *dx, *dy;  
23     ...  
24  
25     m = ... /* rows */; n = ... /* cols */  
26  
27     size_A = sizeof(double)*m*n;    size_x = sizeof(double)*n;  
28     size_y = sizeof(double)*m;  
29  
30     A = (double*)malloc(size_A);    x = (double*)malloc(size_x);    y =  
31         (double*)malloc(size_y);  
32  
33     /* Allocate on the device memory */  
34     cudaMalloc((void **) &dA, size_A);  
35     cudaMalloc((void **) &dx, size_x);  
36     cudaMalloc((void **) &dy, size_y);  
37  
38     /* Initialize A and x */  
39     /* Initialize y */  
40     for(i=0; i<m; i++) y[i] = 0;
```

Host Program

```
40  /* Copy A and x to the device */
41  cudaMemcpy(dA, A, size_A, cudaMemcpyHostToDevice);
42  cudaMemcpy(dx, x, size_x, cudaMemcpyHostToDevice);
43
44  numThreadsPerBlock = 1024;  numBlocks = m/ numThreadsPerBlock;
45
46  dim3 dimGrid(numBlocks);
47  dim3 dimBlock(numThreadsPerBlock);
48  matvec_mul<<< dimGrid, dimBlock >>>(m, n, dA, dx, dy);
49
50  cudaMemcpy(y, dy, size_y, cudaMemcpyDeviceToHost);
51
52  cudaFree(dA);  cudaFree(dx);  cudaFree(dy);
53
54  free(A);  free(x);  free(y);
55
56 }
```


EXAMPLE 1, VERSION 2: ACCESS FROM SHARED MEMORY

```
1  __global__ void matvec_mul( int m, int n, double *A, double *x,
2    double *y)
3  {
4    int row, col;
5    double sum;
6
7    __shared__ int sx[BLOCK_SIZE];
8
9    sx[threadIdx.x] = x[threadIdx.x];
10   __syncthreads();
11
12   row = blockIdx.x*blockDim.x+threadIdx.x;
13
14   sum=0;
15   if (row < m){
16     for (col=0; col<n; col++){
17       sum += A[row*n+col]*sx[col];
18     }
19   }
20   y[row]=sum;
21 }
```

EXAMPLE 2: MATRIX MULTIPLICATION

Matrix Multiplication Example

- Computing the product C of two matrices:
 $A : (wA, hA)$
 $B : (wB, wA)$.
- Each thread block computes one square sub-matrix C_{sub} of C ;
- Each thread within the block computes one element of C_{sub} .



Example 1: Matrix Multiplication

Host matrix multiplication code



```
void Mul(const float* A, const float* B, int hA, int wA, int wB, float* C)
{
    int size;
    // Load Input matrices A and B to the device
    float* Ad;
    size = hA * wA * sizeof(float);
    cudaMalloc((void**)&Ad, size);
    cudaMemcpy(Ad, A, size, cudaMemcpyHostToDevice);

    .
    // Allocate memory for output matrix C on the device
    float* Cd;
    size = hA * wB * sizeof(float);
    cudaMalloc((void**)&Cd, size);
```

Example 1

```
// Compute the execution configuration assuming
// the matrix dimensions are multiples of BLOCK_SIZE
dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);
dim3 dimGrid(wB / dimBlock.x, hA / dimBlock.y);
// Launch the device computation
Muld<<<dimGrid, dimBlock>>>(Ad, Bd, wA, wB, Cd);
// Read Ouput matrix C from the device
cudaMemcpy(C, Cd, size, cudaMemcpyDeviceToHost);
// Free device memory
cudaFree(Ad);
.
}
```

Example 1

Device matrix multiplication function



```
__global__ void Muld ( float* A, float* B, int wA, int wB, float* C )
{
    // Setup aBegin, aEnd, aStep  bBegin, bStep based on Block index and Block size

    // The element of the block sub-matrix that is computed by the thread
    float Csub = 0;
    // Loop over all the sub-matrices of A and B required to compute the block sub-matrix
    for (int a = aBegin, b = bBegin; a <= aEnd; a += aStep, b += bStep) {

        // Shared memory for the sub-matrices of A and B
        __shared__ float As [ BLOCK_SIZE ] [ BLOCK_SIZE ];
        __shared__ float Bs [ BLOCK_SIZE ] [ BLOCK_SIZE ];
```

Example 1

```
// Load the matrices from global memory to shared memory; each thread loads one element of each matrix
As[ty][tx] = A[a + wA * ty + tx];
Bs[ty][tx] = B[b + wB * ty + tx];

// Synchronize to make sure the matrices are loaded
__syncthreads();

// Multiply the two matrices together; each thread computes one element/ of the block sub-matrix
for (int k = 0; k < BLOCK_SIZE; ++k)
    Csub += As[ty][k] * Bs[k][tx];
```

Example 1

```
// Synchronize to make sure that the preceding computation is done before loading two new
// sub-matrices of A and B in the next iteration
__syncthreads();
}

// Write the block sub-matrix to global memory; each thread writes one element
int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;
C[c + wB * ty + tx] = Csub;
}
```

EXAMPLE 2: REDUCTION


Example: Reduction

- Tree based approach used within each thread block
- In this case, partial results need to be communicated across thread blocks
- Hence, global synchronization needed across thread blocks

Reduction

- But CUDA does not have global synchronization –
 - expensive to build in hardware for large number of GPU cores
- Solution
 - Decompose into multiple kernels
 - Kernel launch serves as a global synchronization point

Illustration



Host Code

```
int main(){

    int* h_idata, h_odata; /* host data*/
    Int *d_idata, d_odata; /* device data*/

    /* copying inputs to device memory */
    cudaMemcpy(d_idata, h_idata, bytes, cudaMemcpyHostToDevice) ;
    cudaMemcpy(d_odata, h_idata, numBlocks*sizeof(int),
               cudaMemcpyHostToDevice) ;

    int numThreadsperBlock = (n < maxThreadsperBlock) ? n : maxThreadsperBlock;
    int numBlocks = n / numThreadsperBlock;
    dim3 dimBlock(numThreads, 1, 1); dim3 dimGrid(numBlocks, 1, 1);

    reduce<<< dimGrid, dimBlock >>>(d_idata, d_odata);
```

Host Code

```
int s=numBlocks;
while(s > 1) {
    numThreadsperBlock = (s< maxThreadsperBlock) ? s :
maxThreadsperBlock;  numBlocks = s / numThreadsperBlock;
    dimBlock(numThreads, 1, 1);  dimGrid(numBlocks, 1, 1);
    reduce<<< dimGrid, dimBlock, smemSize >>>(d_idata,
d_odata);
    s = s / numThreadsperBlock;
}
}
```

Device Code

```
__global__ void reduce(int *g_idata, int *g_odata)
{
    extern __shared__ int sdata[];

    // load shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();

    // do reduction in shared mem
    for(unsigned int s=1; s < blockDim.x; s *= 2) {
        if ((tid % (2*s)) == 0)
            sdata[tid] += sdata[tid + s];
        __syncthreads();
    }

    // write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

- For more information...
- CUDA SDK code samples – NVIDIA -
http://www.nvidia.com/object/cuda_get_samples.html

BACKUP

EXAMPLE 3: SCAN

Example: Scan or Parallel-prefix sum

Definition: The all-prefix-sums operation takes a binary associative operator \oplus , and an array of n elements


$$[a_0, a_1, \dots, a_{n-1}],$$

and returns the array


$$[a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-1})].$$

- Using binary tree
- An upward reduction phase (reduce phase or up-sweep phase)
 - Traversing tree from leaves to root forming partial sums at internal nodes
- Down-sweep phase
 - Traversing from root to leaves using partial sums computed in reduction phase

Up Sweep



Down Sweep



Host Code

- int main(){
- const unsigned int num_threads = num_elements / 2;
- /* cudaMalloc d_idata and d_odata */
- cudaMemcpy(d_idata, h_data, mem_size,
cudaMemcpyHostToDevice));
- dim3 grid(256, 1, 1); dim3 threads(num_threads, 1, 1);
- scan<<< grid, threads>>> (d_odata, d_idata);
- cudaMemcpy(h_data, d_odata[i], sizeof(float) *
num_elements, cudaMemcpyDeviceToHost
- /* cudaFree d_idata and d_odata */
- }

Device Code

```
__global__ void scan_workefficient(float *g_odata, float *g_idata, int n)
{
    // Dynamically allocated shared memory for scan kernels
    extern __shared__ float temp[];

    int thid = threadIdx.x;    int offset = 1;

    // Cache the computational window in shared memory
    temp[2*thid] = g_idata[2*thid];
    temp[2*thid+1] = g_idata[2*thid+1];

    // build the sum in place up the tree
    for (int d = n>>1; d > 0; d >>= 1)
    {
        __syncthreads();

        if (thid < d)
        {
            int ai = offset*(2*thid+1)-1;
            int bi = offset*(2*thid+2)-1;

            temp[bi] += temp[ai];
        }

        offset *= 2;
    }
}
```

Device Code

```
// scan back down the tree

// clear the last element
if (thid == 0)      temp[n - 1] = 0;
// traverse down the tree building the scan in place
for (int d = 1; d < n; d *= 2)
{
    offset >>= 1;
    __syncthreads();
    if (thid < d)
    {
        int ai = offset*(2*thid+1)-1;
        int bi = offset*(2*thid+2)-1;

        float t  = temp[ai];
        temp[ai] = temp[bi];
        temp[bi] += t;
    }
}

__syncthreads();

// write results to global memory
g_odata[2*thid]  = temp[2*thid];  g_odata[2*thid+1] = temp[2*thid+1];
}
```