

Datatypes

Feb 2, 2021

Revision: Example 1

```
#define numElements 10

int main( int argc, char *argv[])
{
    int arr[numElements];
    int myrank, size;
    MPI_Status status;
    MPI_Request request;
    int count, recvvarr[numElements];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    MPI_Isend(arr, numElements, MPI_INT, myrank, 99, MPI_COMM_WORLD, &request);
    MPI_Recv(recvvarr, numElements, MPI_INT, myrank, 99, MPI_COMM_WORLD, &status);
    MPI_Wait (&request, &status);

    printf ("done\n");
}
```

Valid code?

Revision: Example 2

```
int main( int argc, char *argv[])
{
    int N = atoi (argv[1]);
    int myrank, count;
    double data[N];
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );

    if (myrank == 0)      /* code for process 0 */
    {
        MPI_Send(data, N, MPI_INT, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1) /* code for process 1 */
    {
        MPI_Recv(data, N, MPI_DOUBLE, 0, 99, MPI_COMM_WORLD, &status);
    }

    MPI_Finalize();
    return 0;
}
```

Revision: Example 3

```
int main( int argc, char *argv[] )
{
    int N = atoi ( argv[1] );
    int myrank, count;
    double data[N];
    int recvdata[N];
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );

    if (myrank == 0)          /* code for process 0 */
    {
        MPI_Send(data, N, MPI_DOUBLE, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1)      /* code for process 1 */
    {
        MPI_Recv(recvdata, N, MPI_INT, 0, 99, MPI_COMM_WORLD, &status);
        MPI_Get_count( &status, MPI_INT, &count );
        printf ("%d\n", count);
    }

    MPI_Finalize();
    return 0;
}
```

Truncate
error

Revision: Example 4

```
int main( int argc, char *argv[])
{
    int N = atoi (argv[1]);
    int myrank, count;
    double data[N];
    int recvdata[2*N];
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );

    if (myrank == 0)      /* code for process 0 */
    {
        MPI_Send(data, N, MPI_DOUBLE, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1) /* code for process 1 */
    {
        MPI_Recv(recvdata, 2*N, MPI_INT, 0, 99, MPI_COMM_WORLD, &status);
        MPI_Get_count( &status, MPI_INT, &count );
        printf ("%d %d\n", count, recvdata[0]);
    }

    MPI_Finalize();
    return 0;
}
```

Output?

Revision: Safety

```
IF (rank.EQ.0) THEN
    MPI_SSEND (buf1, count, MPI_INT, 1, 0, comm, req1)
    MPI_SEND (buf2, count, MPI_INT, 1, 1, comm, req2)
ELSE IF (rank.EQ.1) THEN
    MPI_IRecv (buf1, count, MPI_INT, 0, 0, comm, req1)
    MPI_RECV (buf2, count, MPI_INT, 0, 1, comm, req2)
    MPI_WAIT (req1, status1)
END IF
```

Datatypes

`sizeof (MPI_DOUBLE);`

`MPI_Type_size`

```
#include <stdio.h>
#include <stdlib.h>
#include "mpi.h"

int main( int argc, char *argv[])
{
    int N = atoi (argv[1]);
    int myrank;
    double data[N];
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );

    //initialize data

    if (myrank == 0)      /* code for process 0 */
    {
        MPI_Send(data, N, MPI_DOUBLE, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1) /* code for process 1 */
    {
        MPI_Recv(data, N, MPI_DOUBLE, 0, 99, MPI_COMM_WORLD, &status);
    }

    int size1, size2;
    MPI_Type_size (MPI_DOUBLE, &size1);
    size2 = sizeof (MPI_DOUBLE);
    printf ("size of MPI_DOUBLE %d %d\n", size1, size2);

    MPI_Finalize();
    return 0;
}
```

MPI_PACK

```
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank) ;
MPI_Comm_size(MPI_COMM_WORLD, &size);

// initialize data
for (int i=0; i<M; i++)
    for (int j=0; j<N; j++)
        array2D[i][j] = myrank+i+j;

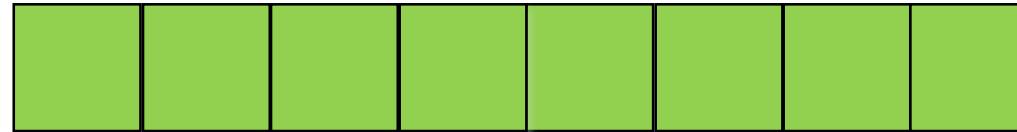
sTime = MPI_Wtime();
if (myrank == 0) {
    // pack the last element of every row (N ints)
    for (int j=0; j<N; j++) {
        MPI_Pack (&array2D[j][M-1], 1, MPI_INT, buffer, 400, &position, MPI_COMM_WORLD);
        printf ("packed %d %d\n", j, position);
    }
    MPI_Send (buffer, position, MPI_PACKED, 1, 1, MPI_COMM_WORLD);
}
else {
    // receive N ints
    if (myrank == 1)
        MPI_Recv (buffer, count, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);
    // verify
    MPI_Get_count (&status, MPI_INT, &count);
}
eTime = MPI_Wtime();
time = eTime - sTime;

printf ("%lf\n", time);
```

source

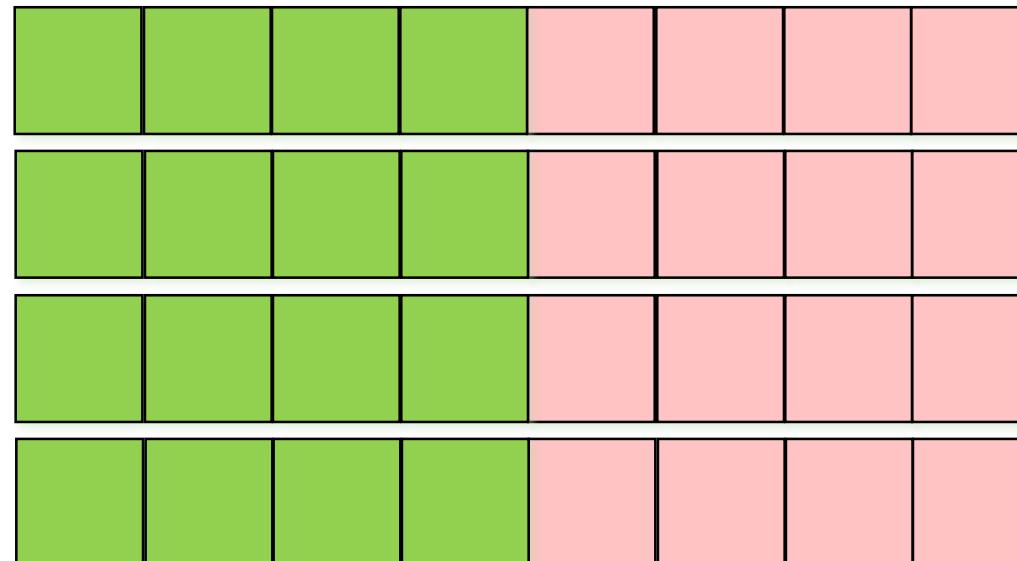
dest

MPI Derived Datatypes



Q: How would you send contiguous data ?

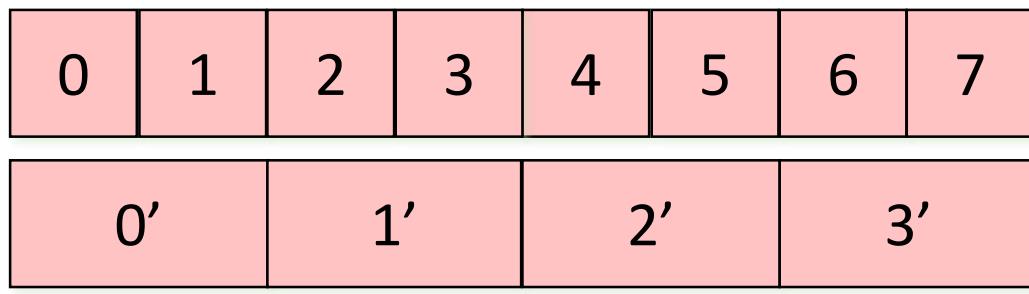
A: MPI_Send (specify starting address, #bytes)



Q: How would you send block of data (non-contiguous data transfer) ?

A: MPI derived datatypes

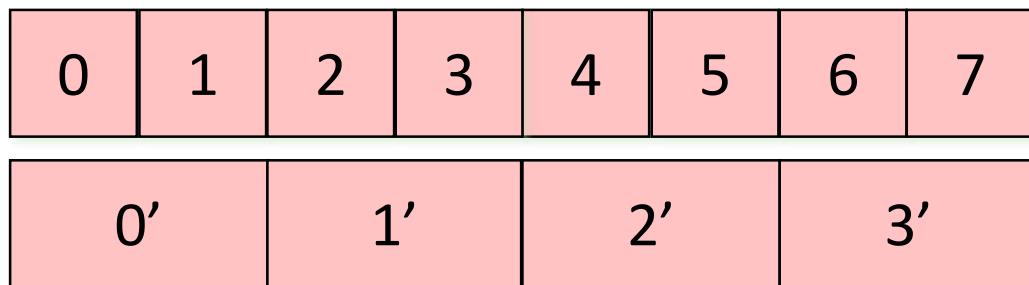
MPI Derived Datatypes



Contiguous

- Replication of datatype into contiguous locations
- `MPI_Type_contiguous(count, oldtype, newtype)`
 - count – replication count

MPI_Type_contiguous



```
int buffer[200];
MPI_Type_contiguous (2, MPI_INT, &twooints)
MPI_Type_commit (&twooints)
MPI_Send (buffer, 10, twooints, ....)
```

MPI_Type_vector



MPI_Type_vector

count = #blocks

blocklength = #elements in each block

stride = #elements between start of each block

count = 3, blocklength = 1, stride = 4

count = 3, blocklength = 2, stride = 4

`MPI_Type_vector (count, blocklength,
stride, oldtype, newtype)`

Code

```
int N = atoi (argv[1]);
int count = atoi (argv[2]);
int blocklen = atoi (argv[3]);
int stride = atoi (argv[4]);
int numVectors = atoi (argv[5]);
int data[N];

MPI_Type_vector (count, blocklen, stride, MPI_INT, &newvtype);
MPI_Type_commit (&newvtype);

//initialize data
for (int i=0; i<N; i++)
    data[i]=0;

if (myrank == 0)      /* code for process 0 */
{
    for (int i=0; i<N; i++)
        data[i]=i;
    MPI_Send(data, numVectors, newvtype, 1, 99, MPI_COMM_WORLD);
}
else if (myrank == 1) /* code for process 1 */
{
    printf("\n");
    MPI_Recv(data, numVectors, newvtype, 0, 99, MPI_COMM_WORLD, &status);
    MPI_Get_count (&status, MPI_INT, &recvcount);
    for (int i=0; i<N; i++)
        printf ("%d ", data[i]);
    printf("\n\n");
}

MPI_Type_free (&newvtype);
```

Output

Code – Sending selected columns

```
int N = atoi (argv[1]);
int column = atoi (argv[2]);
int count = atoi (argv[3]);
int blocklen = atoi (argv[4]);
int stride = atoi (argv[5]);
int data[N][N], received[N*blocklen];

MPI_Type_vector (count, blocklen, stride, MPI_INT, &newvtype);
MPI_Type_commit (&newvtype);

//initialize data
for (int i=0; i<N; i++)
    for (int j=0; j<N; j++)
        data[i][j]=0;

if (myrank == 0)      /* code for process 0 */
{
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++)
            data[i][j]=column+i+j;
    MPI_Send(&data[0][column], 1, newvtype, 1, 99, MPI_COMM_WORLD);
}
else if (myrank == 1) /* code for process 1 */
{
    printf ("\n");
    MPI_Recv(received, count*blocklen, MPI_INT, 0, 99, MPI_COMM_WORLD, &status);
    for (int i=0; i<count*blocklen; i++)
        printf ("%d ", received[i]);
    printf ("\n\n");
}
```

Number of
elements = 1

Output

```
class $ mpirun -np 2 ./vector2D 4 3 4 1 4
```

```
6 7 8 9
```

```
class $ mpirun -np 2 ./vector2D 4 0 4 1 4
```

```
0 1 2 3
```

```
class $ mpirun -np 2 ./vector2D 4 0 4 2 4
```

```
0 1 1 2 2 3 3 4
```

```
class $ mpirun -np 2 ./vector2D 10 8 10 2 10
```

```
16 17 17 18 18 19 19 20 20 21 21 22 22 23 23 24 24 25 25 26
```

0	1	2	3
1	2	3	4
2	3	4	5
3	4	5	6

MPI_Type_indexed

count = #blocks

blocklengths = #elements in each block

displacements = displacement of start of each block



count = 3, blocklengths = 1,2,1
displacements = 0,4,8



count = 3, blocklengths = 2,4,1
displacements = 0,3,8

MPI_Type_indexed

MPI_Type_indexed (count, blocklengths,
displacements, oldtype, newtype)

Code

```
MPI_Datatype newtype;
int N = atoi (argv[1]);
int numElements = atoi (argv[2]);
int count = 3;
int blocklengths[] = {1, 2, 3};
int displacements[] = {0, 3, 6};
int data[N];

MPI_Type_indexed (count, blocklengths, displacements, MPI_INT, &newtype);
MPI_Type_commit (&newtype);

//initialize data
for (int i=0; i<N; i++)
    data[i]=0;

if (myrank == 0)      /* code for process 0 */
{
    for (int i=0; i<N; i++)
        data[i]=i;
    MPI_Send(data, numElements, newtype, 1, 99, MPI_COMM_WORLD);
}

if (myrank == 1)      /* code for process 1 */
{
    MPI_Recv(data, numElements, newtype, 0, 99, MPI_COMM_WORLD, &status);
    for (int i=0; i<N; i++)
        printf ("%d ", data[i]);
    printf ("\n");
}
```

```
class $ mpirun -np 2 ./indexed 30 1
```

0 0 0 3 4 0 6 7 8 0

```
class $ mpirun -np 2 ./indexed 30 2
```

0 0 0 3 4 0 6 7 8 9 0 0 12 13 0 15 16 17 0 0 0 0 0 0 0 0 0 0 0 0

Lower Triangular matrix

```
int N = atoi (argv[1]);
int data[N][N];
int count = N;
int blocklengths[N], displacements[N];

for (int i=0; i<N; i++)
    blocklengths[i] = i+1, displacements[i] = i*N;
MPI_Type_indexed (count, blocklengths, displacements, MPI_INT, &newtype);
MPI_Type_commit (&newtype);

//initialize data
for (int i=0; i<N; i++)
    for (int j=0; j<N; j++)
        data[i][j]=0;

if (myrank == 0)      /* code for process 0 */
{
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++)
            data[i][j]=i+j;
    MPI_Send(data, 1, newtype, 1, 99, MPI_COMM_WORLD);
}
```

```
mpirun -np 2 ./indexed2D 4
```

0	1	2	3
1	2	3	4
2	3	4	5
3	4	5	6

MPI_Type_create_subarray

```
MPI_Datatype newtype;
int ndims = 1;
int array_of_sizes[] = {atoi (argv[1])};
int array_of_subsizes[] = {atoi (argv[2])};
int array_of_starts[] = {atoi (argv[3])};
int N = array_of_sizes[0];
int data[N];

MPI_Type_create_subarray (ndims, array_of_sizes, array_of_subsizes, array_of_starts, MPI_ORDER_C, MPI_INT, &newtype);
MPI_Type_commit (&newtype);

//initialize data
for (int i=0; i<N; i++)
    data[i]=0;

if (myrank == 0) /* code for process 0 */
{
    for (int i=0; i<N; i++)
        data[i]=i;
    MPI_Send(data, 1, newtype, 1, 99, MPI_COMM_WORLD);
}
else if (myrank == 1) /* code for process 1 */
{
    printf("\n");
    MPI_Recv(data, 1, newtype, 0, 99, MPI_COMM_WORLD, &status);
    for (int i=0; i<N; i++)
        printf ("%d ", data[i]);
    printf("\n\n");
}

MPI_Type_free (&newtype);
```

0	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---

0	1	2	3
1	2	3	4
2	3	4	5
3	4	5	6

Output (Array dimension=1)



```
class $ mpirun -np 2 ./subarray 10 4 2  
0 0 2 3 4 5 0 0 0 0  
  
class $ mpirun -np 2 ./subarray 10 4 4  
0 0 0 0 4 5 6 7 0 0  
  
class $ mpirun -np 2 ./subarray 20 10 4  
0 0 0 0 4 5 6 7 8 9 10 11 12 13 0 0 0 0 0
```

Code (2D subarray)

```
MPI_Datatype newtype;
int ndims = 2;
int array_of_sizes[] = {atoi(argv[1]), atoi(argv[2])};
int array_of_subsizes[] = {2, 2};
int array_of_starts[] = {1, 1};
int data[array_of_sizes[0]][array_of_sizes[1]];

MPI_Type_create_subarray (ndims, array_of_sizes, array_of_subsizes, array_of_starts, MPI_ORDER_C,
MPI_Type_commit (&newtype);

//initialize data
for (int i=0; i<array_of_sizes[0]; i++)
    for (int j=0; j<array_of_sizes[1]; j++)
        data[i][j]=0;

if (myrank == 0) /* code for process 0 */
{
    for (int i=0; i<array_of_sizes[0]; i++)
        for (int j=0; j<array_of_sizes[1]; j++)
            data[i][j]=i+j;
    MPI_Send(data, 1, newtype, 1, 99, MPI_COMM_WORLD);
}
else if (myrank == 1) /* code for process 1 */
{
    printf ("\n");
    MPI_Recv(data, 1, newtype, 0, 99, MPI_COMM_WORLD, &status);
    for (int i=0; i<array_of_sizes[0]; i++) {
        for (int j=0; j<array_of_sizes[1]; j++)
            printf ("%d ", data[i][j]);
        printf ("\n");
    }
}
```

0	1	2	3
1	2	3	4
2	3	4	5
3	4	5	6

```
class $ mpirun -np 2 ./subarray2D 4 4
```

```
0 0 0 0
0 2 3 0
0 3 4 0
0 0 0 0
```

```
0 0 0 0
```

MPI Derived Datatypes Summary

MPI_Datatype newtype

- MPI_Type_contiguous (count, oldtype, newtype)
- MPI_Type_vector (count, blocklength, stride, oldtype, newtype)
- MPI_Type_indexed (count, blocklengths, displacements, oldtype, newtype)
- MPI_Type_create_subarray (ndims, array_of_sizes, array_of_subsizes, array_of_starts, order, oldtype, newtype)
- MPI_Type_create_struct (count, array_of_blocklengths, array_of_displacements, array_of_types, newtype)

MPI_Type_commit (newtype)

.....

MPI_Type_free (newtype)

Halo Exchange

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

- Every time step
 - Stencil computation
 - $\text{Val}_{t+1} = \text{Average of } \text{Val}_t \text{ 4 neighboring points}$
 - Communicate halo regions
 - Multiple MPI_Sends
 - MPI_Pack + MPI_Send
 - MPI Derived datatype + MPI_Send

Assignment 1

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Compare the three methods

- Multiple MPI_Sends
- MPI_Pack + MPI_Send
- MPI Derived datatype + MPI_Send
- Send and Recv (Blocking/non-blocking, any mode), Sendrecv, ...

for P in 9, 16, 25, 36,

for N_per_ps in 64, 256, 1024,

mpirun –np P ./halo N_per_ps

General Instructions

- <https://git.cse.iitk.ac.in>
- Create private project named ‘CS633-2020-21-2’ and add pmalakar and TAs (cpathare, fahad, lavleshm, piprotar, mabir, samvid, sharhp, tusharag) and your assignment group member
- Give at least reporter access to the instructor and the TAs
- Create subdirectories for each assignment
 - AssignmentN for assignment #N
- Run your code on CSE lab cluster
 - Notify us if you find most of these are unreachable
 - Do not hard code host names (dynamically generate)
 - Optionally use NodeAllocator

Assignment1 Logistics

- Directory named Assignment1 (not assignment, Assignment- etc.) on git
- Reporter access at least (not guest)
- Send a **submission email** to me and the TAs as soon as you submit
 - [CS633] Assignment1
- Follow the instructions regarding file names **VERY** carefully
- The necessary files should not reside inside subsubdirectories
- Your ‘job script’ should be executable and exactly in the format/location as specified
- Plots should always be boxplots

Assignment Logistics

- Grading based on your effort as well
- Credit for early submissions (+5 / day)
 - Max credit: +15 / assignment
- A total of 2 extra days may be taken
- Score reduction for late submissions (-5 / day)
 - Max 3 late days / assignment
- Indicate in your submission email how many early/extra/late days you have used
- None of the assignments can be completed in a day!
- Start coding early
 - Parallel programs are harder to debug!