

# Communicators

Feb 5, 2021


# Communicator

- Object containing a group of processes
- Representative of communication domain
- Predefined:
  - MPI\_COMM\_WORLD
  - MPI\_COMM\_SELF
- Contains a mapping from MPI process ranks to processor ids
- Memory proportional to #processes in the group
- Several communication contexts may co-exist within a single communicator


# Process Group

- Ordered set of processes
- Ranks are contiguous
- Base group – associated with MPI\_COMM\_WORLD
- MPI\_Group object
  - MPI\_Group\_rank, MPI\_Group\_size
- Unions and intersections of groups
- Not used in communication context


# MPI\_COMM\_WORLD




# Sub-communicator



# Intra-Communicators



# Inter-Communicators




Inter-communicator  
between two groups

# MPI\_COMM\_SPLIT

`MPI_Comm_split (MPI_Comm oldcomm, int color, int key, MPI_Comm *newcomm)`

- Collective call
- Logically divides based on *color*
  - Same color processes form a group
  - Some processes may not be part of newcomm (`MPI_UNDEFINED`)
- Rank assignment based on *key*

# Logical subsets of processes



How do you assign one color to odd processes and another color to even processes ?  
color = rank % 2

# Example code

```
int newrank, newsize, color = myrank%2;  
MPI_Comm newcomm;  
  
MPI_Comm_split (MPI_COMM_WORLD, color, myrank, &newcomm);  
  
MPI_Comm_rank (newcomm, &newrank);  
MPI_Comm_size (newcomm, &newsize);  
printf ("%d: %d of %d\n", myrank, newrank, newsize);  
  
MPI_Comm_free (&newcomm);
```


Output  
for P=8

0: 0 of 4  
1: 0 of 4  
2: 1 of 4  
3: 1 of 4  
4: 2 of 4  
5: 2 of 4  
6: 3 of 4  
7: 3 of 4

# MPI\_Group

```
MPI_Group_incl (  
    MPI_Group g_group,  
    int N,  
    const int ranks[],  
    MPI_Group *new_group);
```

```
MPI_Comm_create_group (  
    MPI_COMM_WORLD,  
    new_group,  
    tag,  
    &new_comm);
```



# MPI\_Group Code

```
int ranks[] = {0,2,4,6,8};
```

```
// Construct a group containing even ranks in g_group
```

```
MPI_Group new_group;
```

```
MPI_Group_incl (g_group, N, ranks, &new_group);
```


```
// Create a new communicator based on the group
```

```
MPI_Comm new_comm;
```

```
MPI_Comm_create_group (MPI_COMM_WORLD, new_group, tag,  
&new_comm);
```


# MPI Sub-communicators

**MPI\_COMM\_WORLD**



Overlapping


**MPI\_COMM\_WORLD**




Non-overlapping

# Usage of Sub-communicators

- Producer processes





- Consumer processes





# Duplicate Communicator

`MPI_Comm_dup (comm, newcomm)`



# Process Mapping



# Groups and Sub-communicators

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31

4 x 8 2D virtual process topology

```
// set up rows and columns in 2D virtual topology
int rows = 4;
int cols = numtasks/rows;

// create new group ranks array
int ranks[cols], i, j=-1;
for (i=(myrank/cols)*cols; i<((myrank/cols)+1)*cols; i++)
    ranks[++j] = i;

// get the world group
MPI_Group g_group;
MPI_Comm_group (MPI_COMM_WORLD, &g_group);

// create new groups
MPI_Group new_group;
MPI_Group_incl (g_group, cols, ranks, &new_group);

// create new communicators
MPI_Comm new_comm;
MPI_Comm_create_group (MPI_COMM_WORLD, new_group, myrank/cols, &new_comm);

// size of new communicators
int new_size, new_rank;
MPI_Comm_size (new_comm, &new_size);
MPI_Comm_rank (new_comm, &new_rank);
printf ("Old rank %d, new rank %d\n", myrank, new_rank);
```

Yet another way: int MPI\_Group\_range\_incl (MPI\_Group group, int n, int ranges[][3], MPI\_Group \*newgroup)

# MPI\_Cart\_create

0	1	2	3
4	5	6	7
8	9	10	11

Cartesian topology

`MPI_Cart_create (MPI_Comm comm_old, int ndims, int *dims, int *periodic, int reorder, MPI_Comm *comm_cart)`

`MPI_Cartdim_get`

`MPI_Cart_rank`

`MPI_Cart_coords`

`MPI_Cart_sub`

# MPI\_Cart\_create Example

```
ndims = 2
dim[0] = 3 /* rows */, dim[1] = 4 /* columns */
wrap_around[0] = 0, wrap_around[1] = 0
reorder = 0
MPI_Cart_create (MPI_COMM_WORLD, ndims, dims, wrap_around,
reorder, &comm2D)
```

# `MPI_Cart_shift`

```
int MPI_Cart_shift (MPI_Comm comm, int direction, int disp, int  
*rank_source, int *rank_dest)
```

*MPI\_Cart\_create*

*MPI\_Comm\_rank*

*MPI\_Cart\_coords*

*MPI\_Cart\_shift*

# MPI\_Cart\_shift Example

```
MPI_Init(&argc, &argv);
MPI_Comm_rank (MPI_COMM_WORLD, &myrank);
MPI_Comm_size (MPI_COMM_WORLD, &size);
MPI_Comm comm2D;

dim[0] = 3 /* rows */, dim[1] = 4 /* columns */;
wrap_around[0] = 0, wrap_around[1] = 0;
reorder = 0;

MPI_Cart_create (MPI_COMM_WORLD, ndims, dim, wrap_around, reorder, &comm2D);

MPI_Comm_rank (comm2D, &newrank);
MPI_Comm_size (comm2D, &newsize);

MPI_Cart_shift (comm2D, 0, 1, &source, &dest);
printf ("Rank %d, new rank %d, source %d dest %d\n", myrank, newrank, source, dest);

MPI_Comm_free (&comm2D);
```

0	1	2	3
4	5	6	7
8	9	10	11

```
class $ mpirun -np 12 ./cart | sort -k2n
Rank 0, new rank 0, source -1 dest 4
Rank 1, new rank 1, source -1 dest 5
Rank 2, new rank 2, source -1 dest 6
Rank 3, new rank 3, source -1 dest 7
Rank 4, new rank 4, source 0 dest 8
Rank 5, new rank 5, source 1 dest 9
Rank 6, new rank 6, source 2 dest 10
Rank 7, new rank 7, source 3 dest 11
Rank 8, new rank 8, source 4 dest -1
Rank 9, new rank 9, source 5 dest -1
Rank 10, new rank 10, source 6 dest -1
Rank 11, new rank 11, source 7 dest -1
```

```
class $ mpirun -np 12 ./cart | sort -k2n
Rank 0, new rank 0, source 4 dest -1
Rank 1, new rank 1, source 5 dest -1
Rank 2, new rank 2, source 6 dest -1
Rank 3, new rank 3, source 7 dest -1
Rank 4, new rank 4, source 8 dest 0
Rank 5, new rank 5, source 9 dest 1
Rank 6, new rank 6, source 10 dest 2
Rank 7, new rank 7, source 11 dest 3
Rank 8, new rank 8, source -1 dest 4
Rank 9, new rank 9, source -1 dest 5
Rank 10, new rank 10, source -1 dest 6
Rank 11, new rank 11, source -1 dest 7
```

# MPI\_Cart\_sub

```
ndims = 2  
  
dim[0] = 3, dim[1] = 3  
  
wrap_around[0] = 0, wrap_around[1] = 0  
  
reorder = 0  
  
MPI_Cart_create (MPI_COMM_WORLD, ndims, dims, wrap_around,  
reorder, &comm2D)  
  
remains[0] = 0, remains[1] = 1 //column dimension  
  
MPI_Cart_sub (comm2D, remains, &comm1D_row)  
  
remains[0] = 1, remains[1] = 0  
  
MPI_Cart_sub (comm2D, remains, &comm1D_col)  
  
MPI_Reduce (&val, &rowmax, 1, MPI_INT, MPI_MAX, 0, comm1D_row);  
MPI_Reduce (&rowmax, &max, 1, MPI_INT, MPI_MAX, 0, comm1D_col);
```

<b>0 (0,0)</b> <b>val=78</b>	<b>1 (0,1)</b> <b>72</b>	<b>2 (0,2)</b> <b>70</b>
<b>3 (1,0)</b> <b>81</b>	<b>4 (1,1)</b> <b>87</b>	<b>5 (1,2)</b> <b>80</b>
<b>6 (2,0)</b> <b>77</b>	<b>7 (2,1)</b> <b>78</b>	<b>8 (2,2)</b> <b>75</b>