


Collective Communication

Feb 9, 2021



P2P and Collective



Collective Communication

- Must be called by all processes that are part of the communicator

Types

- Synchronization (`MPI_Barrier`)
- Global communication (`MPI_Bcast`, `MPI_Gather`, ...)
- Global reduction (`MPI_Reduce`, ..)



Barrier

- `MPI_Barrier (comm)`
- Every rank needs to call this function (for true synchronization)
- Caller returns only after all processes have entered the call

```
printf("Before barrier");
MPI_Barrier (MPI_COMM_WORLD);
printf("After barrier");
```



Barrier


n=4

```
if (myrank != 0)
    MPI_Barrier (MPI_COMM_WORLD);
printf("%d\n", rank);
```



Broadcast


- Root process sends message to all processes
- int `MPI_Bcast` (buffer, count, datatype, `root`, comm)
- “count” is the number of elements in “buffer” – must match
 - “message sizes do not match across processes in the collective routine:
Received 400 but expected 4000”
- “buffer” is input at root process
- Any process can be a root process but has to be the same process when `MPI_Bcast` is called
- Buffer size of “buf” array is not known apriori at non-root processes
- Tag?



Broadcast

Q1: Point-to-point communication for the same?

Q2: Buffer size of "buf" array is not known apriori at non-root processes, how should root broadcast buf?



Bcast Demo




"bcast.c" [dos] 31L, 623C



Gather

- Gathers values from all processes to a root process
- int `MPI_Gather` (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, `root`, comm)
- Arguments `recv*` not relevant on non-root processes
- `recvcount` → size of any (single) receive
- Distinct values (may be same) received from non-root processes at the root process
- Example: Reading a file




Q: Equivalent point-to-point communications for the same?

- `MPI_Recv` at root
- `MPI_Send` at non-root




Scatter

- Scatters values to all processes from a root process
- int MPI_Scatter (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, **root**, comm)
- Arguments send* not relevant on non-root processes
- Output parameter – recvbuf



MPI_Scatter Illustration




Credit: Shun Yan Cheung



Allgather

- All processes gather values from every other process
- int **MPI_Allgather** (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)




Alltoall

- Send data from all processes to all processes
- int `MPI_Alltoall` (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)
- Output parameter – recvbuf


Equivalent collective?

- `MPI_Scatter` at all processes



Scan

- MPI_Scan (inbuf, outbuf, count, datatype, op, comm)
- op: MIN, MAX, SUM, PROD, ...
- Perform a prefix reduction on distributed data
- Reduction of values in the send buffers of processes with ranks 0:i-1 is returned in receive buffer of rank i



DEMO

scan.c



Reduce

- `MPI_Reduce` (`inbuf`, `outbuf`, `count`, `datatype`, `op`, `root`, `comm`)
- Combines element in `inbuf` of each process
- Combined value in `outbuf` of root
- `op`: `MIN`, `MAX`, `SUM`, `PROD`, ...
- Example:

`MPI_Reduce (... , MPI_MAX, ...)` Output: 21

`MPI_Reduce (... , MPI_MIN, ...)` Output: 1

21
5
1
8
3
2
13



Allreduce

- [**MPI_Allreduce**](#) (inbuf, outbuf, count, datatype, op, comm)
- op: MIN, MAX, SUM, PROD, ...
- Combines element in inbuf of each process
- Combined value in outbuf of each process



Homework

1. Broadcast P doubles from rank 0 to all ranks (consider total #processes = 10). You can run on any number of hosts or cores. $P = 10^3, 10^4, 10^5, 10^6, 10^7$.
2. Let total number of processes be P. Write MPI_Allgather using P MPI_Bcast calls (i.e. every process broadcasts its data to the other processes). Compare the performance of both.
3. Compare performance of MPI_Alltoall with multiple MPI_Scatter for different process sizes and data sizes.

