

Collectives

Feb 12, 2021




Blocking Collectives

- MPI_Barrier
- MPI_Bcast
- MPI_Gather
- MPI_Scatter
- MPI_Allgather
- MPI_Alltoall
- MPI_Scan
- MPI_Reduce
- MPI_Allreduce



Parallel Sum



Parallel Sum

```
// local computation at every process  
for i = N/P * rank ; i < N/P * (rank+1) ; i++  
    localsum += a[i] * a[i]  
  
// collect localsum, add up at one of the ranks  
MPI_Reduce (&localsum, ... , MPI_SUM, ...)
```



Using Send/Recv

```
int recvarr[numtasks];

// receive partial sums at rank 0
stime = MPI_Wtime();
if (rank)
{
    MPI_Send(&sum, 1, MPI_INT, 0, rank, MPI_COMM_WORLD);
}
else
{
    for (int r=1; r<numtasks; r++)
    {
        MPI_Recv(&recvarr[r], 1, MPI_INT, r, r, MPI_COMM_WORLD, &status);
    }
}
etime = MPI_Wtime();
ctime = etime - stime;

stime = MPI_Wtime();
// HOMEWORK
// Add the partial sums
for (int r=1; r<numtasks; r++)
    sum += recvarr[r];
etime = MPI_Wtime();
ctime += etime - stime;

if (!rank)
printf ("%d %lf %lf\n", sum, ctime, cotime);
```



Using Reduce

```
// local sum computation
sum=0.0;
stime = MPI_Wtime();
for (i=sidx; i<sidx+N/numtasks ; i++)
    sum += a[i] * a[i];
etime = MPI_Wtime();
ctime = etime - stime;

int globalsum;

stime = MPI_Wtime();
MPI_Reduce (&sum, &globalsum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
etime = MPI_Wtime();
ctime = etime - stime;

if (!rank)
printf ("%d %lf %lf\n", globalsum, ctime, ctime);
```



Timing

```
class $ for i in `seq 1 5`; do mpirun -np 3 ./parsum 6000 ; done
10000 0.000007 0.000301
10000 0.000006 0.000309
10000 0.000010 0.000015
10000 0.000006 0.000281
10000 0.000010 0.000041
class $ for i in `seq 1 5`; do mpirun -np 3 ./parsumreduce 6000 ; done
10000 0.000011 0.000172
10000 0.000010 0.000034
10000 0.000010 0.000033
10000 0.000010 0.000025
10000 0.000010 0.000028
class $ for i in `seq 1 5`; do mpirun -np 30 -hosts csews3:10,csews5:10,csews6:10 ./parsum 6000000 ; done
1711000000 0.001006 0.025789
1711000000 0.001034 0.010492
1711000000 0.001007 0.023931
1711000000 0.001029 0.038722
1711000000 0.001028 0.024971
class $ for i in `seq 1 5`; do mpirun -np 30 -hosts csews3:10,csews5:10,csews6:10 ./parsumreduce 6000000 ; done
1711000000 0.001003 0.011949
1711000000 0.002717 0.001540
1711000000 0.001092 0.009885
1711000000 0.003251 0.003160
1711000000 0.001091 0.012467
```



Vector Variants

- Communicate unequal amount of data to/from each process involved in the collective function call

2B		
2B		
2B		

2B		
3B		
4B		




Gatherv

- Root gathers different amounts of data from the other processes
- int **MPI_Gatherv** (sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, root, comm)
- recvcounts – Number of elements to be received from each process
- displs – Displacement at which to place received data

MPI_Recv (recvbuf+displs[i], recvcounts[i], recvtype, i, i, comm, &status) at root for i^{th} process

MPI_Send at non-root



MPI_Gatherv Code

```
int message[arrSize];
int countArray[numtasks], displArray[numtasks];

int displ = 0; // note that root process is 0 here

// this information is needed by the root
if (!rank)
    for (i = 0; i < numtasks; i++) {
        countArray[i] = arrSize*(i+1); // depends on the counts, root may need to get it from the processes
        displArray[i] = displ;
        displ += countArray[i];
        printf ("%d %d %d\n", i, countArray[i], displArray[i]);
    }

if (!rank)
    printf ("\n");

// every process initializes their local array
srand(time(NULL));
for (i = 0; i < arrSize; i++) {
    message[i] = i; // (double)rand() / (double)RAND_MAX;
}

int recvMessage[displ]; // significant at the root process

// receive different counts of elements from different processes
MPI_Gatherv (message, arrSize, MPI_INT, recvMessage, countArray, displArray, MPI_INT, 0, MPI_COMM_WORLD);

if (!rank)
    for (i = 0; i < displ; i++) {
        printf ("%d %d\n", i, recvMessage[i]);
    }
```



Demo

```
class $ mpirun -np 2 ./gatherv 2
0 2 0
1 4 2

0 0
1 1
2 0
3 1
4 2
5 3
class $ mpirun -np 3 ./gatherv 2
0 2 0
1 4 2
2 6 6

0 0
1 1
2 0
3 1
4 2
5 3
6 0
7 1
8 2
9 3
10 4
11 5
class $
```



Bug

```
Fatal error in PMPI_Gatherv: Message truncated, error stack:  
PMPI_Gatherv(435)..... MPI_Gatherv failed(sbuf=0x7ffd379c4250, scount=2, MPI_INT, rbuf=0x7ffd379c4210, rcnts=0x7ffd379c4240,  
displs=0x7ffd379c4230, MPI_INT, root=0, MPI_COMM_WORLD) failed  
MPIR_GathervImpl(235).....  
MPIR_Gatherv(151).....  
MPIDI_CH3_PktHandler_EagerShortSend(363): Message from rank 1 and tag 4 truncated; 16 bytes received but buffer size is 8  
class $ vi gatherv.c  
class $ !mpicc  
mpicc -o gatherv gatherv.c  
class $ !mpirun  
mpirun -np 2 ./gatherv 2  
0 2 0  
1 4 2  
class $ █
```




Example

```
class $ time mpirun -np 3 ./gatherv 2
real    0m0.007s
user    0m0.004s
sys     0m0.010s
class $ time mpirun -np 30 -hosts csews4:10,csews2:10,csews3:10 ./gatherv 200
real    0m0.695s
user    0m0.026s
sys     0m0.006s
class $ time mpirun -np 30 -hosts csews4:10,csews2:10,csews3:10 ./gatherv 2000
real    0m1.857s
user    0m0.033s
sys     0m0.002s
class $ time mpirun -np 60 -hosts csews4:10,csews2:10,csews3:10,csews5:10,csews6:10,csews7:10 ./gatherv 20000
real    0m2.965s
user    0m12.571s
sys     0m1.559s
class $ time mpirun -np 60 -hosts csews4:10,csews2:10,csews3:10,csews5:10,csews6:10,csews7:10 ./gatherv 800000
real    1m5.621s
user    7m2.519s
sys     1m2.687s
class $ █
```




MPI_Scatterv

- Root scatters different amounts of data to the other processes
- int **MPI_Scatterv** (const void *sendbuf, const int *sendcounts, const int *displs, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)
- sendcounts – Number of elements to be sent to each process
- displs – Displacement (relative to sendbuf) at which the data to be sent resides




Allgatherv

- All processes gather values of different lengths from all processes
- int **MPI_Allgatherv** (sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, comm)
- recvcounts – Number of elements to be received from each process
- displs – Displacement at which to place received data




Alltoallv

- Every process sends data of different lengths to other processes
- int **MPI_Alltoallv** (sendbuf, sendcount, sdispls, sendtype, recvbuf, recvcount, rdispls, recvtype, comm)
- Output parameter – recvbuf
- It's not necessary to receive some data from all processes, i.e. some entries of count and displs may be 0



Practical Use



Non-blocking Collectives

- Introduced in MPI-3
- Benefit of non-blocking point-to-point
- Overlap communication and computation
- Reduce synchronization
- Improve performance for overlapping communicators
- How do we ensure completion?
 - `MPI_Wait` (request, status)



Non-blocking Collectives

- MPI_Ibcast (buffer, count, datatype, root, comm, request)
- MPI_Igather (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm, request)
- MPI_Igatherv (sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, root, comm, request)
- MPI_Ialltoall (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm, request)
- ...



MPI_Ibcast

```
int main( int argc, char *argv[] )
{
    int count = atoi(argv[1]);
    int myrank, buf[count];
    double buf_can_modify[count];
    MPI_Status status;
    MPI_Request request;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );

    // initialize data
    for (int i=0; i<count; i++)
        buf[i] = myrank + i*i;

    // has to be called by all processes
    MPI_Ibcast(buf, count, MPI_INT, 1, MPI_COMM_WORLD, &request);

    // some random expensive computations
    for (int j=0; j<count; j++)
        for (int i=0; i<count; i++)
            buf_can_modify[i] = i*12022021/(myrank+1)*pow(2,i);

    MPI_Wait (&request, &status);

    MPI_Finalize();
    return 0;
}
```



Execution Times

```
class $ for i in `seq 1 5` ; do time mpirun -np 40 -f ./hostfile ./ibcast 100000000 ; done  
  
real    0m14.848s  
user    2m27.926s  
sys     0m5.847s  
  
real    0m14.547s  
user    2m25.309s  
sys     0m6.080s  
  
real    0m13.946s  
user    2m19.674s  
sys     0m5.541s  
  
real    0m14.135s  
user    2m20.602s  
sys     0m5.846s  
  
real    0m13.508s  
user    2m17.507s  
sys     0m5.713s  
  
real    0m16.064s  
user    1m50.779s  
sys     0m5.527s  
  
real    0m20.102s  
user    2m26.007s  
sys     0m6.418s  
  
real    0m20.122s  
user    2m25.546s  
sys     0m6.555s  
  
real    0m20.256s  
user    2m26.897s  
sys     0m6.586s  
  
real    0m18.734s  
user    2m13.618s  
sys     0m7.214s
```

MPICH_ASYNC_PROGRESS=1



MPI Examples

```
If (myrank == 0) {  
    MPI_Ibcast (&buf1, 1, MPI_INT, 0, comm, req1);  
    MPI_Wait (&req1, MPI_STATUS_IGNORE);  
    MPI_Send (buf, count, MPI_INT, 1, tag, comm);  
}  
  
If (myrank == 1) {  
    MPI_Ibcast (&buf1, 1, MPI_INT, 0, comm, req1);  
    MPI_Recv (buf, count, MPI_INT, 0, tag, comm, MPI_STATUS_IGNORE);  
    MPI_Wait (&req1, MPI_STATUS_IGNORE);  
}
```

Will it run for n=2?

Valid code



MPI Examples

```
If (myrank == 0) {  
    MPI_Ibcast (&buf1, 1, MPI_INT, 0, comm, req1[0]);  
    MPI_Ibcast (&buf2, 2, MPI_INT, 0, comm, req1[1]);  
}  
  
If (myrank == 1) {  
    MPI_Ibcast (&buf2, 1, MPI_INT, 0, comm, req1[0]);  
    MPI_Ibcast (&buf1, 2, MPI_INT, 0, comm, req1[1]);  
}  
MPI_Waitall(2, req1, MPI_STATUSES_IGNORE);
```

Correct for n=2?

Valid code

