

Collectives Algorithms

Mar 5, 2021



Communication Cost Model (For now)

- Message transfer time is modeled as $L+n/B$, where L is the latency (or startup time) per message, and $1/B$ is the transfer time per byte, and n the message size in bytes
- All processes can send and receive one message at the same time



Optimization of Collective Communication Operations in MPICH

Rajeev Thakur Rolf Rabenseifner William Gropp

IJHPCA 2005




Broadcast – Naïve/Linear Algorithm

- Root process sends to every other process

Cons

- Root is a bottleneck
- Idling processes
- Communication links are under-utilized



- #Steps for p processes?
 - $p - 1$
- Transfer time for n bytes
 - $T(p) = (p-1) * (L + n/B)$




Broadcast – Binary Tree

- Root process sends to its children in 2 steps
- Every node sends to its children in 2 steps (left and right)


Cons

- Idling processes
- #Steps for p processes?
 - $2 * \log p$
- Transfer time for n bytes
 - $T(p) = 2 * \log p * (L + n/B)$




Broadcast – Pipelined Binary Tree


- Almost similar to binary
- Multiple rounds and smaller message size per round
- #Steps for p processes?
 - $2 * (\log p + R - 1)$
- Transfer time for n bytes
 - $T(p) = 2 * (\log p + R - 1) * (L + (n/R)*1/B)$



Broadcast – Binomial Tree




- #Steps for $p (=2^d)$ processes?
 - $\log p$
- Transfer time for n bytes
 - $T(p) = \log p * (L + n/B)$
 - $T(p^2) = 2 \log p * (L + n/B)$




Broadcast

Q: Which interconnect would most likely exhibit minimum link contention for binomial tree broadcast algorithm?




Q: Equivalent collective?




Scatter

- Scatters values to all processes from a root process
- int MPI_Scatter (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, **root**, comm)
- Arguments send* not relevant on non-root processes
- Output parameter – recvbuf



Scatter – Recursive Halving



Every step the message size halves: $n/2, n/2^2, \dots, n/2^{(\log p)}$

Time for scatter of n bytes from root


$$\log p * L + (p-1)*(n/p)*(1/B)$$



Allgather


- All processes gather values from all processes
- `int MPI_Allgather (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`
- No root process

```
recvbuf = (int *) malloc  
(size*100*sizeof(int));  
  
MPI_Allgather (sendbuf, 100, MPI_INT,  
recvbuf, 100, MPI_INT, comm);
```



Allgather – Ring Algorithm

- Every process sends to and receives from everyone else
- Assume p processes and total n bytes
- Every process sends and receives n/p bytes
- Time
 - $(p - 1) * (L + n/p * (1/B))$
- How can we improve?




Allgather – Recursive Doubling

- Every process sends and receives $(2^{k-1})^* n/p$ bytes at k^{th} step
- $(\log p) * L + (p-1)*n/p*(1/B)$




Bcast – Van de Geijn et al.

- Message is first scattered from the root
- Scattered data is collected at all processes



Bcast – Time Analysis

- Time for broadcasting n bytes from root (binomial tree)
 - $\log p * (L + n/B)$
 - Latency term: $\log p$
 - Bandwidth term: $\log p$
- Time for scatter of n bytes from root
 - $\log p * L + (p-1)*(n/p)*(1/B)$
- Time for allgather (ring) of n/p bytes
 - $(p-1) * L + (p-1)*(n/p)*(1/B)$
- Time for broadcast of n bytes using scatter and allgather
 - $(\log p + p-1) * L + 2((p-1)/p)*(n/B)$




Broadcast Algorithms in MPICH

- Short messages
 - < `MPIR_CVAR_BCAST_SHORT_MSG_SIZE`
 - Binomial
- Medium messages
 - Scatter + Allgather (Recursive doubling)
- Large messages
 - > `MPIR_CVAR_BCAST_LONG_MSG_SIZE`
 - Scatter + Allgather (Ring)



Old vs. New MPI_Bcast



Allgather – Bruck Algorithm

P0	P1	P2	P3	P4	P5
0	1	2	3	4	5

Initial data

P0	P1	P2	P3	P4	P5
0	1	2	3	4	5
1	2	3	4	5	0

After step 0

P0	P1	P2	P3	P4	P5
0	1	2	3	4	5
1	2	3	4	5	0
2	3	4	5	0	1
3	4	5	0	1	2
4	5	0	1	2	3
5	0	1	2	3	4

After step 1

P0	P1	P2	P3	P4	P5
0	1	2	3	4	5
1	2	3	4	5	0
2	3	4	5	0	1
3	4	5	0	1	2
4	5	0	1	2	3
5	0	1	2	3	4

After step 2

P0	P1	P2	P3	P4	P5
0	0	0	0	0	0
1	1	1	1	1	1
2	2	2	2	2	2
3	3	3	3	3	3
4	4	4	4	4	4
5	5	5	5	5	5

After local shift


Shift up by i blocks

Process i
Send $i - 2^k$
Recv $i + 2^k$

Non-power-of-2
Send top $p - 2^{\lfloor \lg p \rfloor}$



Recursive Doubling vs. Bruck




MPI_Allgather in MPICH

- Bruck algorithm
 - $\text{ceil}(\log p) * L + (p-1)*n/p*(1/B)$
 - short messages (< 80 KB) and non-power-of-two numbers of processes
- Recursive doubling
 - $(\log p) * L + (p-1)*n/p*(1/B)$
 - power-of-two numbers of processes and short or medium-sized messages (< 512 KB)
- Ring algorithm
 - $(p-1) * L + (p-1)*n/p*(1/B)$
 - long messages and any number of processes
 - medium-sized messages and non-power-of-two numbers of processes


Why does ring perform better?



Performance Comparison



Performance Comparison




allgather.c

```
tot_bytes = (MPI_Aint) recvcount *comm_size * type_size;
if ((tot_bytes < MPIR_CVAR_ALLGATHER_LONG_MSG_SIZE) && !(comm_size & (comm_size - 1))) {
    mpi_errno =
        MPIR_Allgather_intra_recursive_doubling(sendbuf, sendcount, sendtype, recvbuf,
                                                recvcount, recvtype, comm_ptr, errflag);
} else if (tot_bytes < MPIR_CVAR_ALLGATHER_SHORT_MSG_SIZE) {
    mpi_errno =
        MPIR_Allgather_intra_brucks(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype,
                                    comm_ptr, errflag);
} else {
    mpi_errno =
        MPIR_Allgather_intra_ring(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype,
                                comm_ptr, errflag);
}
```



Adaptive/Dynamic Algorithm Selection



Old vs. new MPI_Allgather times

