

Process Mapping

Mar 13, 2021



Topology-aware task mapping for reducing communication contention on large parallel machines

Tarun Agarwal, Amit Sharma, Laxmikant V. Kale

IPDPS 2006



Effect of Mapping

- $16 \times 16 \times 16$ 3D torus
 - Diameter: 24
 - Average load on links high when #hops large

Message Size	Random Mapping	Optimal Mapping
1KB	56.93ms	46.91ms
10KB	243.64ms	124.56ms
100KB	2247.75ms	914.72ms
500KB	11.62s	4.44s
1MB	23.50s	8.80s

Table 1: Time for 200 iterations of a Jacobi-like program with optimal mapping and random mapping



Network Graphs

- Topology graph
- Task graph
 - c_{ab} → amount of communication (bytes) on e_{ab} between v_a and v_b
- Map: $v_t \in V_t$ placed on $v_p \in V_p$

$$G_p = (V_p, E_p)$$

$$G_t = (V_t, E_t)$$

$$P : V_t \longrightarrow V_p$$



Hop-bytes

Total size of inter-processor communication in bytes weighted by distance between the respective end-processors.

The overall hop-bytes is the sum of hop-bytes due to individual nodes in the task graph.

$$HB(G_t, G_p, P) = \frac{1}{2} \sum_{v_a \in V_t} HB(v_a)$$

$$\text{where } HB(v_a) = \sum_{e_{ab} \in E_t} HB(e_{ab})$$

$$\text{where } HB(e_{ab}) = c_{ab} \times d_p(P(v_a), P(v_b))$$



Hops per byte


Average number of network links a byte has to travel under a task mapping

$$\text{Hops per Byte} = \frac{HB(G_t, G_p, P)}{\sum_{e_{ab} \in E_t} c_{ab}}$$

$$\text{Hops per Byte} = \frac{\sum_{e_{ab} \in E_t} c_{ab} \times d_p(P(v_a), P(v_b))}{\sum_{e_{ab} \in E_t} c_{ab}}$$




Process Mapping



8 x 2 2D virtual process topology

Hop-bytes: bytes weighed by distance

Hops per byte: average links traversed by a byte




4 x 4 2D mesh



Process Mapping

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31


8 x 4 2D virtual process topology



Topology-aware Process Mapping

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31

8 x 4 2D virtual process topology



This Paper

A process mapping strategy that minimizes the impact of topology by heuristically minimizing the total number of hop-bytes communicated



Two Phases

- Partitioning
 - Partitioning compute objects into p groups (assume p processors)
 - METIS*
- Mapping
 - Map the p groups to p processors such that the heavily communicating groups are placed on nearby processors
 - Where should the next process be placed?

**"A fast and high quality multilevel scheme for partitioning irregular graphs", George Karypis and Vipin Kumar, International Conference on Parallel Processing (ICPP), pp. 113-122, 1995



Estimation Function

- $f_{est}(t, p, M)$ = Cost of estimating the placement of a task t onto processor p under current task mapping M
- Estimate how critical it is to place a task in the current cycle, select the task with maximum criticality
- T_k is the set of tasks yet to be placed
- P_k is the set of processors that are available

$$T_k \cup \bar{T}_k = \emptyset$$
$$P_k \cup \bar{P}_k = \emptyset$$



Estimation Function

- $f_{est}(t, p, M)$ = Cost of estimating the placement of a task t onto processor p under current task mapping M

- First order approximation w.r.t placed tasks

- Hop-bytes

$$f_{est}(t_i, p, M) = \sum_{t_j \in \bar{T}_k} c_{ij} d_p(p, M(t_j))$$

- Second order approximation w.r.t all tasks

$$d_p(p, M(t_j)) \approx \frac{\sum_{p_j \in V_p} d_p(p, p_j)}{|V_p|}$$




Mapping Heuristic

- For each task, find the best processor, the one where it costs least to place it.
- For each task, the estimation function gives the cost of placing a task on its best processor and the expected cost when placed on an arbitrary processor.
- Determine how critical it is to place the task in the current cycle and select the most critical task for placement in the current cycle.
- The estimation function f (using hop-bytes) is maintained in a matrix $p \times p$.
- Select the task that maximizes $f_{avg}[t] - fm_{in}[t]$.
- Find the processor that minimizes f .




Results – LeanMD on 2D Torus



34% reduction
in hops-per-byte



Results – LeanMD on 3D Torus



40% reduction
in hops-per-byte

