

Parallelization

Mar 16, 2021



Parallelization Steps

1. *Decomposition* of computation into tasks

- Identifying portions of the work that can be performed concurrently

2. *Assignment* of tasks to processes

- Assigning concurrent pieces of work onto multiple processes running in parallel

3. *Orchestration* of data access, communication and synchronization among processes


- Distributing the data associated with the program
- Managing access to data shared by multiple processes
- Synchronizing at various stages of the parallel program execution

4. *Mapping* of processes to processors

- Placement of processes in the physical processor topology



Illustration of Parallelization Steps



Q: What if number of processes != available processors?

Source: Culler et al.




Performance Goals

- Expose concurrency
- Reduce inter-process communications
- Load-balance
- Reduce synchronization
- Reduce idling
- Reduce management overhead
- Preserve data locality
- Exploit network topology



Matrix Vector Multiplication – Decomposition



P = 9 ?

P = 3 ?


Decomposition

Identifying portions of the work that can be performed concurrently

Assignment



Matrix Vector Multiplication – Orchestration



- What is the initial communication step?
- Ways to distribute vector (if its present in 1 process)?
- What are the differences between distribution and parallel reads?



Distribute using Bcast vs. Allgather

```
MPI_Init( &argc, &argv );
MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
MPI_Comm_size( MPI_COMM_WORLD, &commsize );

time = MPI_Wtime();
MPI_Bcast (buf, N, MPI_FLOAT, 0, MPI_COMM_WORLD);
etime = MPI_Wtime() - time;
MPI_Reduce (&etime, &maxtime, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
if (!myrank) printf ("Time to bcast: %11.3lf\n", maxtime);

int count = N/commsize;
time = MPI_Wtime();
MPI_Allgather (buf, count, MPI_FLOAT, recvbuf, count, MPI_FLOAT, MPI_COMM_WORLD);
etime = MPI_Wtime() - time;
MPI_Reduce (&etime, &maxtime, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
if (!myrank) printf ("Time to allgather: %7.3lf\n", maxtime);
```



Bcast vs. Allgather

```
class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 10000; echo ; done
Time to bcast: 0.014
Time to allgather: 0.021

Time to bcast: 0.018
Time to allgather: 0.009

Time to bcast: 0.012
Time to allgather: 0.007

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 100000; echo ; done
Time to bcast: 0.034
Time to allgather: 0.011

Time to bcast: 0.027
Time to allgather: 0.023

Time to bcast: 0.026
Time to allgather: 0.011

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 1000000; echo ; done
Time to bcast: 0.187
Time to allgather: 0.347

Time to bcast: 0.176
Time to allgather: 0.111

Time to bcast: 0.155
Time to allgather: 0.112

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 10000000; echo ; done
Time to bcast: 1.421
Time to allgather: 1.121

Time to bcast: 1.618
Time to allgather: 1.282

Time to bcast: 1.674
Time to allgather: 1.583


class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 100000000; echo ; done
Time to bcast: 18.061
Time to allgather: 15.616

Time to bcast: 23.447
Time to allgather: 17.005

Time to bcast: 16.875
Time to allgather: 11.085
```



Matrix Vector Multiplication – Decomposition




$P = 3$

Decomposition
Assignment
Orchestration


What is the advantage of column-wise partitioning ?



Matrix Vector Multiplication – Mapping



Stencils



Five-point stencil



Virtual Topology


0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31

8 x 4 2D virtual process topology

- Communication pattern of MPI processes
 - Graphical representation of communications
 - Nearest neighbor in a mesh
 - All-to-all
 - ...
 - Convenient way to represent communications
- Note: Virtual topology is set up before execution




Physical Topology



- Connections between allocated nodes
- Mapping: Placement of ranks onto cores





Physical Topology



- Connections between allocated cores
- Mapping: Placement of ranks onto cores
- Default placement of ranks based on node IDs




Rank placement




- May place ranks anywhere in the allocated nodes based on the communication pattern
- **Topology-aware mapping:** Mapping that minimizes all communication times taking into account the physical topology



Process-to-processor Mapping




4 x 4 x 2 3D torus



8 x 4 2D virtual process topology



Topology-aware Mapping




4 x 4 x 2 3D torus

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31

8 x 4 2D virtual process topology



1D Domain decomposition



Communications?

2 sends()
2 recvs()

N grid points
P processes
N/P points per process

#Communications?


2

#Computations?
N/P

Communication to computation ratio=?



1D Domain decomposition



Q: One drawback?


N grid points
P processes
N/P points per process

#Communications?
2vN (assuming square grid)
#Computations?
N/P (assuming square grid)

Communication to computation ratio=?




2D Domain decomposition



- + Several parallel communications
- + Lower communication volume/process



2D Domain decomposition



#Communications?

2 Sends()
2 Recvs()

#Communications?

$2\sqrt{N}/\sqrt{P}$ (assuming square grid)


#Computations?

N/P (assuming square grid)

Communication to computation ratio=?



2D Domain decomposition



#Communications?

4 Sends()
4 Recvs()

#Communications?

$4\sqrt{N}/\sqrt{P}$ (assuming square grid)


#Computations?

N/P (assuming square grid)

Communication to computation ratio=?



Communication to Computation Ratio (1D)




#Communications = 12

#Computations = 12

Communication to
computation ratio = 1



Communication to Computation Ratio (2D)



#Communications = 6

#Computations = 6

Communication to
computation ratio = 1

