


Parallel I/O - II

Apr 9, 2021




Data Sieving



- Make large I/O requests and extract the data that is really needed
- Huge benefit of reading large, contiguous chunks



Data Sieving – Interleaved data



Q: What is the problem here?

Solution – Lock the relevant portions in the file



User-controlled MPI-IO Parameters

- `ind_rd_buffer_size` - Buffer size for data sieving for read
- `ind_wr_buffer_size` - Buffer size for data sieving for write
- `romio_ds_read` - Enable or not data sieving for read
- `romio_ds_write` - Enable or not data sieving for write



MPI_Info – Example

```
MPI_Info_create (&info);
```

```
MPI_Info_set (info, "ind_rd_buffer_size", "2097152");
```

```
MPI_Info_set (info, "ind_wr_buffer_size", "1048576");
```

```
MPI_File_open (MPI_COMM_WORLD, filename, amode, info, &fh);
```



Independent I/O – Summary

- Individual file pointers
 - Explicit offsets or not
 - Shared file pointers
 - Read/write starting from the current location of file pointer
 - All processes share the same file view
- [`MPI_File_read_at/MPI_File_read`](#)
- [`MPI_File_read_shared/MPI_File_write_shared`](#)




Parallel I/O Classification


- Independent I/O (we saw till now)
- Collective I/O (we will see next)



Non-contiguous Accesses



Multiple Non-contiguous Accesses



- Every process' local array is non-contiguous in file
- Every process needs to make small I/O requests
- Can these requests be merged?



MPI Collective I/O

```
MPI_File_open (MPI_COMM_WORLD, "/scratch/largefile", MPI_MODE_RDONLY,  
MPI_INFO_NULL, &fh)
```

```
MPI_File_read_at_all (fh, offset, buffer, count, MPI_INT, status)
```

```
MPI_File_read_all (fh, buffer, count, MPI_INT, status)
```

```
MPI_File_close (&fh)
```

```
class mpirun -np 10 ./2.indepIOnoncontig 100000  
time diff 0.089348 1.067375 MB/s  
class mpirun -np 10 ./2.indepIOnoncontig 100000  
time diff 0.147186 0.647937 MB/s  
class mpirun -np 10 ./2.indepIOnoncontig 100000  
time diff 0.535093 0.178226 MB/s  
class  
class  
class mpirun -np 10 ./3.colIOnoncontig 100000  
time diff 0.005033 18.947468 MB/s  
class mpirun -np 10 ./3.colIOnoncontig 100000  
time diff 0.002475 38.531933 MB/s  
class mpirun -np 10 ./3.colIOnoncontig 100000  
time diff 0.002582 36.931031 MB/s
```



Two-phase I/O


Entire access pattern must be known before making file accesses

- Phase 1
 - Processes request for a single large contiguous chunk
 - Reduced file I/O cost due to large accesses
- Phase 2
 - Processes redistribute data among themselves
 - Additional inter-process communications




Two-phase I/O

- Phase 1
 - Processes analyze their own I/O requests
 - Create list of offsets and list of lengths
 - Everyone broadcasts start offset and end offset to others
 - Each process reads its own *file domain*




Two-phase I/O

- Phase 2
 - Processes analyze the file domains
 - Processes exchange data with the corresponding process





File domain – Example



P0 and P1 exchange, P2 and P3 exchange




File domain – Example



Everyone needs data from every other process




Collective I/O



Communication may become bottleneck?



Collective I/O Aggregators



- Multiple small non-contiguous I/O requests from different processes are combined
- A subset of processes, I/O aggregators, access their file domains (I/O phase)
- Data redistributed among all processes (communication phase)
- **Cons - Synchronization**



Aggregators

- Too few aggregators
 - Large buffer size required per aggregator and multiple I/O iterations
 - Underutilization of the full bandwidth of the storage system
- Too many aggregators
 - Request for large number of small chunks → suboptimal file system performance
 - Increased cost of data exchange operations

In MPICH

- Buffer size in aggregators = 16 MB
- Default number of aggregators – #unique hosts which open the file
- Placement – Specific to file system
 - `mpich/src/mpi/romio/adio/ad_gpfs/pe/ad_pe_aggrs.c` (GPFS)



I/O Aggregators – Limited buffer

Total number of processes = 1024

Let each process read 2^{20} doubles (= 1 MB)

Total number of aggregators = 16

Temporary buffer in each aggregator process = 4 MB

- Collective I/O may be done in several iterations
- Double buffering may help

$$\text{I/O data size per aggregator (D)} = \frac{1024 * 1}{16} \text{ MB}$$

$$\text{Number of times each aggregator needs to do the I/O} = \frac{D}{4} = 16$$



User-controlled Parameters

- Number of aggregators
- Buffer size in aggregators



User-controlled Parameters

- Number of aggregators (cb_nodes, default=1 per node)
- Placement of aggregators (cb_config_list)
- Buffer size in aggregators (cb_buffer_size, default=16MB)
- ...

- Can be set via hints (ROMIO_HINTS file)
- MPI_Info object is used to pass hints



Set Hints via MPI_Info

```
MPI_Info_create(&info);
```

```
MPI_Info_set(info, "cb_nodes", "8");
```

```
MPI_File_open(MPI_COMM_WORLD, filename, amode, info, &fh);
```



Set Hints via Environment Variable

- `export ROMIO_HINTS=$PWD/romio_hints`

```
romio_cb_read enable
romio_cb_write enable
ind_rd_buffer_size 4194304
ind_wr_buffer_size 1048576
cb_config_list=*:1
```



MPIIO Hints

```
MPI_File_open (MPI_COMM_WORLD, "/pfs/datafile", MPI_MODE_CREATE |  
MPI_MODE_RDWR, MPI_INFO_NULL, &fh);  
  
MPI_File_get_info (fh, &info_used);  
  
MPI_Info_get_nkeys (info_used, &nkeys);  
  
for (i=0; i<nkeys; i++) {  
    MPI_Info_get_nthkey (info_used, i, key);  
    MPI_Info_get (info_used, key, MPI_MAX_INFO_VAL, value, &flag);  
    printf ("Process %d, Default: key = %s, value = %s\n", rank, key, value);  
}
```

export ROMIO_PRINT_HINTS=1



Performance

```
$ mpirun -np 6 -hosts csews2:2,csews20:2,csews30:2 ./collOnoncontig 10485760  
key = cb_buffer_size          value = 16777216  
key = romio_cb_read           value = enable  
key = romio_cb_write          value = enable  
key = cb_nodes                value = 3  
36.014944 MB/s  
$ mpirun -np 6 -hosts csews2:3,csews20:3 ./collOnoncontig 10485760  
key = cb_buffer_size          value = 16777216  
key = romio_cb_read           value = enable  
key = romio_cb_write          value = enable  
key = cb_nodes                value = 2  
27.474843 MB/s
```




Performance

```
$ mpirun -np 6 -hosts csews2:2,csews20:2,csews30:2 ./collOnoncontig 10485760  
key = cb_buffer_size          value = 16777216  
key = romio_cb_read           value = enable  
key = romio_cb_write          value = disable  
key = cb_nodes                value = 3  
5.106368 MB/s
```

```
$ mpirun -np 6 -hosts csews2:3,csews20:3 ./collOnoncontig 10485760  
key = cb_buffer_size          value = 16777216  
key = romio_cb_read           value = enable  
key = romio_cb_write          value = disable  
key = cb_nodes                value = 2  
6.028663 MB/s
```



I/O Aggregators



- Data sieving
- Independent I/O may be called if there is no benefit of collective I/O



Recap

Independent I/O

- Individual file pointers `MPI_File_read`
- Explicit offsets `MPI_File_read_at`
- Shared file pointers `MPI_File_read_shared`

Collective I/O

- Individual file pointers `MPI_File_read_all`
- Explicit offsets `MPI_File_read_at_all`
- Shared file pointers `MPI_File_read_ordered`



Non-blocking I/O

```
MPI_Request request;  
MPI_File_iwrite_at (fh, offset, buf, count, datatype, &request);  
MPI_File_iwrite_at_all (fh, buf, count, datatype, &request);  
...  
/* computation */  
  
MPI_Wait (&request, &status);
```



Parallel I/O approaches

- Shared file
 - Independent I/O
 - Collective I/O
- File per process
- File per group of processes

- Challenging for analysis and visualization codes
- Restriction on #processes while restarting

File system locking overhead is high

Numerous files, large overhead

Locking overhead nil



Aggregator Selection

How many aggregators, and their placements?

- Depends on the architecture, file system, data size, access pattern
- Depends on the network topology, number of nodes and their placements, etc.

