


Parallel I/O – III

Apr 20, 2021



Recap – Parallel I/O

Uniform access pattern



Each process reads a big chunk of data (one-fourth of the file) from a **common** file

Independent I/O

```
MPI_File_set_view (fh, rank*file_size_per_proc, ...)  
MPI_File_read (fh, data, datacount, MPI_INT, status)
```

fh: individual
file pointer




Recap – File View

```
MPI_File_set_view (           file handle  
    fh,                      displacement (bytes to be skipped from the start of the file)  
    rank*file_size_per_proc,  etype (any basic or derive MPI datatype)  
    MPI_INT,                 filetype (constructed out of etype, specifies visible portions)  
    MPI_INT,                 data representation (native → same as memory)  
    "native",                file info hints  
    info  
)
```



Using File Views




Each process reads a **small** chunk of data from a common file

```
MPI_File_set_view (fh, displacement, etype, filetype, "native", info)  
MPI_File_read_all (fh, data, datacount, MPI_INT, status)
```



Discontiguous Access Pattern



Each process reads **small discontiguous** chunks of data from a common file



Collective I/O – Selecting Aggregators


Mohamad Chaarawi and Edgar Gabriel, Automatically Selecting the Number of Aggregators for Collective I/O Operations, IEEE Cluster 2011

Three approaches for collective I/O aggregation

- Two-phase
 - Shuffle phase
 - I/O phase
- Dynamic segmentation
 - Fixed #processes per aggregator
- Static segmentation
 - Extension of dynamic segmentation
 - Fixed #bytes per cycle per process per aggregator



Collective I/O – Selecting Aggregators



Source: Chaarawi et al., Automatically Selecting the Number of Aggregators for Collective I/O Operations, Cluster'11



Collective I/O – Selecting Aggregators

Select

- Data size k (write saturation point)
 - Using a benchmark
 - Determined by network topology, I/O network, parallel file system, etc.
- Initial #aggregators
- Refine #aggregators
 - If total bytes per group $> k$, then split
 - If total bytes per group $< k$, then join


Group 1	0	1	2	3
Group 2	4	5	6	7
Group 3	8	9	10	11
Group 4	12	13	14	15

Group 1	0	1	2	3	Group 2
Group 3	4	5	6	7	Group 4
Group 5	8	9	10	11	Group 6
Group 7	12	13	14	15	Group 8

Group 1	0	1	2	3
	4	5	6	7
Group 2	8	9	10	11
	12	13	14	15



Tuning for Bandwidth Saturation




Shark

- 29 nodes
- SDR Infiniband and Gigabit Ethernet
- PVFS2



Tuning for Bandwidth Saturation



Deimos




- 724 nodes
- 11 I/O servers
- via SDR
- Inifiniband
- Lustre



Dynamic vs. Static vs. Two-phase on Shark

The number of aggregators chosen is equal to the total number of processes, because each process is requesting to write data larger than the value of k per function call.

64 B, 1 KB, 1 MB





MPI Tile IO benchmark with 81 (9 X 9) processes




Dynamic vs. Static vs. Two-phase on Deimos

Size of a segment




- The number of aggregators chosen was good enough for small size, but communication costs dominated the cost of large segment size.
- 2 MB being a multiple of stripe size did not affect performance



Latency IO benchmark with 64 processes



I/O Stack




High Data Throughput

How?

- I/O forwarding from compute to I/O nodes
- Multiple I/O servers to manage the data storage
- A large file may be striped across several disks



I/O Forwarding




Source: Ohta et al., “Optimization Techniques at the I/O Forwarding Layer”

- I/O requests forwarded to dedicated I/O nodes by compute node kernels
- I/O nodes redirect I/O requests to the backend parallel file systems
- Reduces the number of clients accessing the file systems
- Can reduce the file system traffic by aggregating and reordering I/O requests
- I/O forwarding scheduler can exploit the global view of parallel applications to sort and merge I/O requests more effectively



IBM Blue Gene/P


- Compute nodes are partitioned into subsets that map to an I/O node
- Compute node kernel forwards all I/O and socket requests to the I/O node
- A dedicated control and I/O daemon running on the I/O node performs I/O on behalf of the compute nodes
- I/O forwarding can potentially reduce the file system traffic by aggregating, caching the I/O requests at the I/O nodes



Source: Ali et al., “Scalable I/O Forwarding Framework for High-Performance Computing Systems”



BG/Q – I/O Node Architecture



Compute node rack

1024 compute nodes

16 bridge nodes

I/O nodes

2 bridge nodes

connect to 1 I/O node

Q: Where should the aggregators be placed?




Lustre File System

- Parallel file system
- Used in 15/30 top 500 supercomputers
- POSIX-compliant file system
 - presents a unified file system interface to the user
- Object-based filesystem
 - “A storage object is a logical collection of bytes on a storage device”¹
 - Composed of data, attributes, metadata
 - Files distributed across multiple objects
- Scalability due to object storage and division of labor
- No file server bottleneck

¹ Mesnier et al., Object-Based Storage, IEEE Communications Magazine, 2003



Lustre Scalable Storage Architecture



“Lustre can deliver more than a terabyte-per-second of combined throughput.” --
<http://wiki.lustre.org/images/6/64/LustreArchitecture-v4.pdf>




Lustre

Three components

- Metadata servers (MDS)
- Object storage servers (OSS)
 - Object storage targets (OST)
- Clients

Lustre clients access and concurrently use data through the standard POSIX I/O system calls.



Source: Understanding Lustre Internals



Lustre Components

Metadata Server (MDS)

- File operations (create, open, read etc.) require metadata stored on MDS
- Handles metadata requests - file lookups, file and directory attribute manipulation
- Maintains a transactional record of file system changes

Object Storage Server (OSS) and Object Storage Targets (OST)

- Each file is composed of data objects striped on one or more OSTs
- Responsible for actual file system I/O
- Responsible for interfacing with storage devices








Lustre Components

Lustre Client

- Queries MDS
- Retrieves the list of OSTs
- Sends request to the OSTs



Lustre Striping



- Stripe size
- Stripe count/width

Obj 1 => OST A
Obj 2 => OST B
Obj 3 => OST C
Obj 4 => OST D
Obj 5 => OST E



Lustre Striping Example

```
[pmalakar@cn364 testq]$ lfs setstripe -c 10 testmpio.out
[pmalakar@cn364 testq]$ lfs getstripe testmpio.out
testmpio.out
lmm_stripe_count: 10
lmm_stripe_size: 1048576
lmm_pattern: 1
lmm_layout_gen: 0
lmm_stripe_offset: 1
      obdidx      objid      objid      group
        1        4673479    0x474fc7        0
        0        4600893    0x46343d        0
       20        4551236    0x457244        0
        3        4701254    0x47bc46        0
       21        4479152    0x4458b0        0
       19        4696884    0x47ab34        0
        5        4704057    0x47c739        0
        7        4647142    0x46e8e6        0
       16        4640736    0x46cf0        0
       18        4595400    0x461ec8        0
```

Example: File striped across 10 OSTs. Each OST stores 1 MB objects.



Lustre Stripping Parameters

lfs setstripe -S <size> -c <count> filename

```
[pmalakar@cn364 testq]$ rm testmpio.out
[pmalakar@cn364 testq]$ time dd if=/dev/zero of=testmpio.out bs=10M count=1000
1000+0 records in
1000+0 records out
10485760000 bytes (10 GB) copied, 18.2025 s, 576 MB/s

real    0m18.205s
user    0m0.004s
sys     0m10.042s
[pmalakar@cn364 testq]$ rm testmpio.out
[pmalakar@cn364 testq]$ lfs setstripe -S 2M -c 10 testmpio.out
[pmalakar@cn364 testq]$ time dd if=/dev/zero of=testmpio.out bs=10M count=1000
1000+0 records in
1000+0 records out
10485760000 bytes (10 GB) copied, 10.4116 s, 1.0 GB/s

real    0m10.420s
user    0m0.003s
sys     0m10.406s
```



Striping Benefit

```
8 MB  Time 0.010138  
       Time 0.013419  
       Time 0.027182  
       Time 0.075958  
       Time 0.219819  
256 MB Time 0.333267
```

Stripe count = 1

```
8 MB  Time 0.020716  
       Time 0.025181  
       Time 0.035053  
       Time 0.063688  
       Time 0.220986  
256 MB Time 0.223855
```

Stripe count = 18



Striping Benefit

1.5 GB

```
Time 0.102997
Time 0.099457
Time 0.245411
Time 0.439744
Time 0.692367
Time 0.954562
```

Stripe count = 18

1.5 GB

```
Time 0.053866
Time 0.158203
Time 0.330782
Time 0.744611
Time 1.399156
Time 2.390323
```

Stripe count = 1

